



DISTRIBUTED SYSTEMS


with ZeroMQ and gevent


Jeff Lindsay
@progrum


Why distributed systems?


- Harness more CPUs and resources
- Run faster in parallel
- Tolerance of individual failures
- Better separation of concerns

Most web apps evolve into
distributed systems





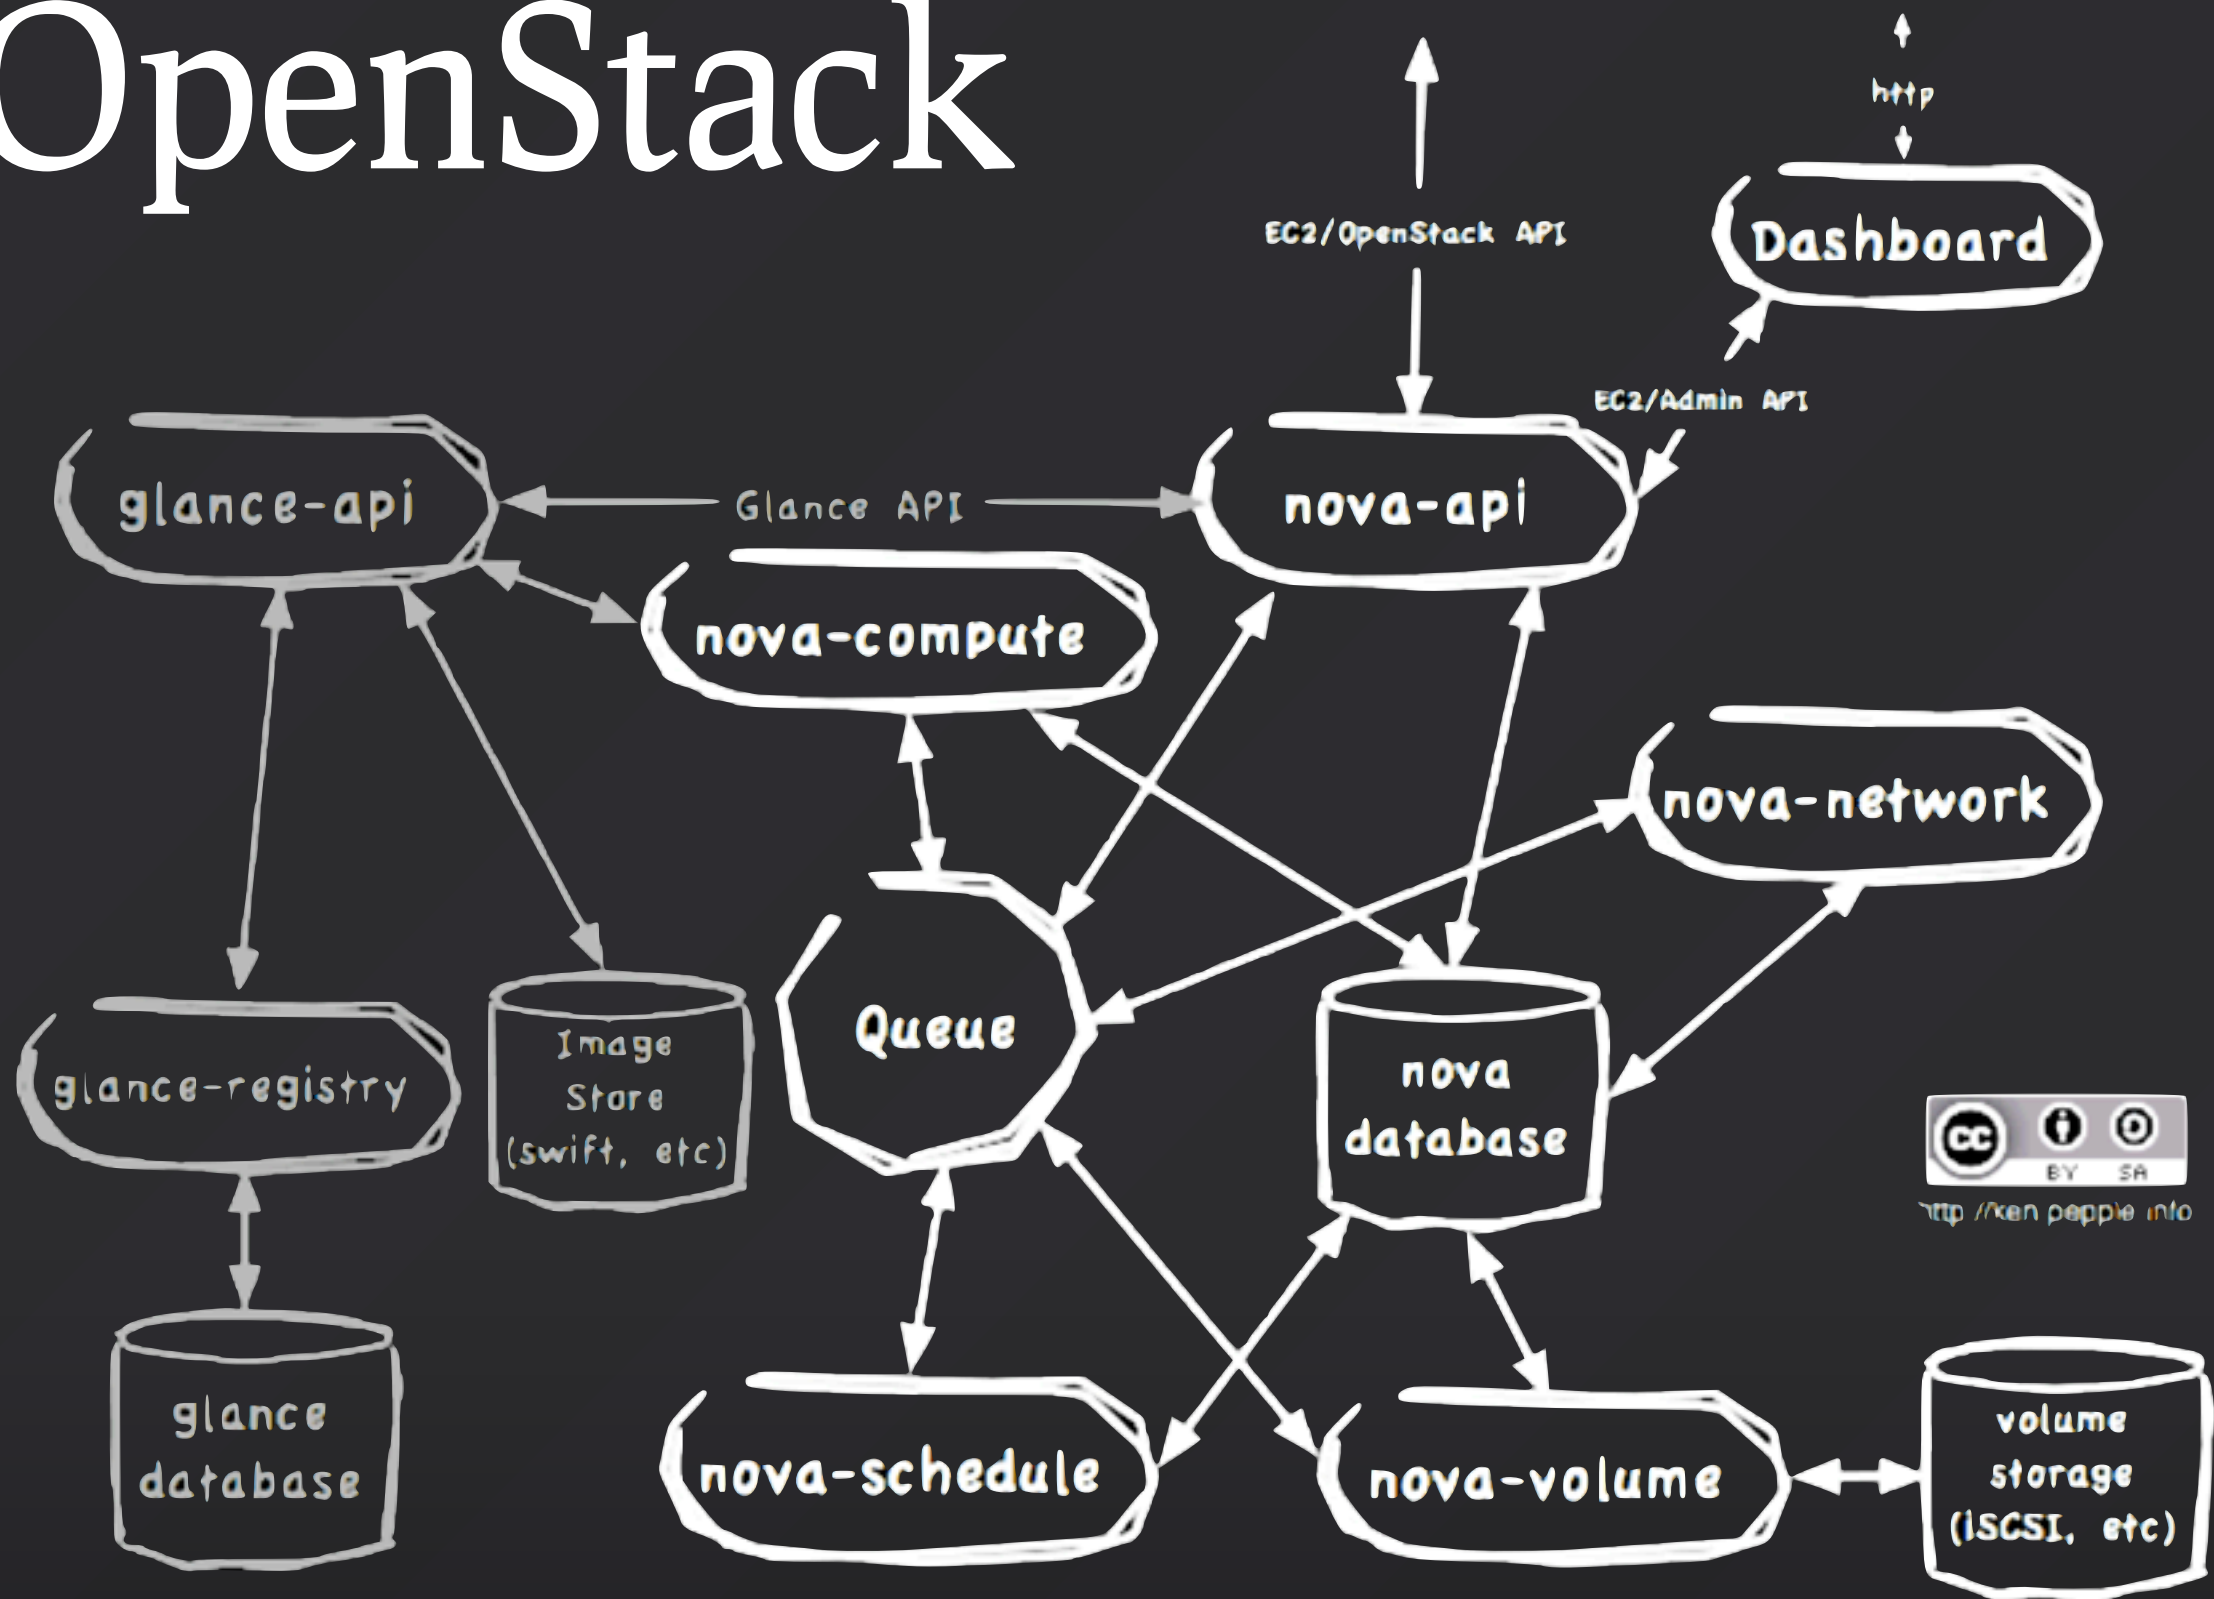








OpenStack





ZeroMQ + gevent


Two powerful and misunderstood tools

CONCURRENCY


Heart of Distributed Systems

Distributed computing
is just another flavor of
local concurrency

Multithreading



Distributed system



Concurrency models

- Execution model
 - Defines the “computational unit”
- Communication model
 - Means of sharing and coordination

Concurrency models

- Traditional multithreading
 - OS threads
 - Shared memory, locks, etc
- Async or Evented I/O
 - I/O loop + callback chains
 - Shared memory, futures
- Actor model
 - Shared nothing “processes”
 - Built-in messaging

Examples

- Erlang
 - Actor model
- Scala
 - Actor model
- Go
 - Channels, Goroutines
- Everything else (Ruby, Python, PHP, Perl, C/C++, Java)
 - Threading
 - Evented

Erlang is special.

Normally, the networking of distributed systems is tacked on to the local concurrency model.

MQ, RPC, REST, ...

Why not always use Erlang?

Why not always use Erlang?

- Half reasons
 - Weird/ugly language
 - Limited library ecosystem
 - VM requires operational expertise
 - Functional programming isn't mainstream

Why not always use Erlang?

- Half reasons
 - Weird/ugly language
 - Limited library ecosystem
 - VM requires operational expertise
 - Functional programming isn't mainstream
- Biggest reason
 - It's not always the right tool for the job



Service Oriented Architecture

Multiple languages
Heterogeneous cluster

RPC

RPC

- Client / server

RPC

- Client / server
- Mapping to functions

RPC

- Client / server
- Mapping to functions
- Message serialization

RPC

- Client / server
- Mapping to functions
- Message serialization

Poor abstraction of what you really want

What you want are tools to help you get distributed actor model concurrency like Erlang ... without Erlang.
Even better if they're decoupled and optional.

Rarely will you build an application as part of a distributed system that does not also need local concurrency.

COMMUNICATION MODEL

How do we unify communications in local concurrency
and distributed systems across languages?

EXECUTION MODEL

How do we get Erlang-style local concurrency without interfering with the language's idiomatic paradigm?

ZeroMQ

COMMUNICATION MODEL

Misconceptions

Misconceptions

- It's just another MQ, right?

Misconceptions

- It's just another MQ, right?
 - Not really.

Misconceptions

- It's just another MQ, right?
 - Not really.

Misconceptions

- It's just another MQ, right?
 - Not really.
- Oh, it's just sockets, right?

Misconceptions

- It's just another MQ, right?
 - Not really.
- Oh, it's just sockets, right?
 - Not really.

Misconceptions

- It's just another MQ, right?
 - Not really.
- Oh, it's just sockets, right?
 - Not really.

Misconceptions

- It's just another MQ, right?
 - Not really.
- Oh, it's just sockets, right?
 - Not really.
- Wait, isn't messaging a solved problem?


Misconceptions

- It's just another MQ, right?
 - Not really.
- Oh, it's just sockets, right?
 - Not really.
- Wait, isn't messaging a solved problem?
 - *sigh* ... maybe.

Regular Sockets


Regular Sockets

- Point to point




Regular Sockets

- Point to point
- Stream of bytes




Regular Sockets

- Point to point
- Stream of bytes
- Buffering




Regular Sockets

- Point to point
- Stream of bytes
- Buffering
- Standard API



Regular Sockets

- Point to point
- Stream of bytes
- Buffering
- Standard API
- TCP/IP or UDP, IPC



Messaging


Messaging

Messages are atomic



Messaging

Messages are atomic




Messaging

Messages are atomic




Messaging

Messages are atomic



Messaging

Messages are atomic



Messages can be routed




Messaging

Messages are atomic




Messages can be routed



Messaging

Messages are atomic




Messages can be routed



Messaging

Messages are atomic



Messages can be routed




Messages may sit around



Messaging

Messages are atomic



Messages can be routed




Messages may sit around



Messaging

Messages are atomic



Messages can be routed




Messages may sit around



Messaging

Messages are atomic



Messages can be routed




Messages may sit around



Messaging

Messages are atomic



Messages can be routed




Messages may sit around



Messaging

Messages are atomic



Messages can be routed




Messages may sit around



Messaging

Messages are atomic



Messages can be routed



Messages may sit around




Messages are delivered



Messaging

Messages are atomic



Messages can be routed



Messages may sit around

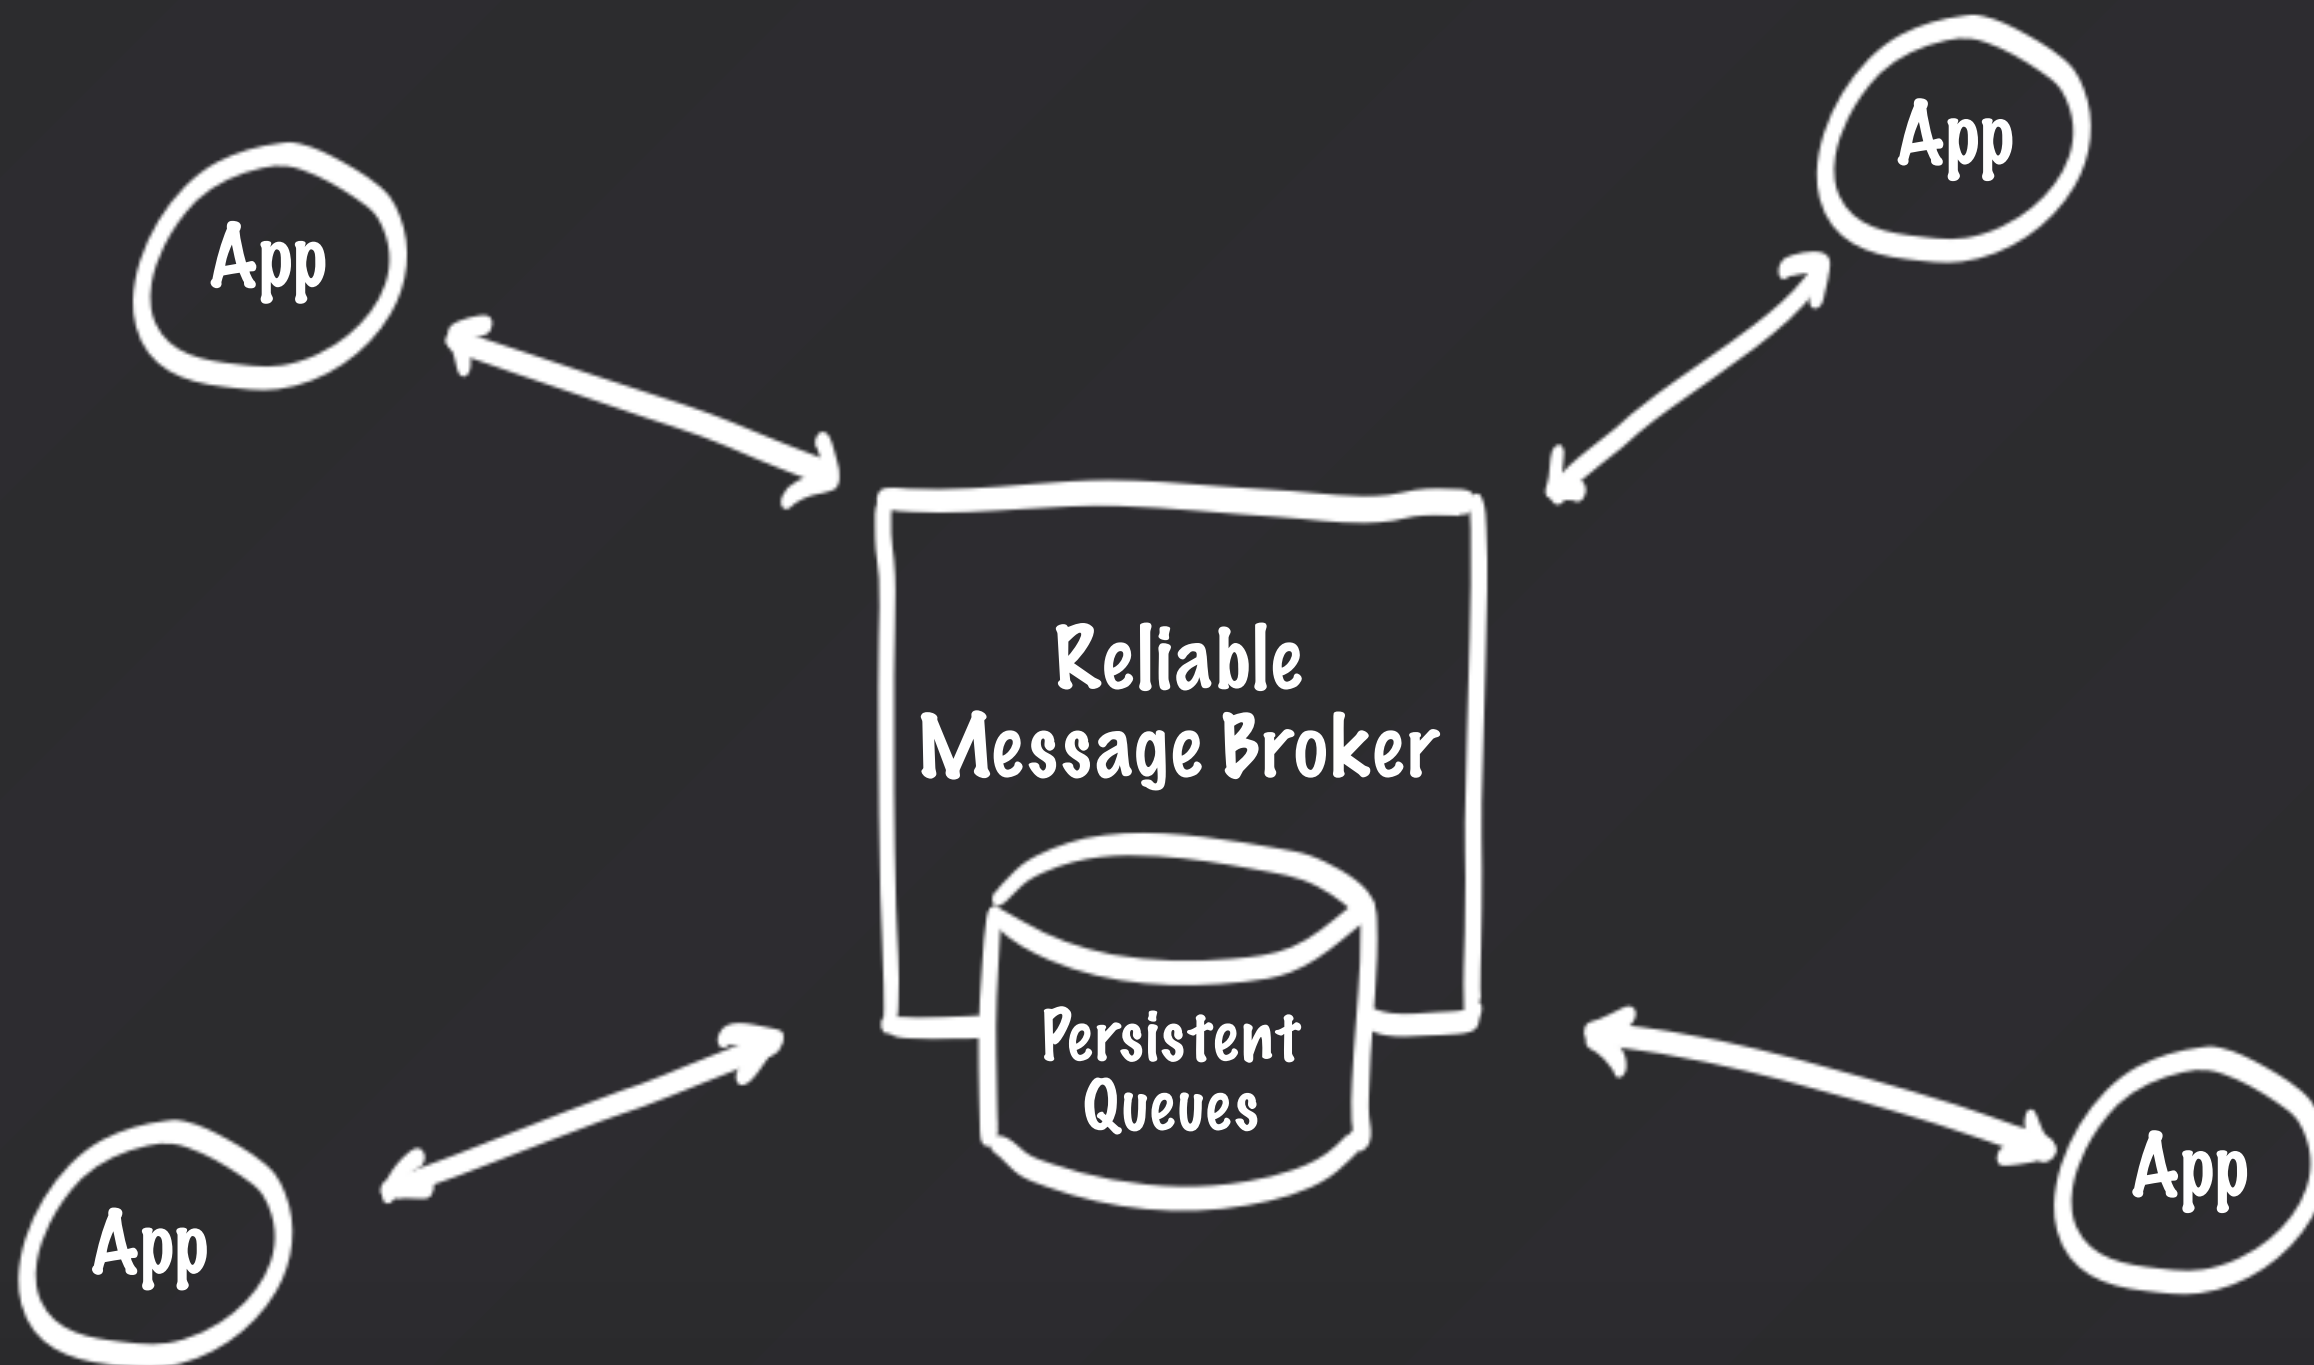



Messages are delivered



A faint, stylized illustration of a book with a face. The book is represented by a vertical rectangle with horizontal lines indicating pages. In the center of the book, there is a simple face with two eyes and a small mouth, giving it a personified appearance.

Rise of the Big MQ







AMQP



AMQP



AMQP



AMQP Recipes

AMQP Recipes

Work queues

Distributing tasks among workers



AMQP Recipes

Work queues

Distributing tasks among workers



Publish/Subscribe

Sending to many consumers at once



AMQP Recipes

Work queues

Distributing tasks among workers




Publish/Subscribe

Sending to many consumers at once



Routing

Receiving messages selectively



AMQP Recipes

Work queues

Distributing tasks among workers




Publish/Subscribe

Sending to many consumers at once



Routing

Receiving messages selectively



RPC

Remote procedure call implementation



Drawbacks of Big MQ

- Lots of complexity
- Queues are heavyweight
- HA is a challenge
- Poor primitives

Enter ZeroMQ

“Float like a butterfly, sting like a bee”

Echo in Python

Server

```
1 import zmq
2 context = zmq.Context()
3 socket = context.socket(zmq.REP)
4 socket.bind("tcp://127.0.0.1:5000")
5
6 while True:
7     msg = socket.recv()
8     print "Received", msg
9     socket.send(msg)
```

Client

```
1 import zmq
2 context = zmq.Context()
3 socket = context.socket(zmq.REQ)
4 socket.connect("tcp://127.0.0.1:5000")
5
6 for i in range(10):
7     msg = "msg %s" % i
8     socket.send(msg)
9     print "Sending", msg
10    reply = socket.recv()
```

Echo in Ruby

Server

```
1 require "zmq"
2 context = ZMQ::Context.new(1)
3 socket = context.socket(ZMQ::REP)
4 socket.bind("tcp://127.0.0.1:5000")
5
6 loop do
7     msg = socket.recv
8     puts "Received #{msg}"
9     socket.send(msg)
10 end
```

Client

```
1 require "zmq"
2 context = ZMQ::Context.new(1)
3 socket = context.socket(ZMQ::REQ)
4 socket.connect("tcp://127.0.0.1:5000")
5
6 (0..10).each do |i|
7     msg = "msg #{i}"
8     socket.send(msg)
9     puts "Sending #{msg}"
10    reply = socket.recv
11 end
```


Echo in PHP

Server

```
1 <?php
2 $context = new ZMQContext();
3 $socket = $context->getSocket(ZMQ::SOCKET_REP);
4 $socket->bind("tcp://127.0.0.1:5000");
5
6 while (true) {
7     $msg = $socket->recv();
8     echo "Received {$msg}";
9     $socket->send($msg);
10 }
11 ?>
```

Client

```
1 <?php
2 $context = new ZMQContext();
3 $socket = $context->getSocket(ZMQ::SOCKET_REQ);
4 $socket->connect("tcp://127.0.0.1:5000");
5
6 foreach (range(0, 9) as $i) {
7     $msg = "msg {$i}";
8     $socket->send($msg);
9     echo "Sending {$msg}";
10    $reply = $socket->recv();
11 }
12 ?>
```

Bindings

ActionScript, Ada, Bash, Basic, C, Chicken
Scheme, Common Lisp, C#, C++, D, Erlang,
F#, Go, Guile, Haskell, Haxe, Java, JavaScript,
Lua, Node.js, Objective-C, Objective Caml,
ooc, Perl, PHP, Python, Racket, REBOL,
Red, Ruby, Smalltalk

Plumbing


Plumbing



Plumbing



Plumbing



Plumbing



Plumbing

- inproc
- ipc
- tcp
- multicast



Plumbing


- inproc
- ipc
- tcp
- multicast



```
socket.bind("tcp://localhost:5560")
socket.bind("ipc:///tmp/this-socket")
socket.connect("tcp://10.0.0.100:9000")
socket.connect("ipc:///tmp/another-socket")
socket.connect("inproc://another-socket")
```

Plumbing


- inproc
- ipc
- tcp
- multicast



```
socket.bind("tcp://localhost:5560")
socket.bind("ipc:///tmp/this-socket")
socket.connect("tcp://10.0.0.100:9000")
socket.connect("ipc:///tmp/another-socket")
socket.connect("inproc://another-socket")
```

Plumbing

- inproc
- ipc
- tcp
- multicast



```
socket.bind("tcp://localhost:5560")  
socket.bind("ipc:///tmp/this-socket")  
socket.connect("tcp://10.0.0.100:9000")  
socket.connect("ipc:///tmp/another-socket")  
socket.connect("inproc://another-socket")
```

Message Patterns


Message Patterns

Request-Reply




Message Patterns

Request-Reply




Message Patterns

Request-Reply




Message Patterns

Request-Reply




Message Patterns

Request-Reply




Message Patterns

Request-Reply




Publish-Subscribe




Message Patterns


Request-Reply



Publish-Subscribe




Push-Pull (Pipelining)




Message Patterns


Request-Reply



Publish-Subscribe




Push-Pull (Pipelining)




Message Patterns


Request-Reply



Publish-Subscribe




Push-Pull (Pipelining)




Message Patterns


Request-Reply



Publish-Subscribe




Push-Pull (Pipelining)




Message Patterns


Request-Reply



Publish-Subscribe



Push-Pull (Pipelining)



Pair



Devices

Queue

Forwarder

Streamer



Design architectures around devices.

Devices

Queue

Forwarder

Streamer



Design architectures around devices.

Devices

Queue

Forwarder

Streamer



Design architectures around devices.

Devices

Queue

Forwarder

Streamer



Design architectures around devices.

Performance

Performance

- Orders of magnitude faster than most MQs

Performance

- Orders of magnitude faster than most MQs
- Higher throughput than raw sockets

Performance

- Orders of magnitude faster than most MQs
- Higher throughput than raw sockets
 - Intelligent message batching

Performance

- Orders of magnitude faster than most MQs
- Higher throughput than raw sockets
 - Intelligent message batching
 - Edge case optimizations

Concurrency?

"Come for the messaging, stay for the easy concurrency"

Hintjens' Law of Concurrency

$$E = MC^2$$

E is effort, the pain that it takes


M is mass, the size of the code

C is conflict, when C threads collide

Hintjens' Law of Concurrency



Hintjens' Law of Concurrency



Hintjens' Law of Concurrency

ZeroMQ:

$E=MC^2$, FOR $C=1$

ZeroMQ

- Easy ... familiar socket API
- Cheap ... lightweight queues in a library
- Fast ... higher throughput than raw TCP
- Expressive ... maps to your architecture

Messaging toolkit for concurrency and distributed systems.

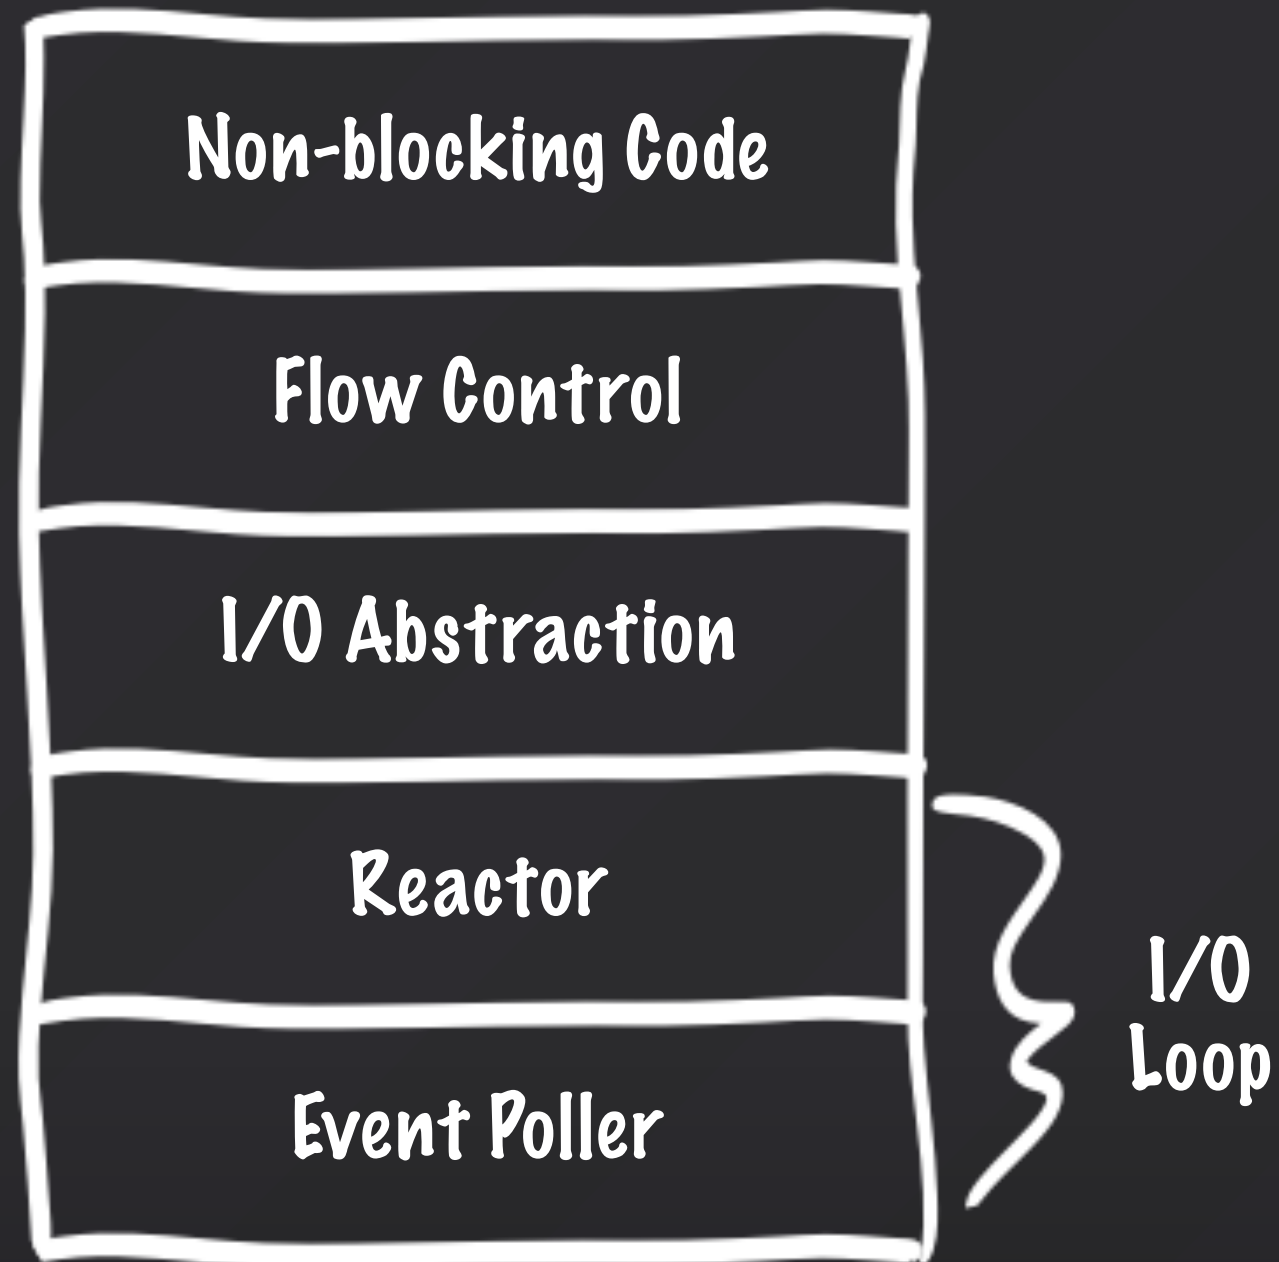
gevent

EXECUTION MODEL

Threading vs Evented

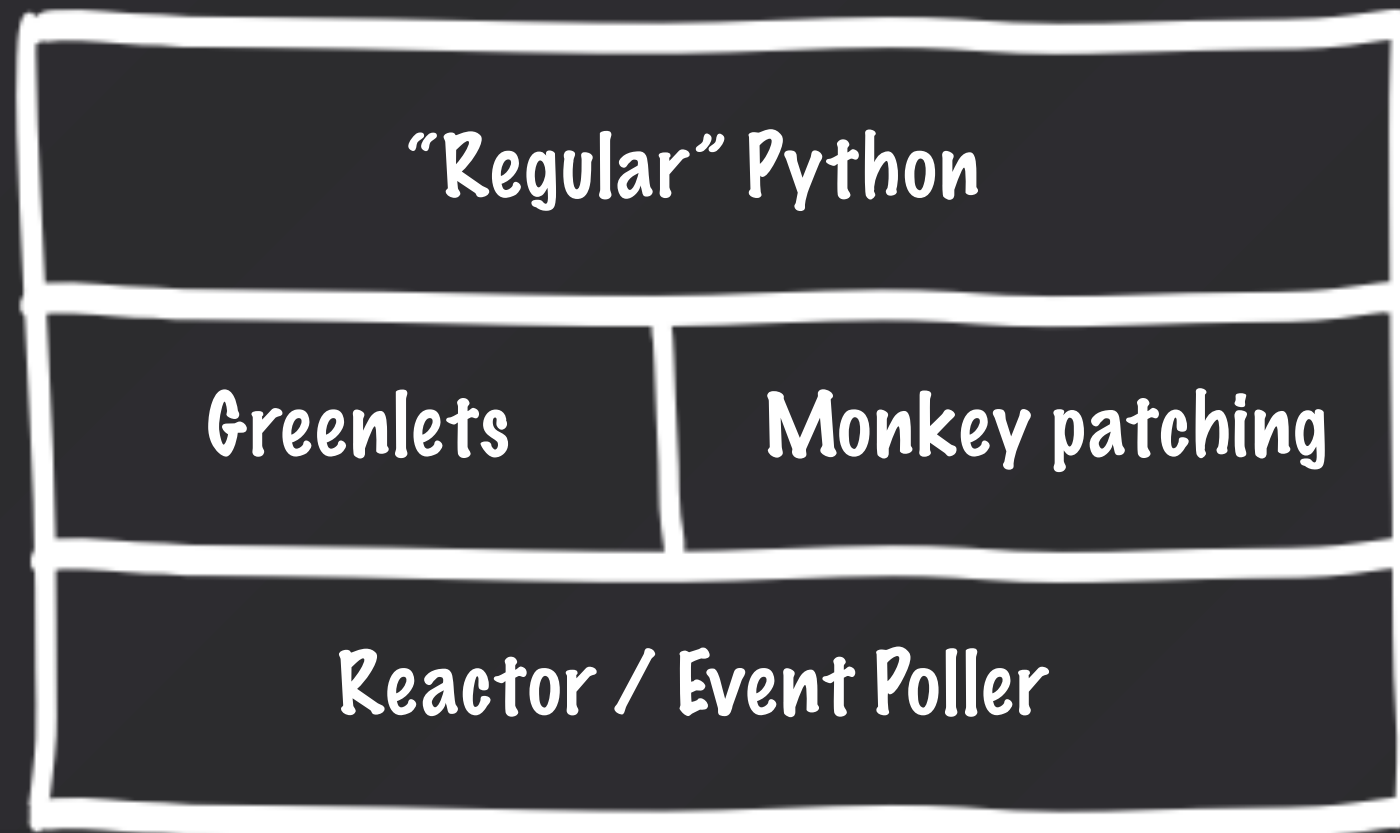
Evented seems to be preferred for scalable I/O applications

Evented Stack



```
1 def lookup(country, search_term):
2     main_d = defer.Deferred()
3
4     def first_step():
5         query = "http://www.google.%s/search?q=%s" % (country, search_term)
6         d = getPage(query)
7         d.addCallback(second_step, country)
8         d.addErrback(failure, country)
9
10    def second_step(content, country):
11        m = re.search('<div id="?res.*?href="(P<url>http://[^"]+)"',
12                      content, re.DOTALL)
13        if not m:
14            main_d.callback(None)
15            return
16        url = m.group('url')
17        d = getPage(url)
18        d.addCallback(third_step, country, url)
19        d.addErrback(failure, country)
20
21    def third_step(content, country, url):
22        m = re.search("<title>(.*?)</title>", content)
23        if m:
24            title = m.group(1)
25            main_d.callback(dict(url = url, title = title))
26        else:
27            main_d.callback(dict(url=url, title="{not-specified}"))
28
29    def failure(e, country):
30        print ".%s FAILED: %s" % (country, str(e))
31        main_d.callback(None)
32
33    first_step()
34    return main_d
```

gevent



Green threads

“Threads” implemented in user space (VM, library)

Monkey patching

socket, ssl, threading, time

Twisted

Twisted




~400 modules

gevent

25 modules

Performance


Performance



Performance


Successful Connection Rate

on an increasing amount of requests (more is better)




Response Time

on an increasing amount of requests (less is better)




Error Rate

on an increasing amount of requests (less is better)




Successful Reply Rate

on an increasing amount of requests (more is better)




Response Time

on an increasing amount of requests (less is better)



Error Rate

on an increasing amount of requests (less is better)



Building a Networking App

```
1  #===
2  # 1. Basic gevent TCP server
3
4  from gevent.server import StreamServer
5
6  def handle_tcp(socket, address):
7      print 'new tcp connection!'
8      while True:
9          socket.send('hello\n')
10         gevent.sleep(1)
11
12  tcp_server = StreamServer(('127.0.0.1', 1234), handle_tcp)
13  tcp_server.serve_forever()
```

```
1 #===
2 # 2. Basic gevent TCP server and WSGI server
3
4 from gevent.pywsgi import WSGIServer
5 from gevent.server import StreamServer
6
7 def handle_http(env, start_response):
8     start_response('200 OK', [('Content-Type', 'text/html')])
9     print 'new http request!'
10    return ["hello world"]
11
12 def handle_tcp(socket, address):
13     print 'new tcp connection!'
14     while True:
15         socket.send('hello\n')
16         gevent.sleep(1)
17
18 tcp_server = StreamServer(('127.0.0.1', 1234), handle_tcp)
19 tcp_server.start()
20
21 http_server = WSGIServer(('127.0.0.1', 8080), handle_http)
22 http_server.serve_forever()
```

```
1  from gevent.pywsgi import WSGIServer
2  from gevent.server import StreamServer
3  from gevent.socket import create_connection
4
5  def handle_http(env, start_response):
6      start_response('200 OK', [('Content-Type', 'text/html')])
7      print 'new http request!'
8      return ["hello world"]
9
10 def handle_tcp(socket, address):
11     print 'new tcp connection!'
12     while True:
13         socket.send('hello\n')
14         gevent.sleep(1)
15
16 def client_connect(address):
17     sockfile = create_connection(address).makefile()
18     while True:
19         line = sockfile.readline() # returns None on EOF
20         if line is not None:
21             print "<<<", line,
22         else:
23             break
24
25 tcp_server = StreamServer(('127.0.0.1', 1234), handle_tcp)
26 tcp_server.start()
27
28 gevent.spawn(client_connect, ('127.0.0.1', 1234))
29
30 http_server = WSGIServer(('127.0.0.1', 8080), handle_http)
31 http_server.serve_forever()
```

```
1  from gevent.pywsgi import WSGIServer
2  from gevent.server import StreamServer
3  from gevent.socket import create_connection
4
5  def handle_http(env, start_response):
6      start_response('200 OK', [('Content-Type', 'text/html')])
7      print 'new http request!'
8      return ["hello world"]
9
10 def handle_tcp(socket, address):
11     print 'new tcp connection!'
12     while True:
13         socket.send('hello\n')
14         gevent.sleep(1)
15
16 def client_connect(address):
17     sockfile = create_connection(address).makefile()
18     while True:
19         line = sockfile.readline() # returns None on EOF
20         if line is not None:
21             print "<<<", line,
22         else:
23             break
24
25 tcp_server = StreamServer(('127.0.0.1', 1234), handle_tcp)
26 http_server = WSGIServer(('127.0.0.1', 8080), handle_http)
27 greenlets = [
28     gevent.spawn(tcp_server.serve_forever),
29     gevent.spawn(http_server.serve_forever),
30     gevent.spawn(client_connect, ('127.0.0.1', 1234)),
31 ]
32 gevent.joinall(greenlets)
```

ZeroMQ in gevent?



😊 traviscline / gevent-zeromq

Unwatch

Fork

79

18

Source

Commits

Network

Pull Requests (3)

Issues (7)

Graphs

Branch: master

Switch Branches (1)

Switch Tags (6)

Branch List

pyzeromq gevent compat. lib — [Read more](#)

Downloads

HTTP

Git Read-Only

<https://github.com/traviscline/gevent-zeromq.git>



Read-Only access

Clone in Mac

bumped version number to represent _state_event refactor



traviscline authored June 29, 2011

commit ead6ce80f4

gevent-zeromq /

name	age	message	history
examples/	January 25, 2011	changed to respect edge-triggered nature of zmq.FD... [traviscline]	
gevent_zeromq/	June 29, 2011	_state_event refactor [traviscline]	
.gitignore	April 19, 2011	updated setup.py in prep for pypi [traviscline]	
AUTHORS	March 02, 2011	added AUTHORS [traviscline]	
LICENSE	January 25, 2011	added LICENSE, added to README(.rst) [traviscline]	
MANIFEST.in	April 20, 2011	Changed to include core.c in MANIFEST.in so buildi... [traviscline]	
		Slight comment wording updates, version bump. [traviscline]	

```
1  from gevent import spawn
2  from gevent_zeromq import zmq
3
4  context = zmq.Context()
5
6  def serve():
7      socket = context.socket(zmq.REP)
8      socket.bind("tcp://localhost:5559")
9      while True:
10         message = socket.recv()
11         print "Received request: ", message
12         socket.send("World")
13
14  server = spawn(serve)
15
16  def client():
17      socket = context.socket(zmq.REQ)
18      socket.connect("tcp://localhost:5559")
19      for request in range(10):
20         socket.send("Hello")
21         message = socket.recv()
22         print "Received reply ", request, "[", message, "]"
23
24  spawn(client).join()
```

Actor model?

Easy to implement, in whole or in part,
optionally with ZeroMQ

1 Answer

active

oldest

votes

▲
2



To make actors with `gevent`, use a `Greenlet` subclass with embedded `gevent.queue.Queue` instance used as an inbox. To read a message from the inbox, simply `get()` from the queue. To send a message to an actor, `put` it into that actor's queue.

[Read about subclassing Greenlet here.](#)

If you need help with writing the Actor class, feel free to [ask the mailing list](#).

[link](#) | [edit](#) | [flag](#)

answered **Aug 8 '10 at 16:36**

community wiki
[Denis Bilenko](#)

Thank you very much! I will dive into `gevent`, and try it extensively – [daitangio](#) Aug 9 '10 at 12:53

[add comment](#)

What is gevent missing?

What is gevent missing?

- Documentation

What is gevent missing?

- Documentation
- Application framework



gservice

Application framework for gevent


```
1 from gevent.pywsgi import WSGIServer
2 from gevent.server import StreamServer
3 from gevent.socket import create_connection
4
5 def handle_http(env, start_response):
6     start_response('200 OK', [('Content-Type', 'text/html')])
7     print 'new http request!'
8     return ["hello world"]
9
10 def handle_tcp(socket, address):
11     print 'new tcp connection!'
12     while True:
13         socket.send('hello\n')
14         gevent.sleep(1)
15
16 def client_connect(address):
17     sockfile = create_connection(address).makefile()
18     while True:
19         line = sockfile.readline() # returns None on EOF
20         if line is not None:
21             print "<<<", line,
22         else:
23             break
24
25 tcp_server = StreamServer(('127.0.0.1', 1234), handle_tcp)
26 http_server = WSGIServer(('127.0.0.1', 8080), handle_http)
27 greenlets = [
28     gevent.spawn(tcp_server.serve_forever),
29     gevent.spawn(http_server.serve_forever),
30     gevent.spawn(client_connect, ('127.0.0.1', 1234)),
31 ]
32 gevent.joinall(greenlets)
```

```
1  from gevent.pywsgi import WSGIServer
2  from gevent.server import StreamServer
3  from gevent.socket import create_connection
4
5  from gservice.core import Service
6
7  def handle_http(env, start_response):
8      start_response('200 OK', [('Content-Type', 'text/html')])
9      print 'new http request!'
10     return ["hello world"]
11
12 def handle_tcp(socket, address):
13     print 'new tcp connection!'
14     while True:
15         socket.send('hello\n')
16         gevent.sleep(1)
17
18 def client_connect(address):
19     sockfile = create_connection(address).makefile()
20     while True:
21         line = sockfile.readline() # returns None on EOF
22         if line is not None:
23             print "<<<", line,
24         else:
25             break
26
27 app = Service()
28 app.add_service(StreamServer(('127.0.0.1', 1234), handle_tcp))
29 app.add_service(WSGIServer(('127.0.0.1', 8080), handle_http))
30 app.add_service(TcpClient(('127.0.0.1', 1234), client_connect))
31 app.serve_forever()
```

```
1 from gservice.core import Service
2 from gservice.config import Setting
3
4 class MyApplication(Service):
5     http_port = Setting('http_port')
6     tcp_port = Setting('tcp_port')
7     connect_address = Setting('connect_address')
8
9     def __init__(self):
10         self.add_service(WSGIServer(('127.0.0.1', self.http_port), self.handle_http))
11         self.add_service(StreamServer(('127.0.0.1', self.tcp_port), self.handle_tcp))
12         self.add_service(TcpClient(self.connect_address, self.client_connect))
13
14     def client_connect(self, address):
15         sockfile = create_connection(address).makefile()
16         while True:
17             line = sockfile.readline() # returns None on EOF
18             if line is not None:
19                 print "<<<", line,
20             else:
21                 break
22
23     def handle_tcp(self, socket, address):
24         print 'new tcp connection!'
25         while True:
26             socket.send('hello\n')
27             gevent.sleep(1)
28
29     def handle_http(self, env, start_response):
30         start_response('200 OK', [('Content-Type', 'text/html')])
31         print 'new http request!'
32         return ["hello world"]
```

```
1 # example.conf.py
2
3 pidfile = 'example.pid'
4 logfile = 'example.log'
5 http_port = 8080
6 tcp_port = 1234
7 connect_address = ('127.0.0.1', 1234)
8
9 def service():
10     from example import MyApplication
11     return MyApplication()
```

```
# Run in the foreground
gservice -C example.conf.py
```

```
# Start service as daemon
gservice -C example.conf.py start
```

```
# Control service
gservice -C example.conf.py restart
gservice -C example.conf.py reload
gservice -C example.conf.py stop
```

```
# Run with overriding configuration
gservice -C example.conf.py -X 'http_port = 7070'
```

Generalizing

gevent proves a model that can be implemented in almost
any language that can implement an evented stack

gevent

- Easy ... just normal Python
- Small ... only 25 modules
- Fast ... top performing server
- Compatible ... works with most libraries

Futuristic evented platform for network applications.

Raiden

Lightning fast, scalable messaging

<https://github.com/progrium/raiden>

Concurrency models

- Traditional multithreading
- Async or Evented I/O
- Actor model

Conclusion

Two very simple, but very powerful tools
for distributed / concurrent systems

Thanks

@progrium