



A new programming model for the cloud

Cyprien Noel

ObjectFabric

- Founded February 2011 in Mountain View



- Technology derives from financial software




- 4-year development
- Open source, 3 contributors
- .NET apps, Web sites
- Previously xstm.org

Plan


1. Challenges: Concurrency & Distribution
2. Transactional Memories
3. Distributed Transactional Memories
4. Our solution
 - Overview
 - Code samples
5. Thoughts (divagations!)
6. Questions

Challenges

Concurrency



Distribution



Models

Concurrency

- Locks
- "Lock Free"
 - CAS, Memory Barriers
- Message Passing
 - Message Loops, Actors
- Functional Programming
- Data Parallelism, GPUs
- Transactions
 - Transactional Memories

Distribution

- Message Passing
 - RPC, SOA
 - Queues, Actors
- Transactions
 - Databases

Transactional Memory

- Using transactions when modifying memory
 - Like a database but for in-memory objects
 - Sort of generalized CAS

Strengths


Avoids locks pitfalls like deadlocks
Composable
Rollback, no error recovery code

Weaknesses

Performance
Interactions with IO, locks
Debugging, language support

Transactional Memory

- More than 500 papers since first (Tim Harris, 2010)
 - *Transactional Memory: Architectural Support for Lock-Free Data Structures*, Maurice Herlihy, J. Eliot B. Moss, 1993



2004 to 2007 – Lots talks & impl.

2008 – First negative papers

2009 – Sun cancels Rock

2010 – Microsoft stops STM.net

- Clojure, Multiverse, Scala, Fortress, TBoost.STM, CloudTM
- IBM BlueGene/Q processor

Transactional Memory

- Our implementation
 - For each tradeoff: pick easiest for developers
 - MVCC, View-Isolation & Opacity
 - No spurious aborts: can do IO
 - Strong isolation through type system
 - Lock free: guaranteed progress (starvation possible)
 - No byte code, JVM or lang. modif: easy debugging
- Modeled after source control systems
 - Intuitive to developers


Thread 1

Time





Thread 2

Take snapshot

Update




Merge



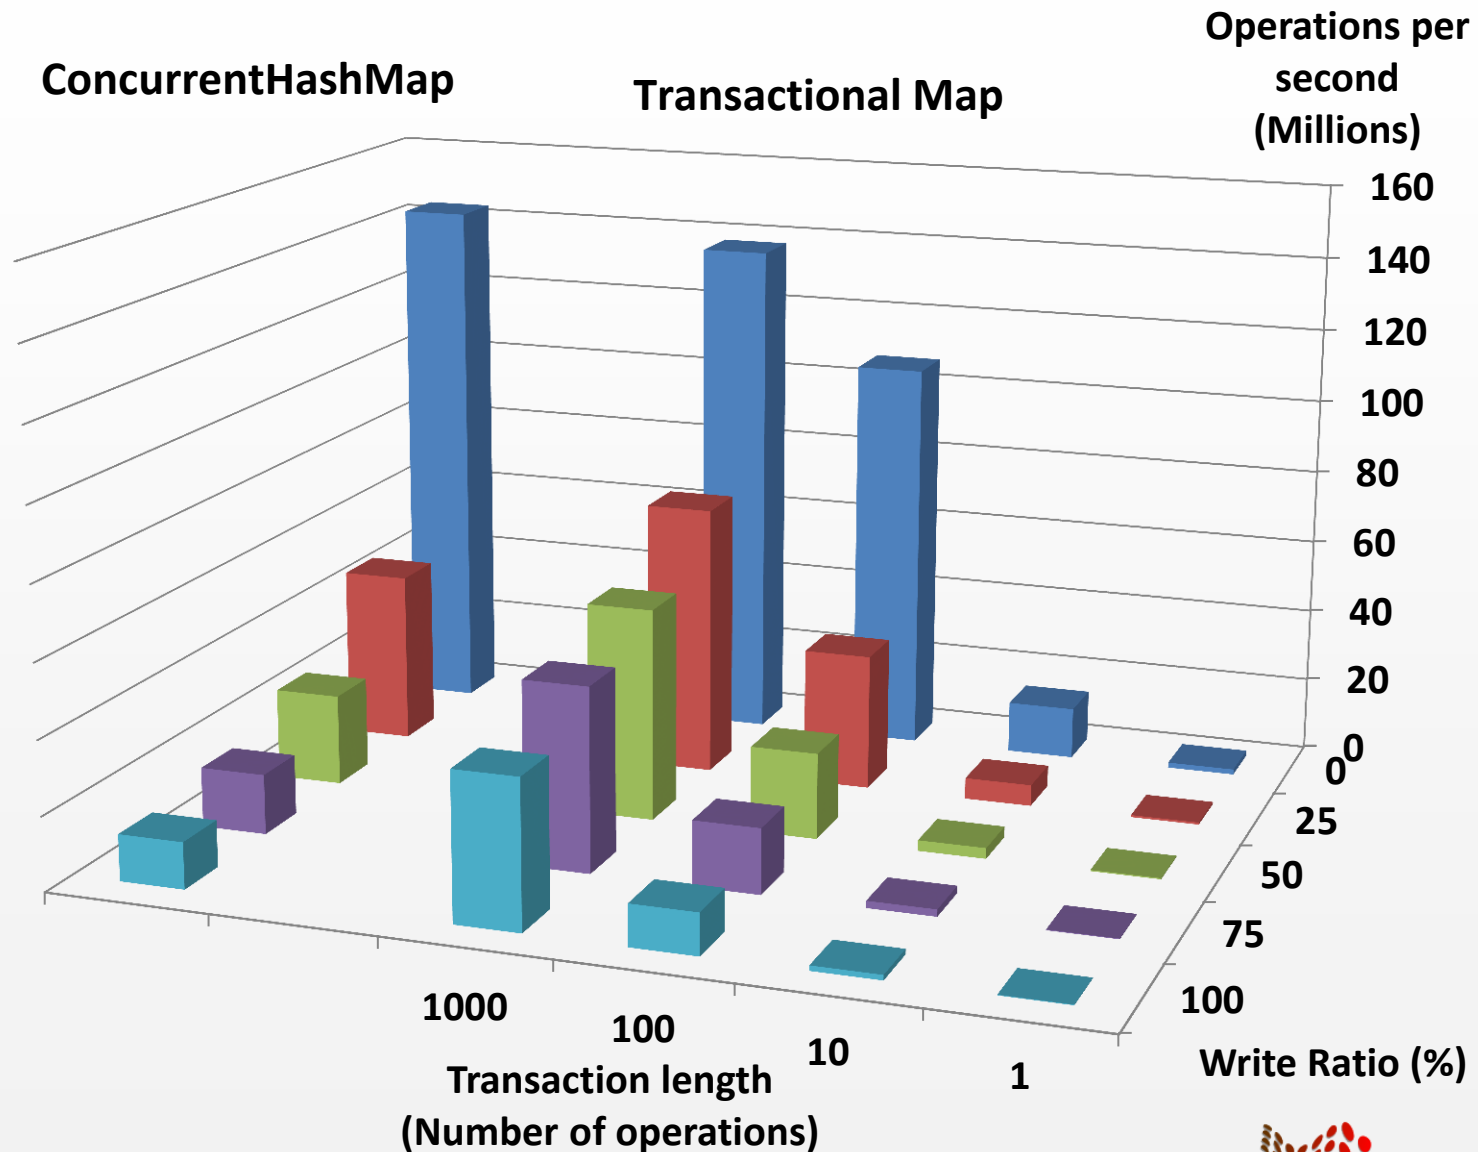
Take snapshot

Update



Merge

Performance



What else can be done with a TM?

- Transactions have read & write sets
 - Iterate write set -> e.g. logging, change notification

Thread A

```
// Register a listener on the map
map.addListener(new KeyListener() {

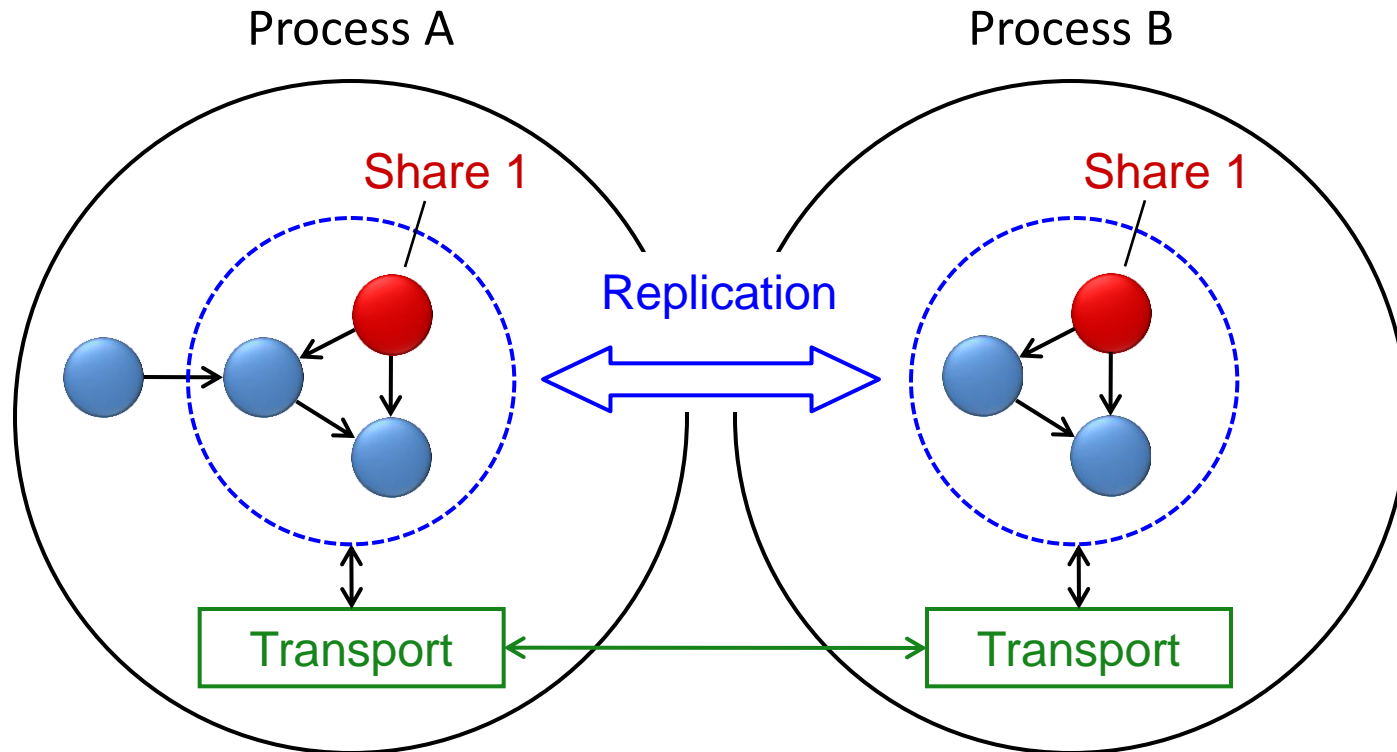
    // A key has been put in the map
    public void onPut(Object key) {
        // Returns "value"
        map.get(key);
    }
}
```

Thread B


```
map.put("key", "value");
```

- Persist writes to a store
- Serialize and commit transactions remotely
- Extensions have small impact on performance
 - Extensions can skip to last update

Distributed Transactional Memory



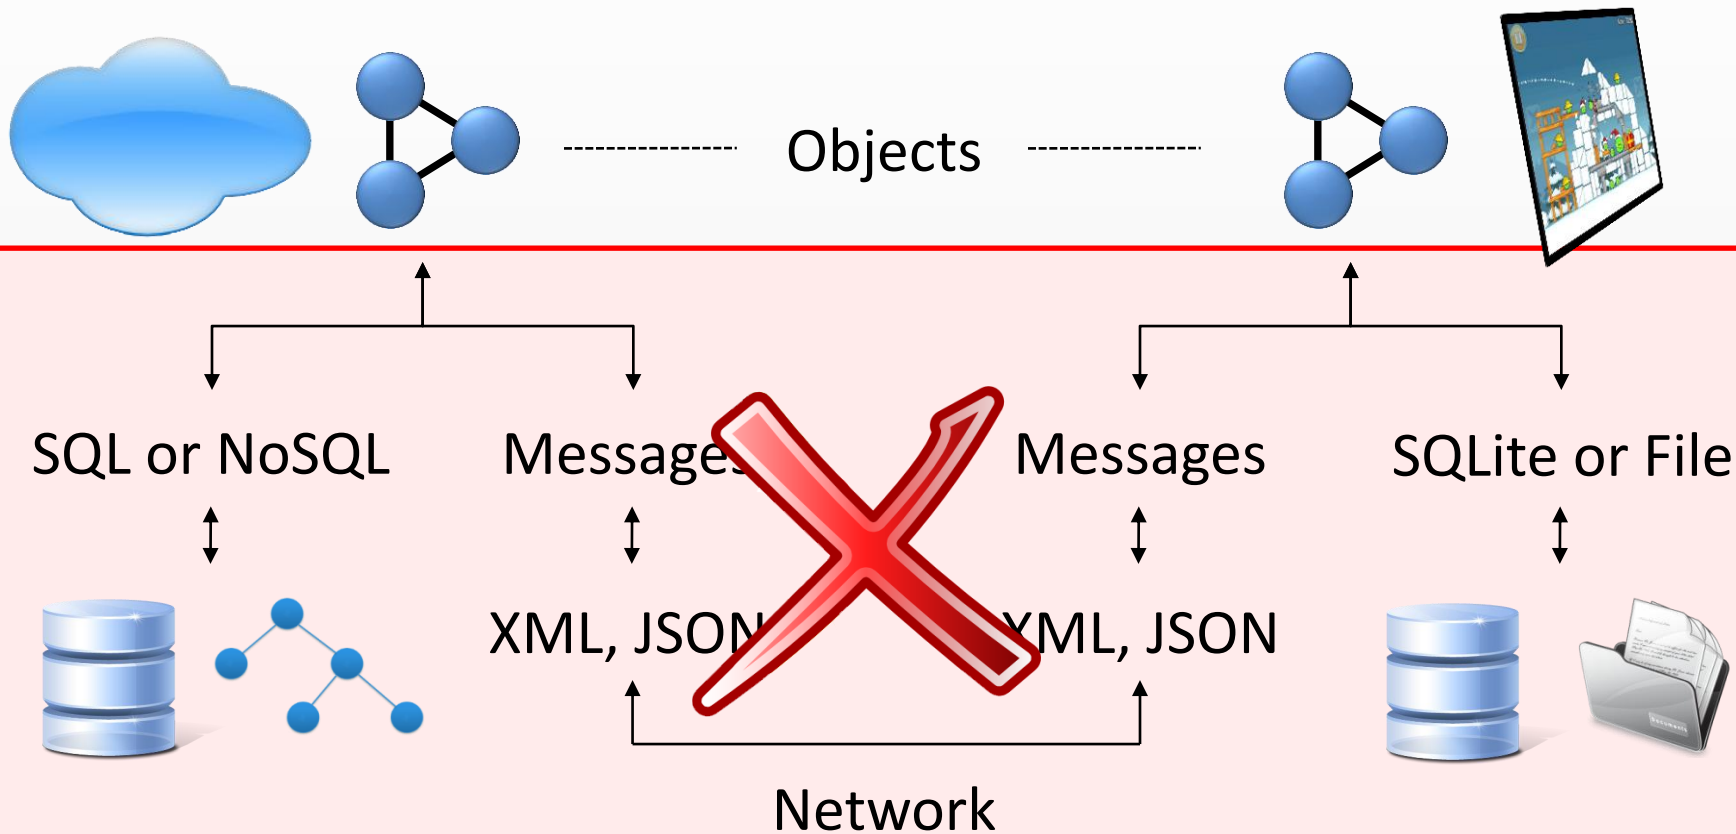
Traditional Architecture



Traditional Architecture

Cloud

Application




Distributed Transactional Memory

Cloud



Application



Developers only declare intents

- E.g. "Replicate this with server"
- E.g. "Make this durable"

Benefits


- A lot less or no code
- 10X faster than REST

Distributed Transactional Memory

- Lots of interest over past 3 or 4 years
 - Many papers and implementations from academia
 - **Fénix, D²STM** (*Technical University of Lisbon*)
 - **Arjuna** (*Newcastle University*)
 - **DMV** (*University of Toronto*)
 - **Isis²** (*Cornell University*)
 - **DiSTM** (*University of Manchester*)
 - EU recently funded the **CloudTM** initiative in which **RedHat** is involved



ObjectFabric



- Lightweight Java server
- Web, mobile & desktop
- Advanced stuff opt. in
- No configuration
- Java SE (Not EE, 1 jar)
- Secure (TLS, Shiro)
- Scalable (NIO)
- Fast (>100K message/s)

Get Started

Java

- objectfabric.examples [objectfabric master]
 - src
 - part01.helloworld
 - Client.java
 - Server.java
 - part02.objectmodel
 - part03.replication
 - part04.store
 - part05.stm
 - part06.methods
 - part07.tls
 - part08.versioning
 - part09.images
 - part10.trading
 - part11.bench
 - .settings
 - temp
 - .classpath
 - .gitignore
 - .project

Google Web Toolkit

- of4gwt.examples01.helloworld [objectfabric master]
 - src
 - helloworld
 - client
 - Main.java
 - HelloWorld.gwt.xml
 - .settings
 - war
 - WEB-INF
 - HelloWorld.html
 - .classpath
 - .gitignore
 - .project
 - of4gwt.examples02.objectmodel [objectfabric master]
 - of4gwt.examples03.images [objectfabric master]

Server

```
// Open a port and listen for connections (Socket, or Comet)
SocketServer server = new SocketServer(8080);

server.setCallback(new Callback() {

    public void onConnection(SocketConnection session) {
        // Client connected, send an object
        session.send(new TMap ());
    }

    public void onDisconnection(SocketConnection session, Throwable t) {
        // Client disconnected
    }

    public void onReceived(SocketConnection session, Object object) {
        // Received an object from client
    }
}
```

```
client = new SocketClient("myCompany.com", 8080);

client.setCallback(new Callback() {

    public void onReceived(Object object) {
        // Received object from server
        TMap map = (TMap) object;
    }

    public void onDisconnected() {
        // Lost connection
    }
});

client.connectAsync(null);
```

Client

- Supported objects
 - Basic objects: "Immutable" (String, Integer, Date etc.)
 - Transactional collections: TMap, TList, TSet, LazyMap
 - Custom objects (Java or C# generated from description in XML or SQL)
- Once a collection or custom object is sent to another process, its state is replicated automatically
- Listeners work on replicas as in previous example

Process A (Client or server)

```
// Register a listener on the map
map.addListener(new KeyListener() {

    // A key has been put in the map
    public void onPut(Object key) {
        // Returns "value"
        map.get(key);
    }
}
```

Process B


```
map.put("key", "value");
```














Transactional Memory

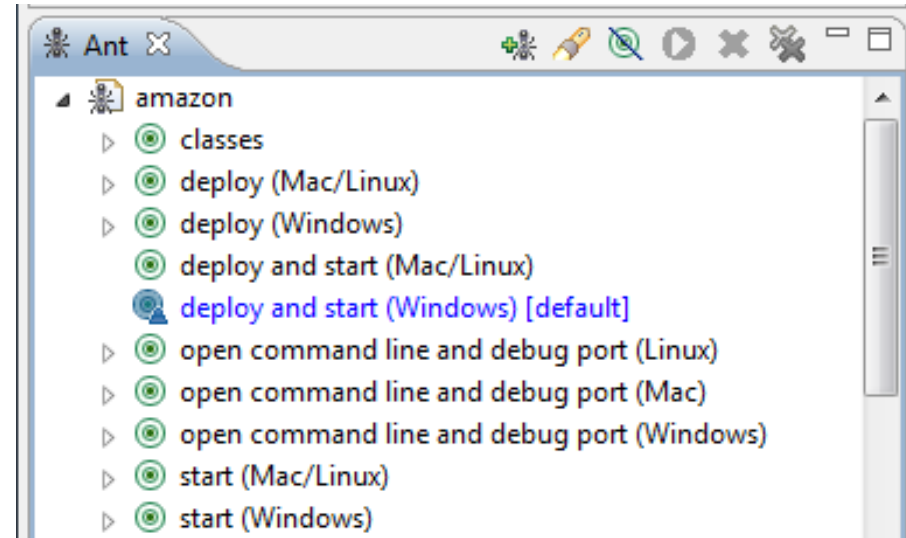
- Optional: transaction are started and committed automatically if needed when accessing collections or custom objects
- E.g. 2 objects **a** and **b**. This runs in parallel on several threads:

```
for (int i = 0; i < WRITE_COUNT; i++) {  
    Transaction.run(new Runnable() {  
  
        public void run() {  
            assert a.getInt() == b.getInt();  
  
            a.setInt(a.getInt() + 1);  
            b.setInt(b.getInt() + 1);  
        }  
    });  
}
```

Deploy to Amazon EC2

 > objectfabric.examples.amazon [objectfabric master]

-  .settings
-  cwRsync
-  deployed
 -  classes
 -  lib
-  boot.sh
-  install_daemon.sh
-  start.sh
-  > .gitignore
-  .project
-  > build.xml
-  debug.launch
-  known_hosts



Real-time interop. across entire cloud ecosystem



BlackBerry



ANDROID



Microsoft®
Silverlight™




 **Windows
phone**



Java™

Sun Microsystems



Microsoft®
.NET



Asynchronous Javascript And XML



mono™



 **iPhone**

MySQL

Memcache

Apache


PHP

HTML

SQL

CSS

Load Balancing




S3

HTML

CSS

ObjectFabric

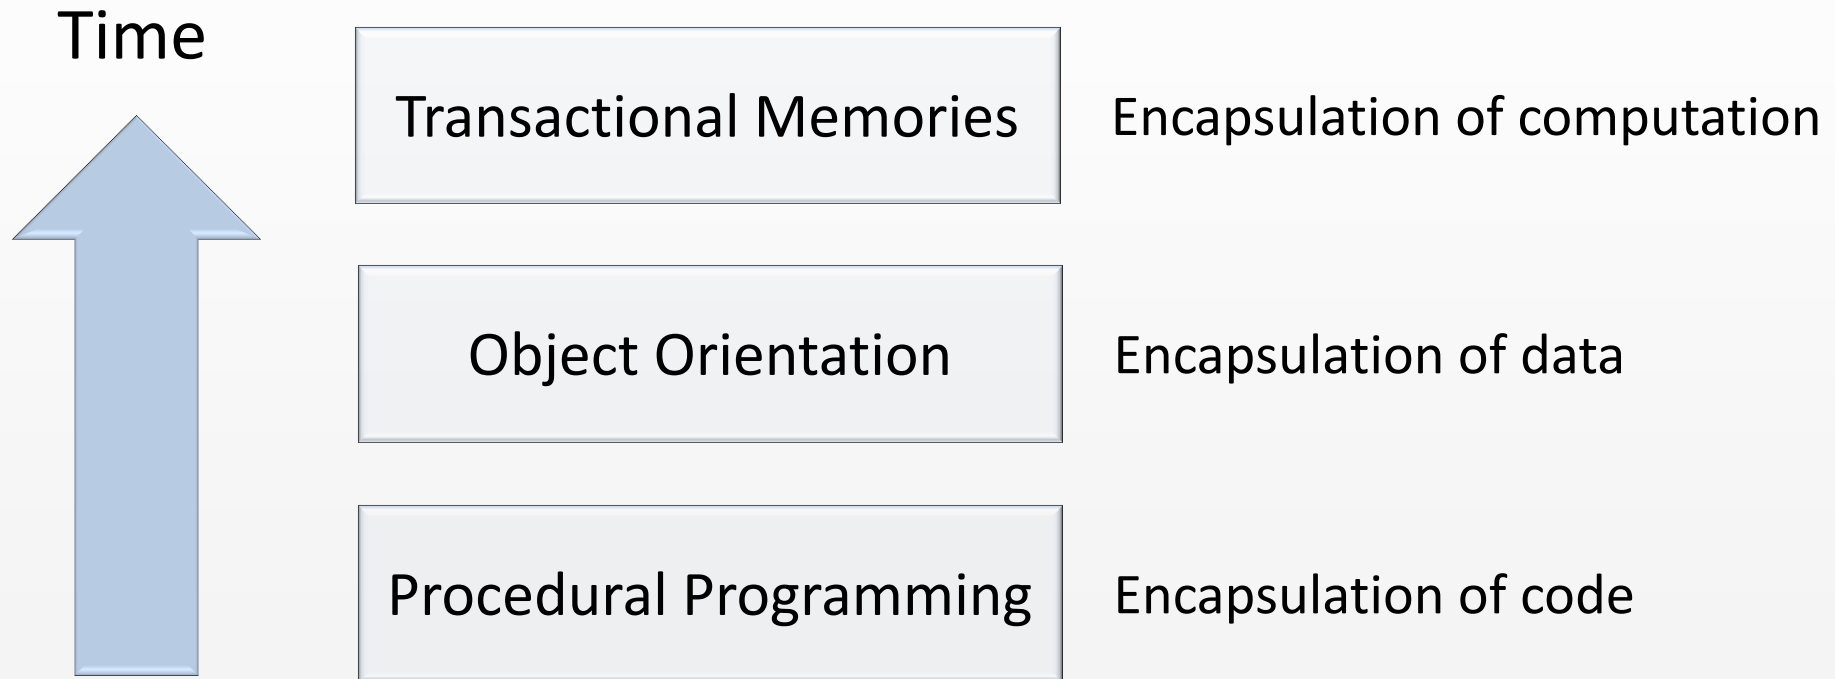
Java/.NET



Divagations!

- Transactions are the natural next step
- There are two right ways to model the world
- Boundaries between objects and messages should not be imposed by technology

Transactions are the natural next step



Two right ways to model the world

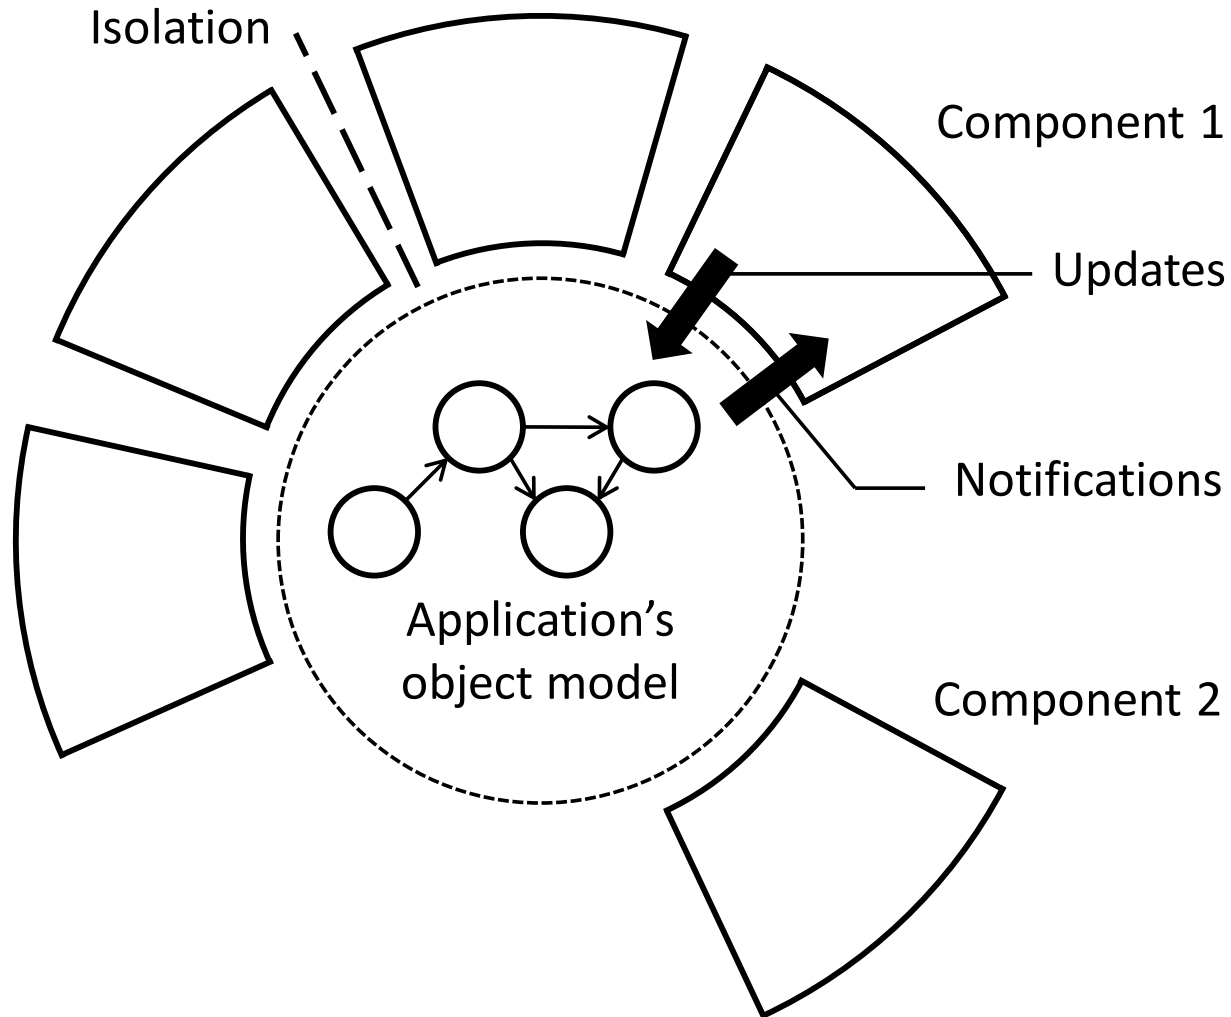
- Shared state mutated by transactions
 - Databases
 - Mainframes
 - Source controls
 - Transactional memories
 - Functional Programming
- Immutable messages between stateless processors
 - Actor Model
 - Message Bus, Queues etc.
 - Functional Programming

Two right ways to model the world

- "Objects" and "Messages"
 - Like frequency and time domains for sound
- The choice should be by domain or preference
 - Instead objects are confined to one process
 - Messages to "stitch" objects-based processes
 - In many applications domain modeled twice
- Previous example of UI connected to a server
 - Conversions Objects \leftrightarrow Messages \leftrightarrow Objects is wasteful and doesn't reduce coupling


Objects as a Bus

- Strengths of both shared state (OO, perf.) and Actors (safety)



Status / Road Map

- Java & GWT OK, .NET & Javascript in Beta
- Storage extensions
 - File based BTree (Ready)
 - SQL connector (Beta)
 - NoSQL connector (Alpha)
 - Offline Sync – HTML5/Mobile (Alpha)
- Scalability
 - Multi-Master, Eventual consistency
 - SimpleDB, ZooKeeper or Doozer
 - Goal: Hosted version (OaaS)
- Other
 - Specialized SDKs (Gaming, Chats)
 - REST connector (Alpha)



Thank You!

<https://github.com/ObjectFabric>