



: Reloaded

The Problem:

Programming is **HARD**....

- correct highly concurrent systems
- truly scalable systems
- fault tolerant systems that self-heal

... Let's go shopping!

Simpler

- Concurrency
- Scalability
- Fault-Tolerance

Vision

A single unified....

- programming model
- runtime service

Manage System Overload



Scale up & Scale out



Replicate and Distribute for fault tolerance





Transparent load balancing

What was...

Remote Actors

Remote Server

```
// use host & port in config  
Actor.remote.start()  
Actor.remote.start("localhost", 2552)
```

Scalable implementation based on NIO (Netty) & Protobuf

Remote Actor

```
import Actor._  
remote register ("service:id", actorOf[MyService])
```

server-side

Remote Actor

```
val service = remote actorFor (  
    "service:id",  
    "darkstar",  
    9999)
```

```
service ! message
```

client-side

We can do better!

Does **not** meet the vision

- Deployment (local vs remote) is a development decision
- We get a fixed and hard-coded topology
- Can't change it dynamically and adaptively

Needs to be
a **deployment** & **runtime** decision

Introducing ...

Clustered **Actors**

Address

```
val actor = actorOf[MyActor]
```

Bind the actor to a **virtual** address

Address

```
val actor = actorOf[MyActor]("my-service")
```

Bind the actor to a **virtual** address

Deployment Configuration

```
akka {  
  actor {  
    deployment {  
      my-service {  
        router = "least-cpu"  
        clustered {  
          replicas = 3  
          stateless = on  
        }  
      }  
    }  
  }  
}
```

Deployment Configuration

```
akka {  
  actor {  
    deployment {  
      my-service {  
        router = "least-cpu"  
        clustered {  
          replicas = 3  
          stateless = on  
        }  
      }  
    }  
  }  
}
```



Address

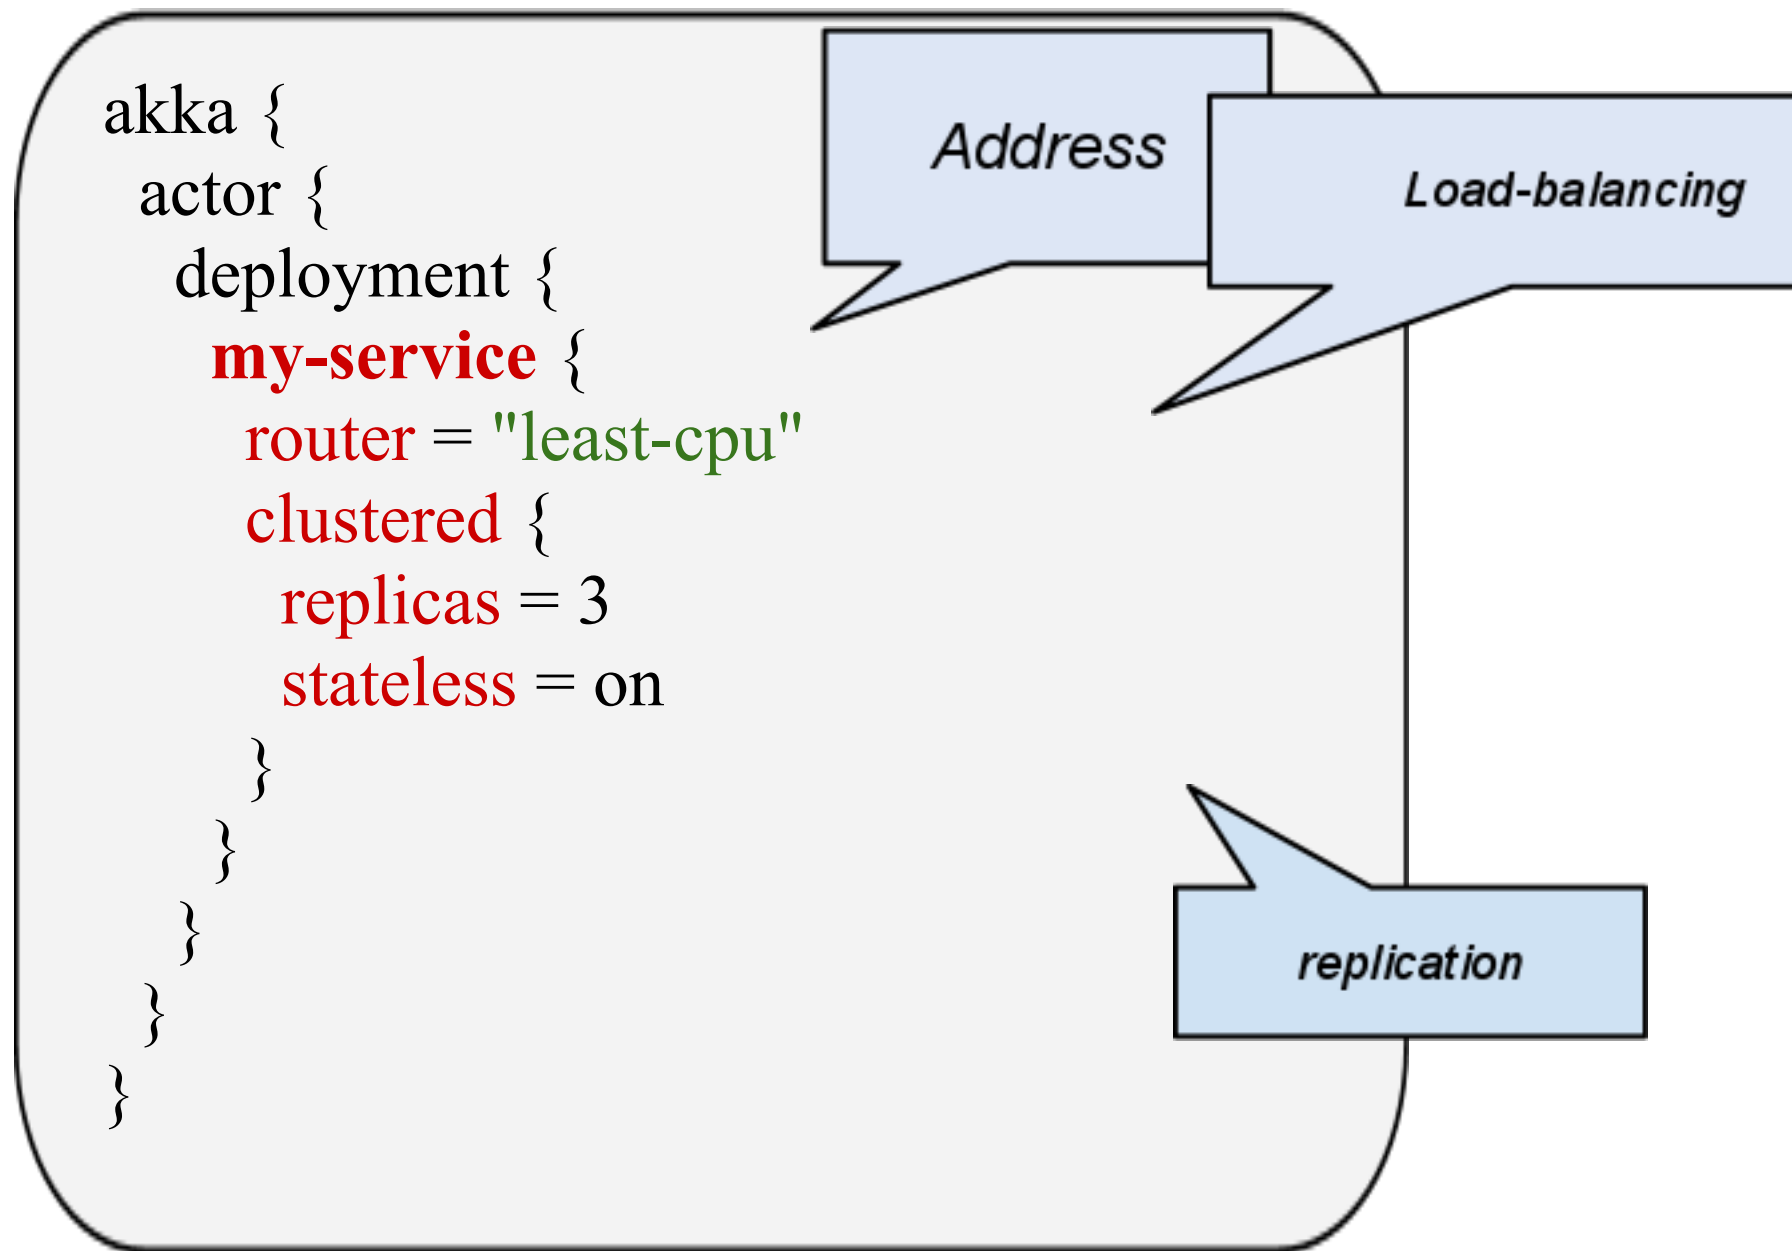
Deployment Configuration

```
akka {  
  actor {  
    deployment {  
      my-service {  
        router = "least-cpu"  
        clustered {  
          replicas = 3  
          stateless = on  
        }  
      }  
    }  
  }  
}
```

Address

Load-balancing

Deployment Configuration



Deployment

- Actor **address** is *decoupled* from **location** and **deployment**.

Deployment

- Actor **address** is *decoupled* from **location** and **deployment**.
- If **no configuration** is found, actor is **deployed locally**

Deployment

- Actor **address** is *decoupled* from **location** and **deployment**.
- If **no configuration** is found, actor is **deployed locally**
- The **same system** can be **configured for distribution** *without code change*

Deployment

- Actor **address** is *decoupled* from **location** and **deployment**.
- If **no configuration** is found, actor is **deployed locally**
- The **same system** can be **configured for distribution** *without code change*.
- Write and test **locally**. Test and deploy in the cloud **with confidence**.

Deployment

- Actor **address** is *decoupled* from **location** and **deployment**.
- If **no configuration** is found, actor is **deployed locally**
- The **same system** can be **configured for distribution** *without code change*.
- Write and test **locally**. Test and deploy in the cloud **with confidence**.
- Modify distribution **at runtime**.

Deployment

- Actor **address** is *decoupled* from **location** and **deployment**.
- If **no configuration** is found, actor is **deployed locally**
- The **same system** can be **configured for distribution** *without code change*.
- Write and test **locally**. Test and deploy in the cloud **with confidence**.
- Modify distribution **at runtime**.
- Runtime will **dynamically and adaptively change topology**

The runtime provides...

- Subscription based **cluster membership** service

The runtime provides...

- Subscription based **cluster membership** service
- Highly available **cluster registry** for actors

The runtime provides...

- Subscription based **cluster membership** service
- Highly available **cluster registry** for actors
- Automatic **cluster-wide deployment**

The runtime provides...

- Subscription based **cluster membership** service
- Highly available **cluster registry** for actors
- Automatic **cluster-wide deployment**
- **Automatic replication** with fail-over

The runtime provides...

- Subscription based **cluster membership** service
- Highly available **cluster registry** for actors
- Automatic **cluster-wide deployment**
- **Automatic replication** with **fail-over**
- Transparent and user configurable **load balancing**

The runtime provides...

- Subscription based **cluster membership** service
- Highly available **cluster registry** for actors
- Automatic **cluster-wide deployment**
- **Automatic replication** with **fail-over**
- Transparent and user configurable **load balancing**
- Transparent **adaptive cluster rebalancing**

The runtime provides...

- Subscription based **cluster membership** service
- Highly available **cluster registry** for actors
- Automatic **cluster-wide deployment**
- **Automatic replication** with **fail-over**
- Transparent and user configurable **load balancing**
- Transparent **adaptive cluster rebalancing**
- **Leader election**

The runtime provides...

- Subscription based **cluster membership** service
- Highly available **cluster registry** for actors
- Automatic **cluster-wide deployment**
- **Automatic replication** with **fail-over**
- Transparent and user configurable **load balancing**
- Transparent **adaptive cluster rebalancing**
- **Leader election**
- Durable Mailboxes - **guaranteed delivery**

The runtime provides...

- Subscription based **cluster membership** service
- Highly available **cluster registry** for actors
- Automatic **cluster-wide deployment**
- **Automatic replication** with **fail-over**
- Transparent and user configurable **load balancing**
- Transparent **adaptive cluster rebalancing**
- **Leader election**
- Durable Mailboxes
- Highly available centralized **configuration service**

The runtime provides...


- Subscription based **cluster membership** service
- Highly available **cluster registry** for actors
- Automatic **cluster-wide deployment**
- **Automatic replication** with **fail-over**
- Transparent and user configurable **load balancing**
- Transparent **adaptive cluster rebalancing**
- **Leader election**
- Durable Mailboxes
- Highly available centralized **configuration service**
- ... and more

Clustering of Stateless Actors

Classic Example

```
val ping = actorOf[Ping]("ping")  
val pong = actorOf[Pong]("ping")  
  
ping ! Ball(pong)
```

Akka Node



Akka Cluster Node

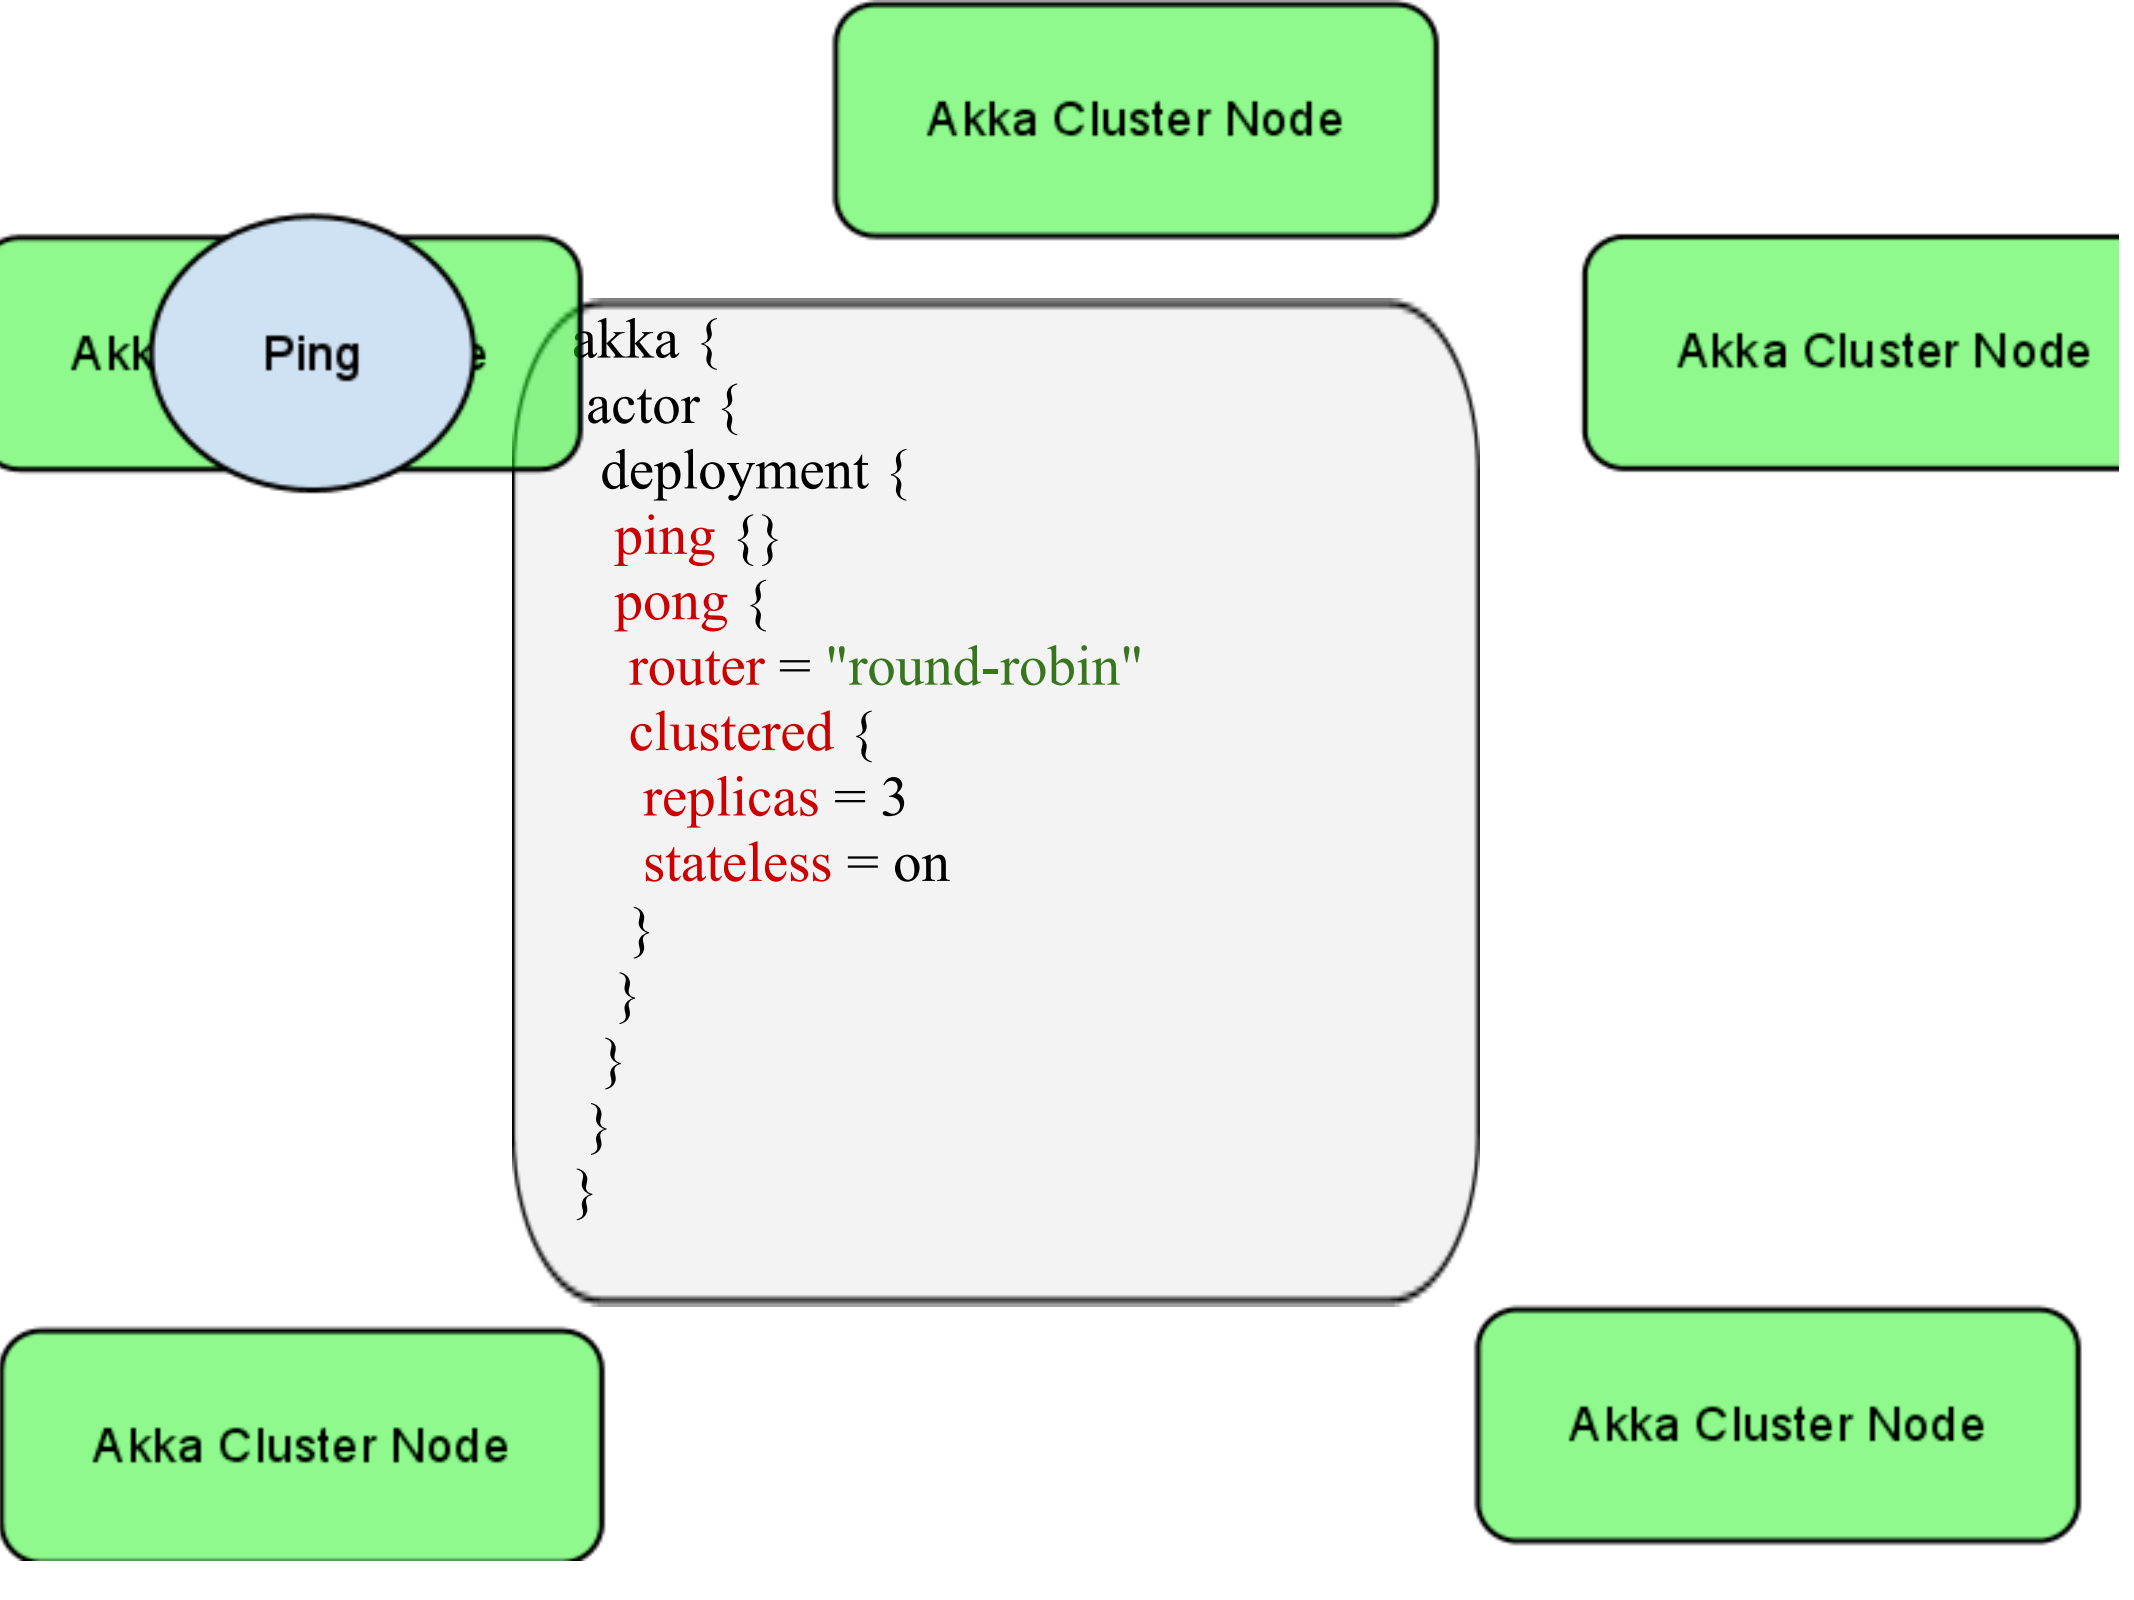
Akka Cluster Node

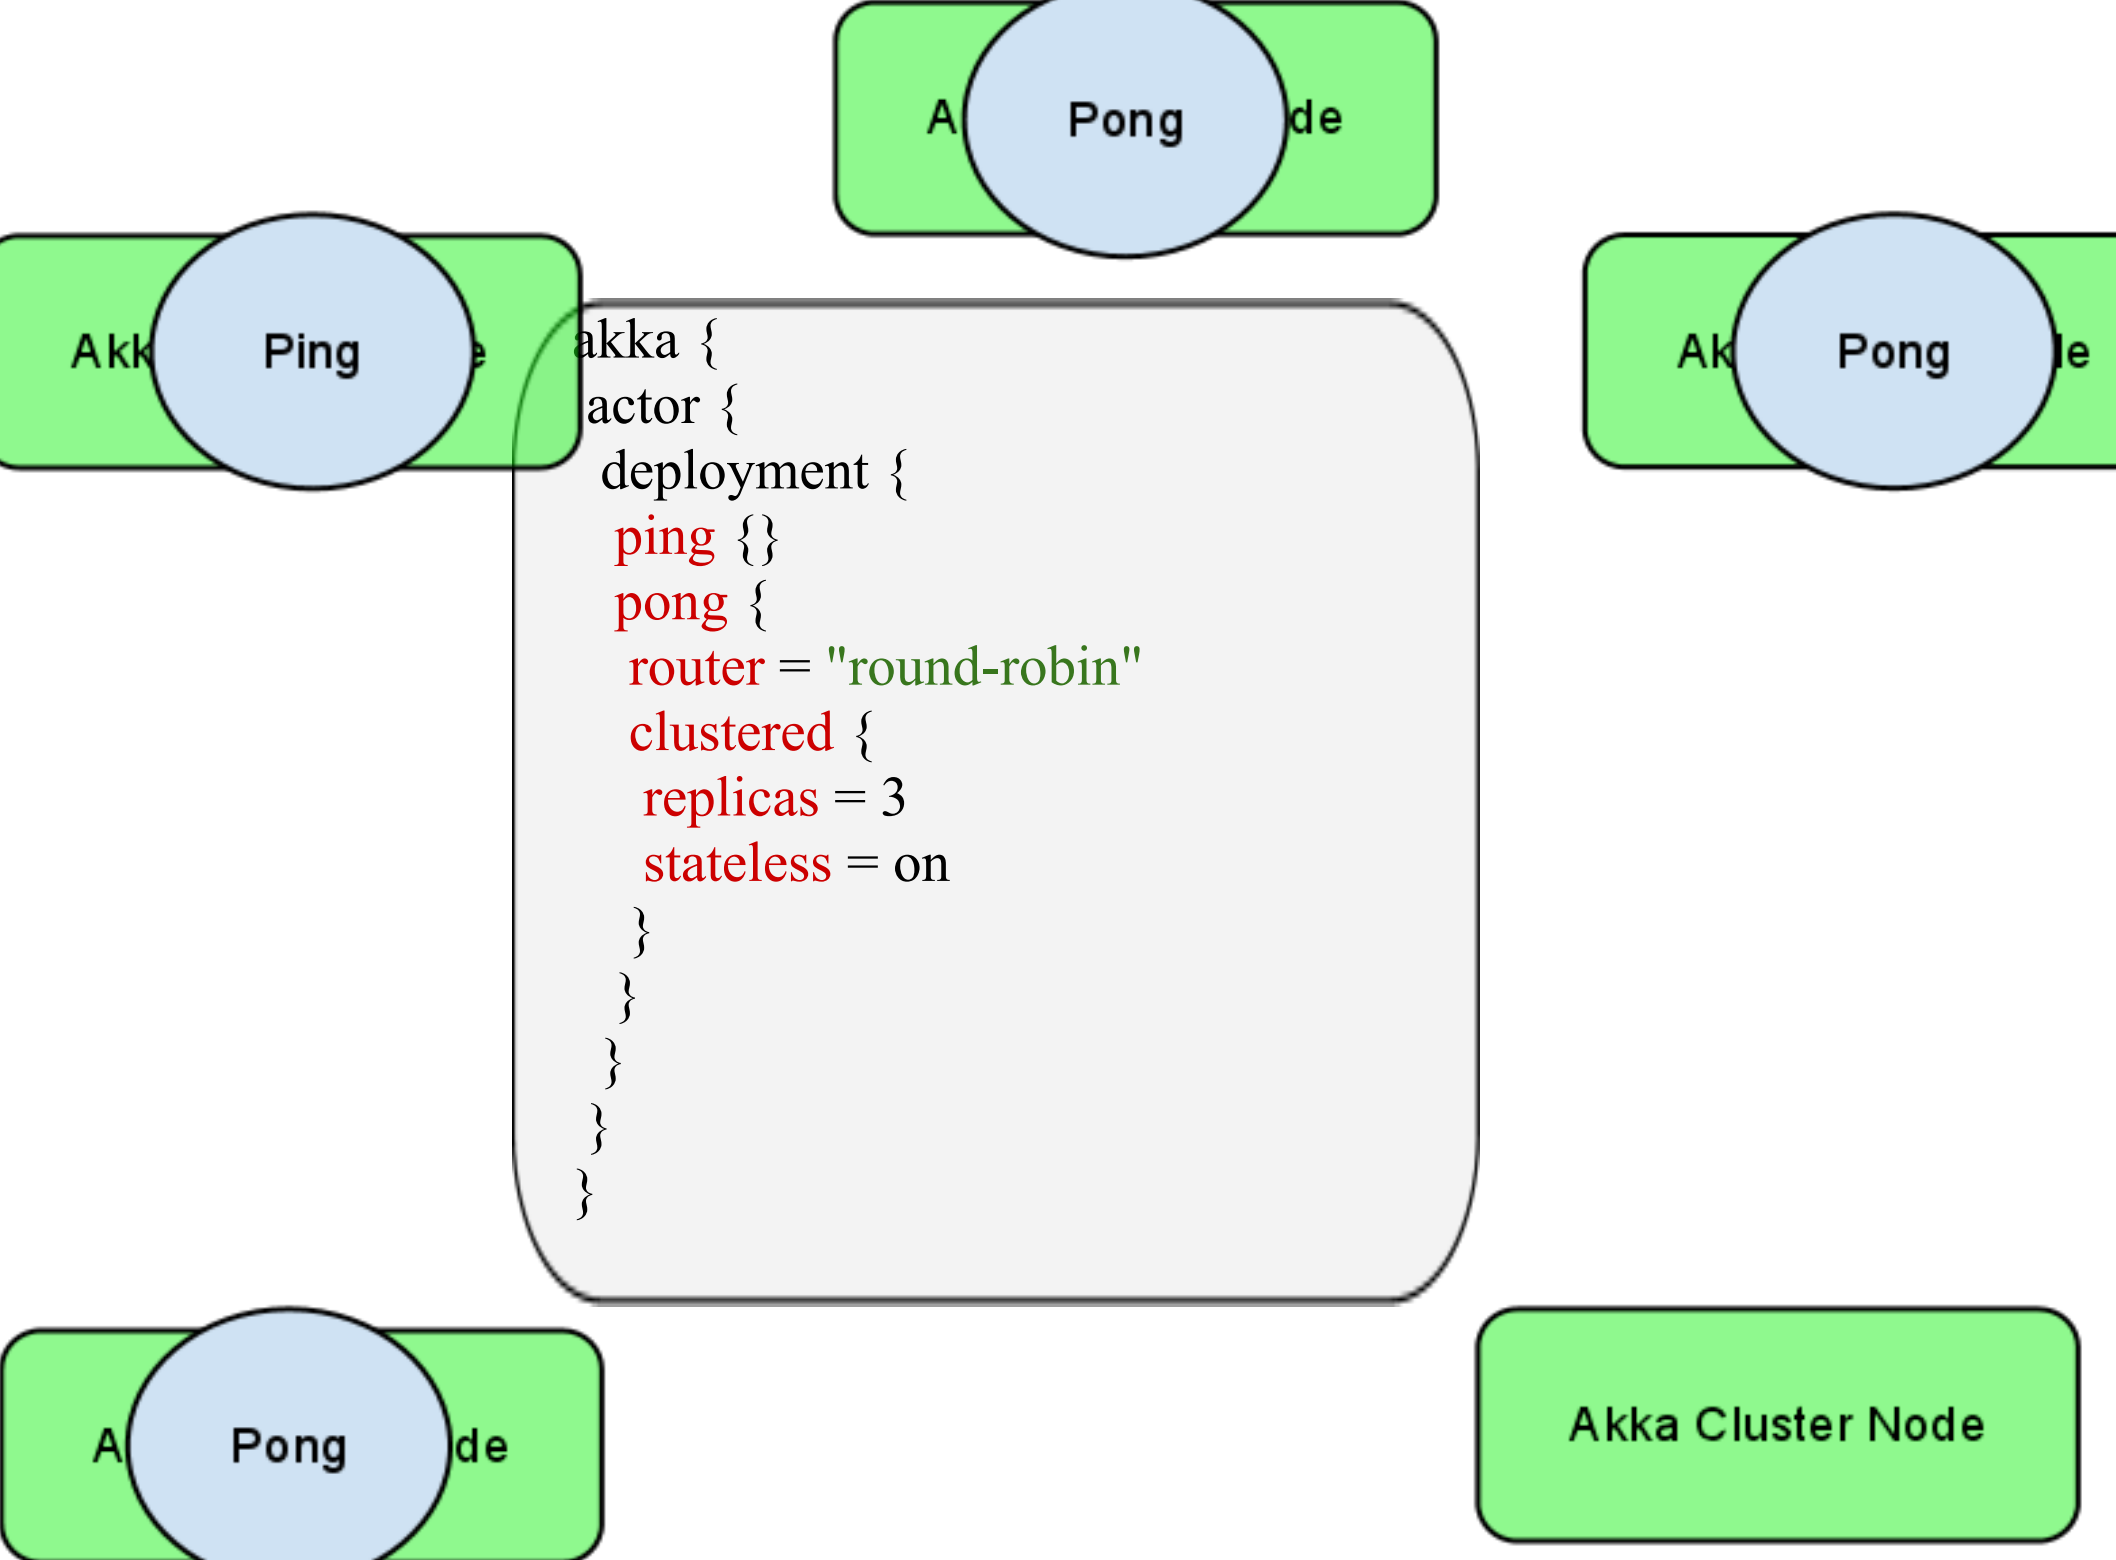
Akka Cluster Node

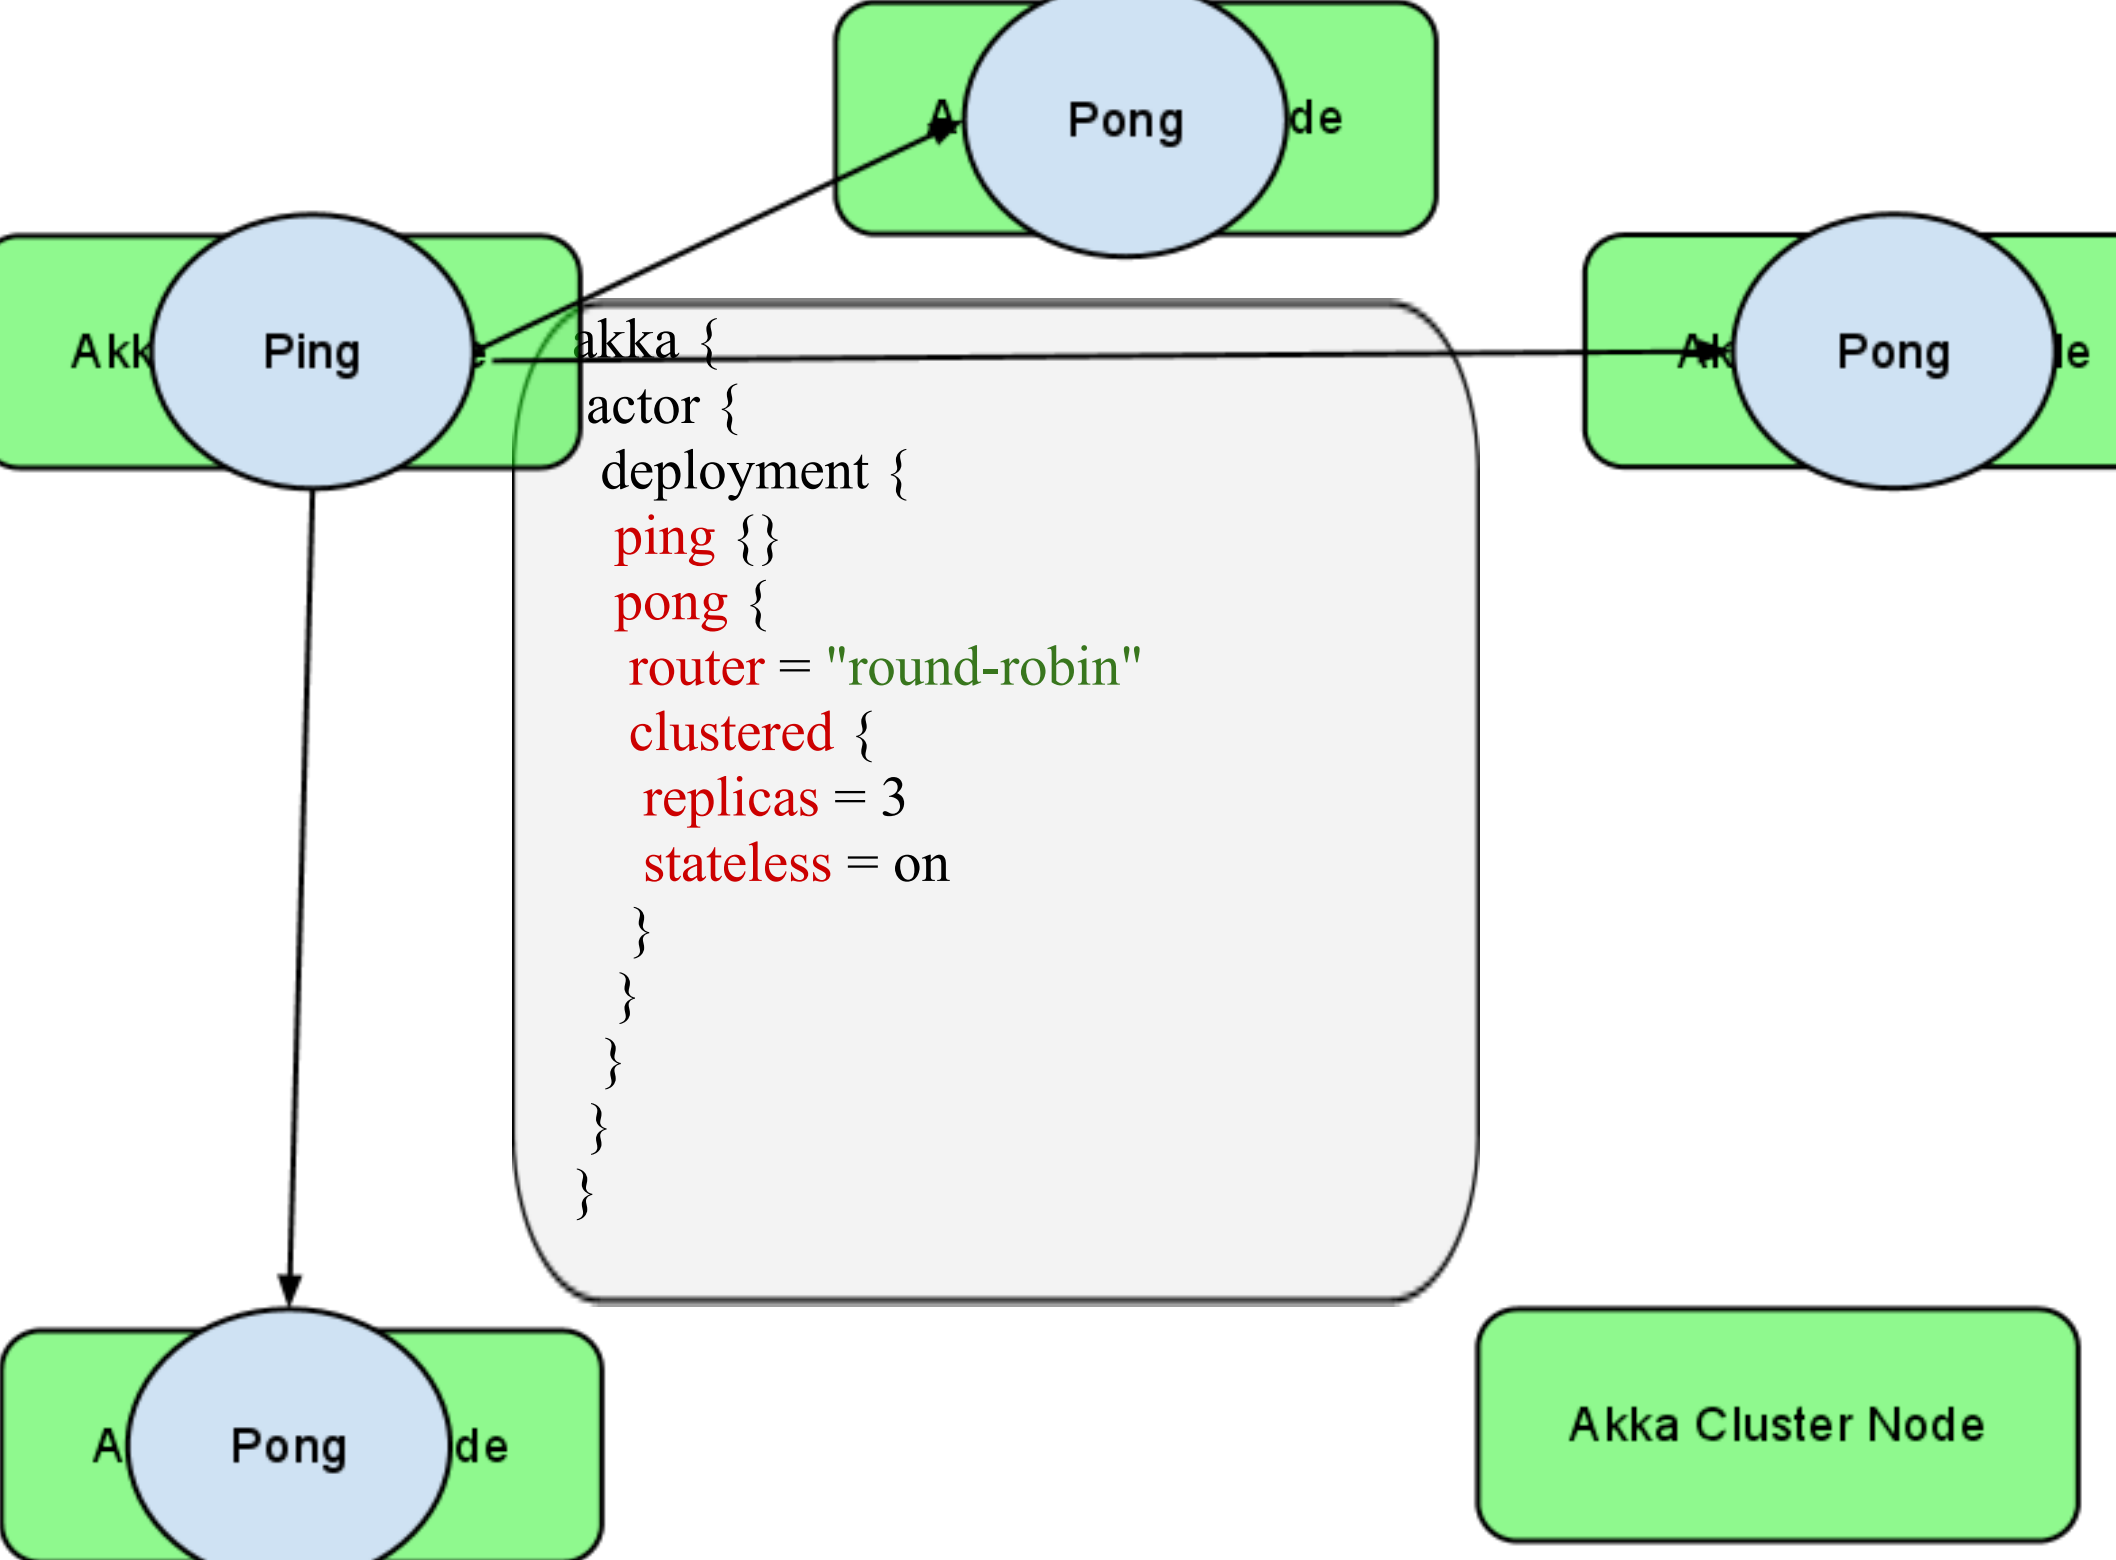
```
akka {  
  actor {  
    deployment {  
      ping {}  
      pong {  
        router = "round-robin"  
        clustered {  
          replicas = 3  
          stateless = on  
        }  
      }  
    }  
  }  
}
```

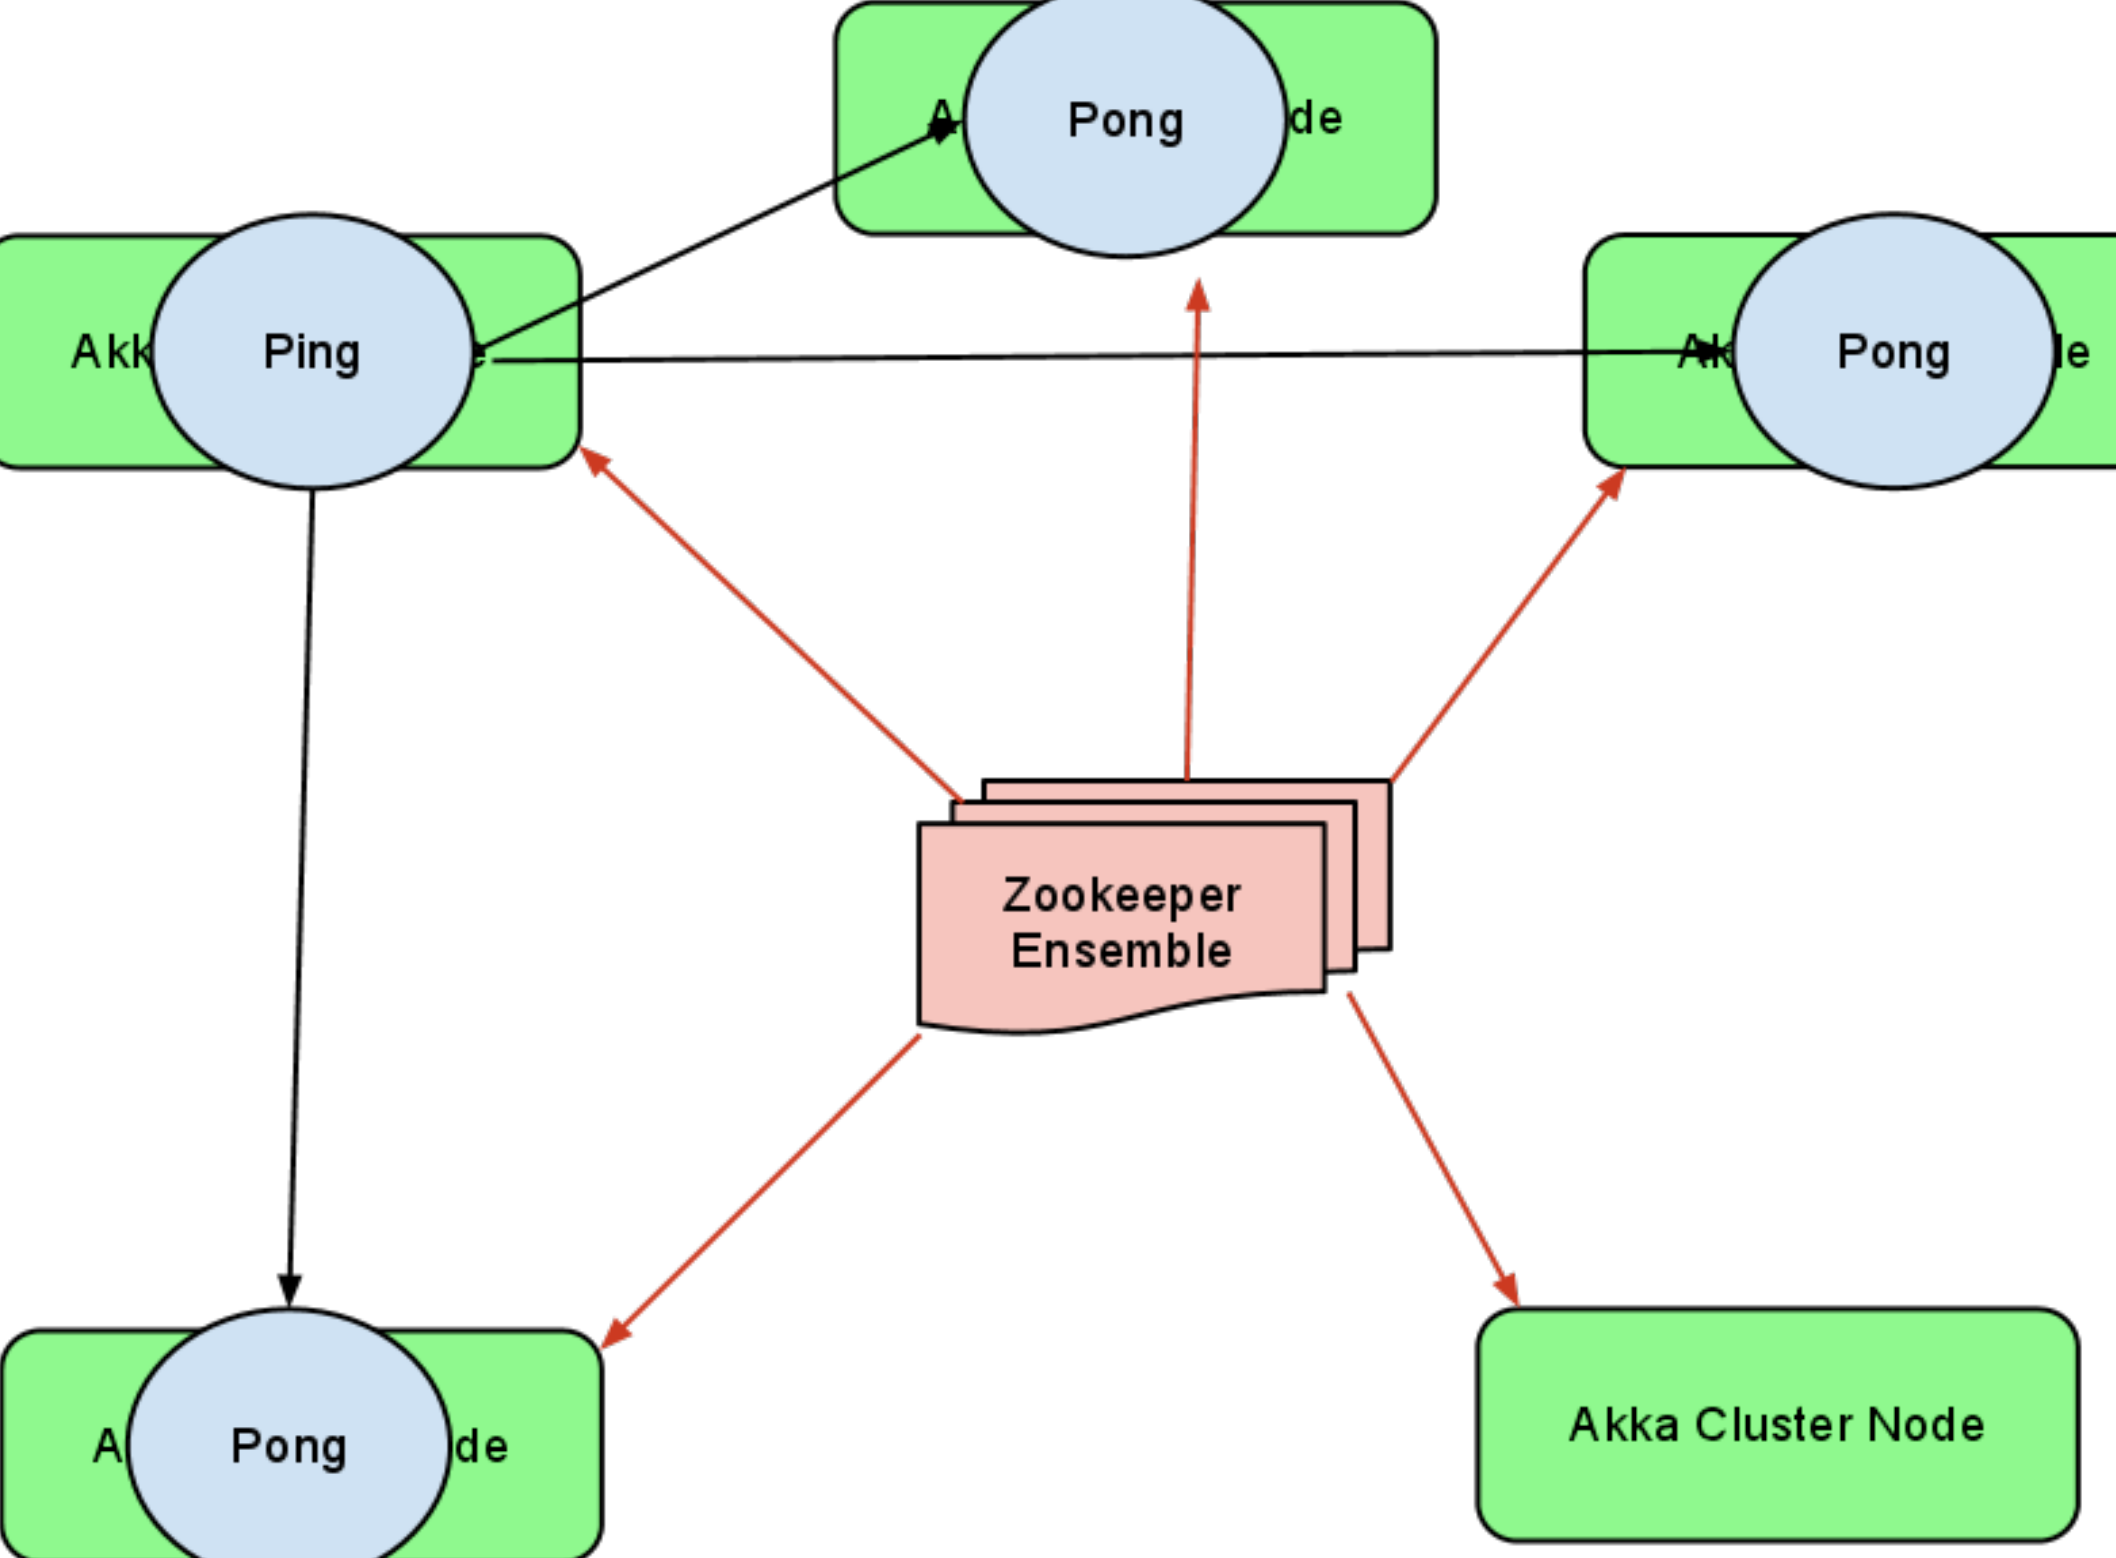
Akka Cluster Node

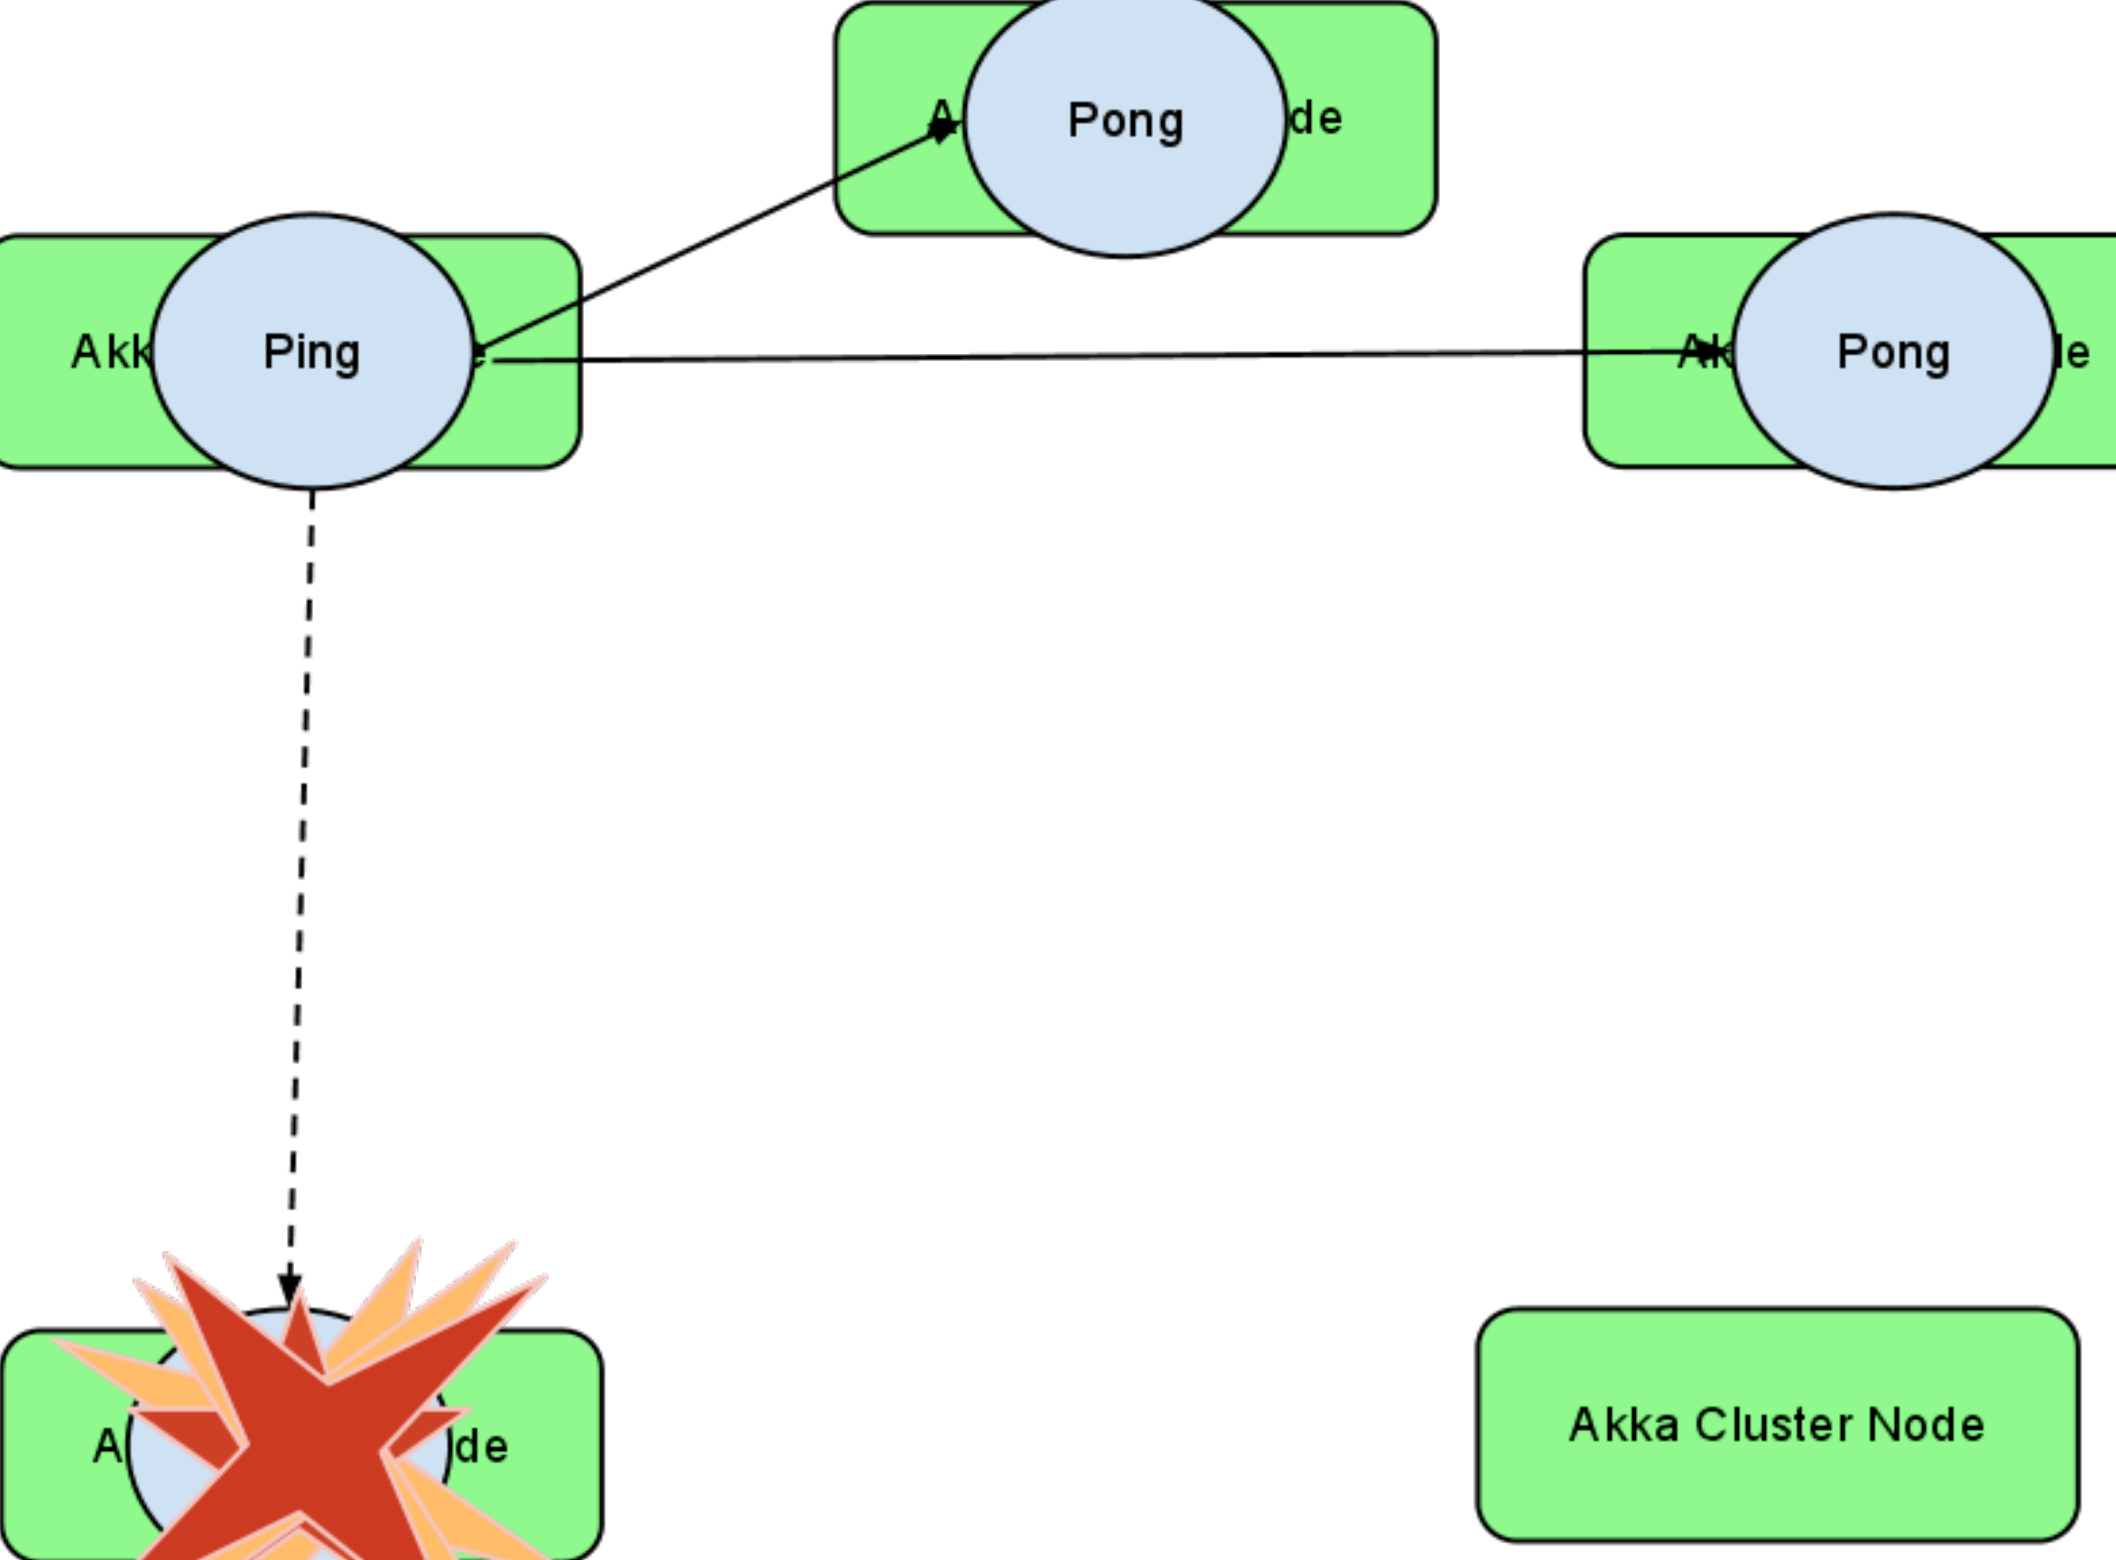
Akka Cluster Node

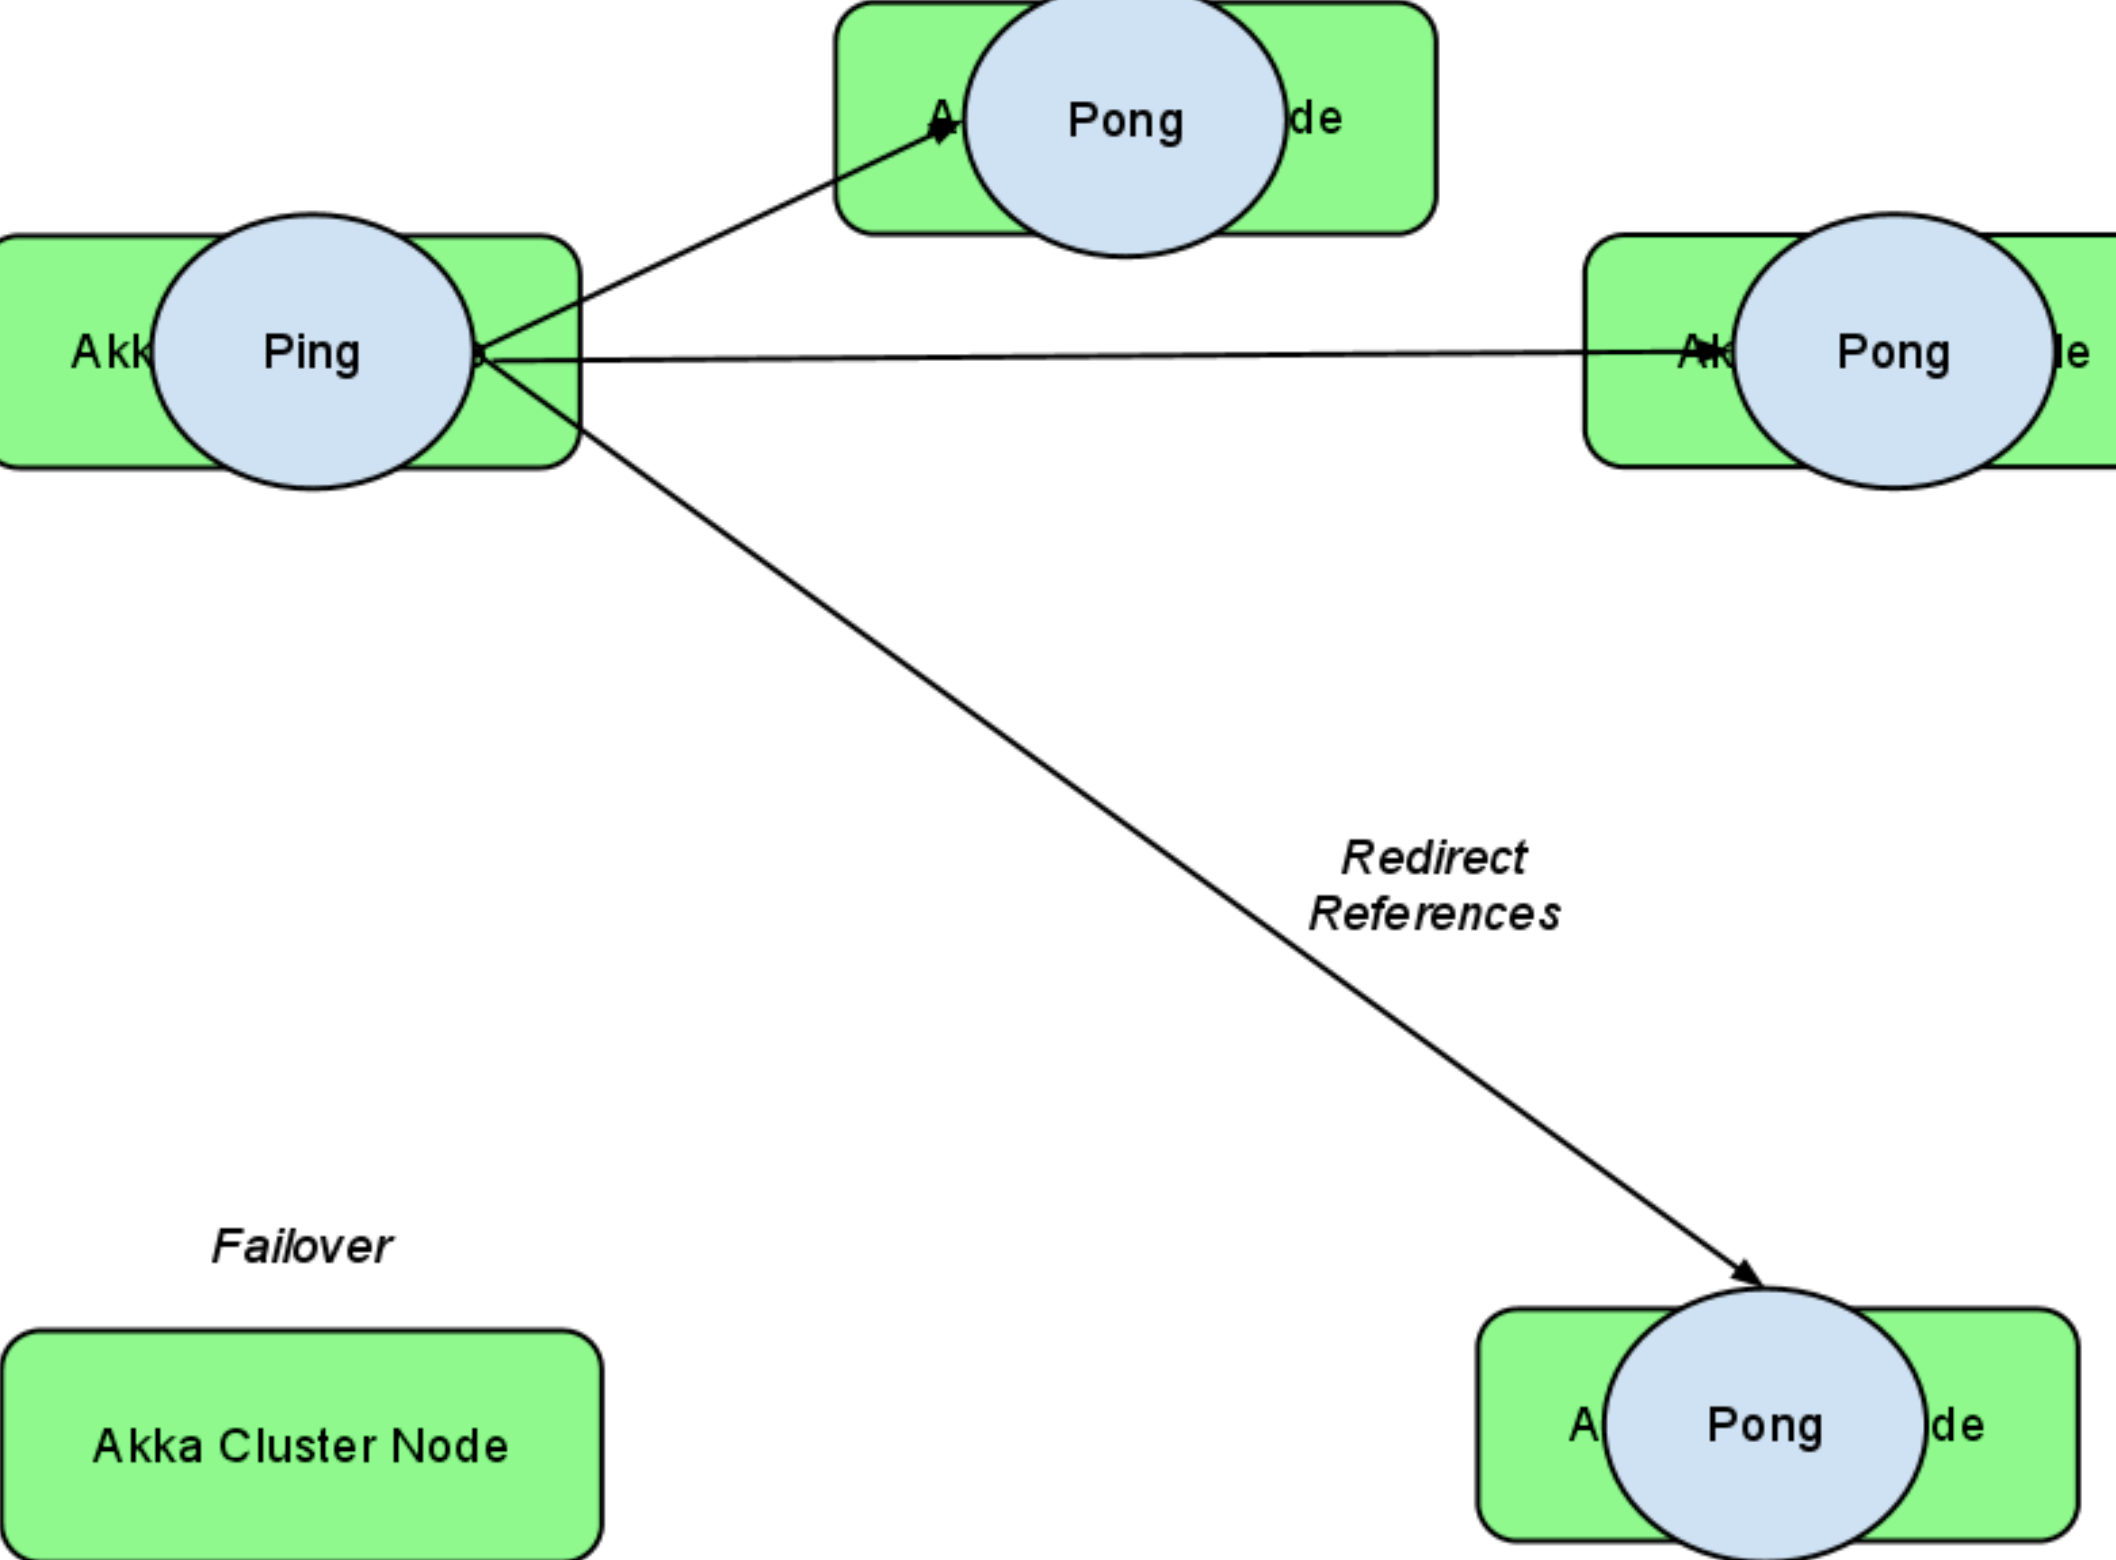












Clustering of Stateful Actors

Replication

1

Transaction Log

Deployment Configuration

```
akka {  
  actor {  
    deployment {  
      carts {  
        clustered {  
          home = "node:test-node-1"  
          stateless = off  
        }  
      }  
    }  
  }  
}
```

Deployment Configuration

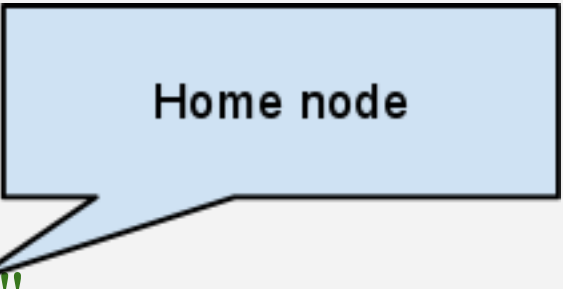
```
akka {  
  actor {  
    deployment {  
      carts {  
        clustered {  
          home = "node:test-node-1"  
          stateless = off  
        }  
      }  
    }  
  }  
}
```



Home node

Deployment Configuration

```
akka {  
  actor {  
    deployment {  
      carts {  
        clustered {  
          home = "node:test-node-1"  
          stateless = off  
        }  
      }  
    }  
  }  
}
```



Home node



Stateful

Akka node

Client

Cart



Akka Cluster Node




Akka Cluster Node



Akka Cluster Node

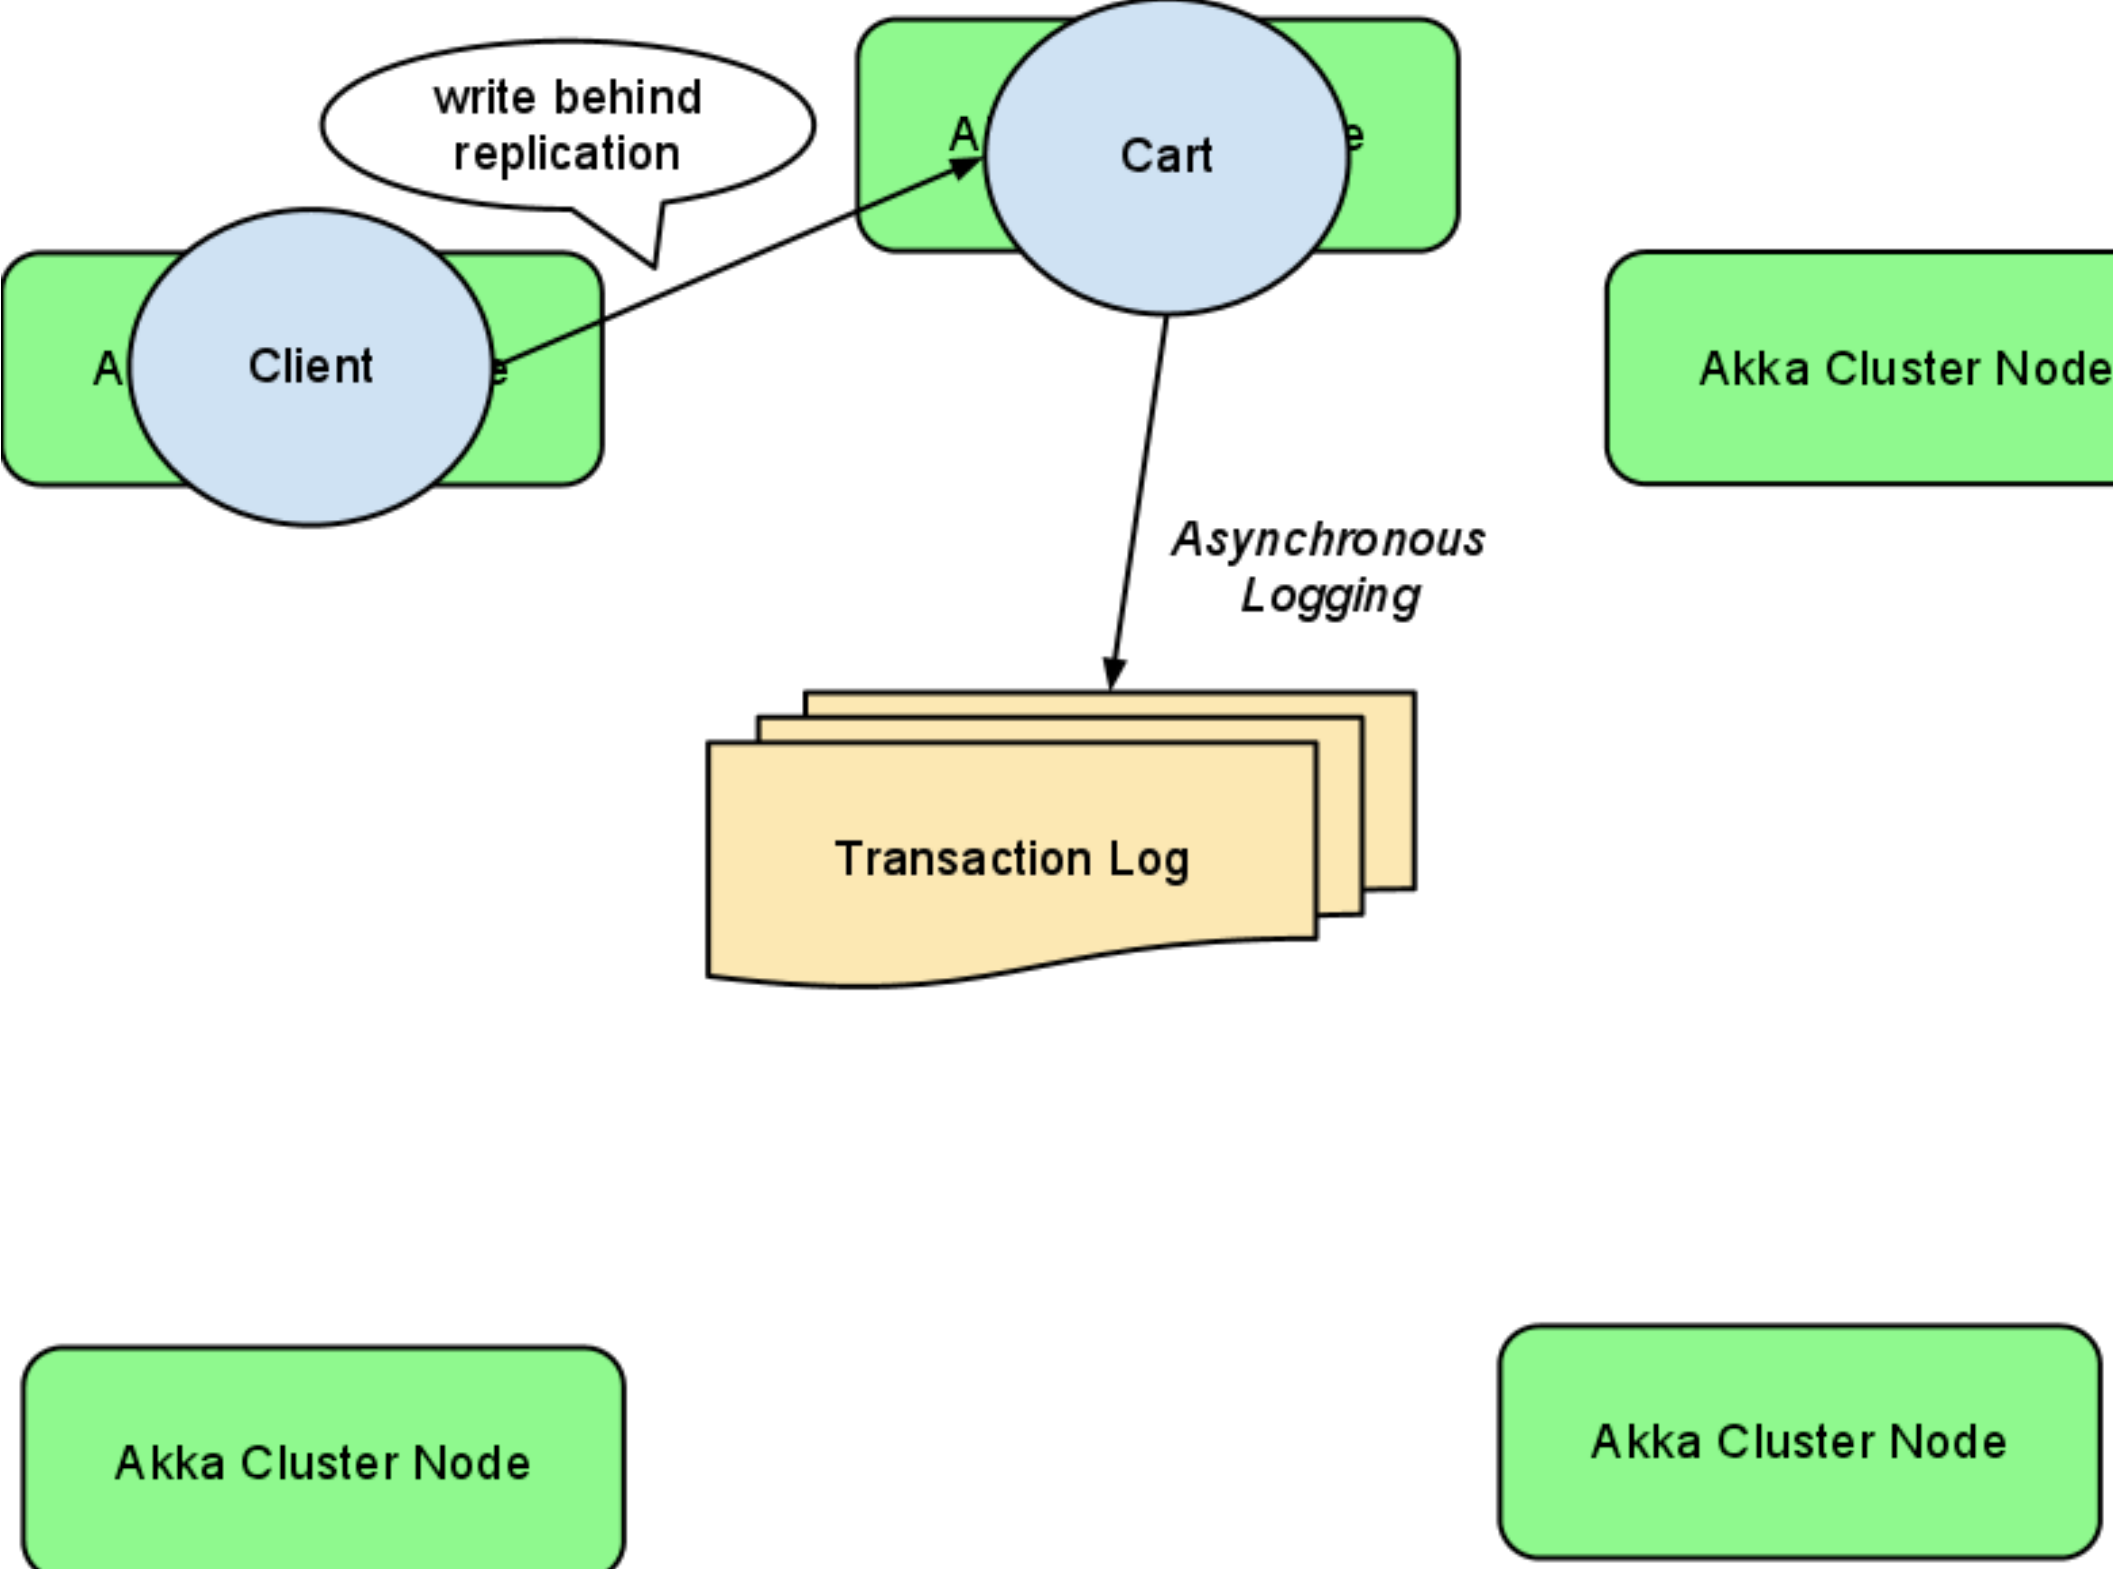
```
akka {  
  actor {  
    deployment {  
      carts {  
        clustered {  
          home = "node:test-node-1"  
          stateless = off  
        }  
      }  
    }  
  }  
}
```

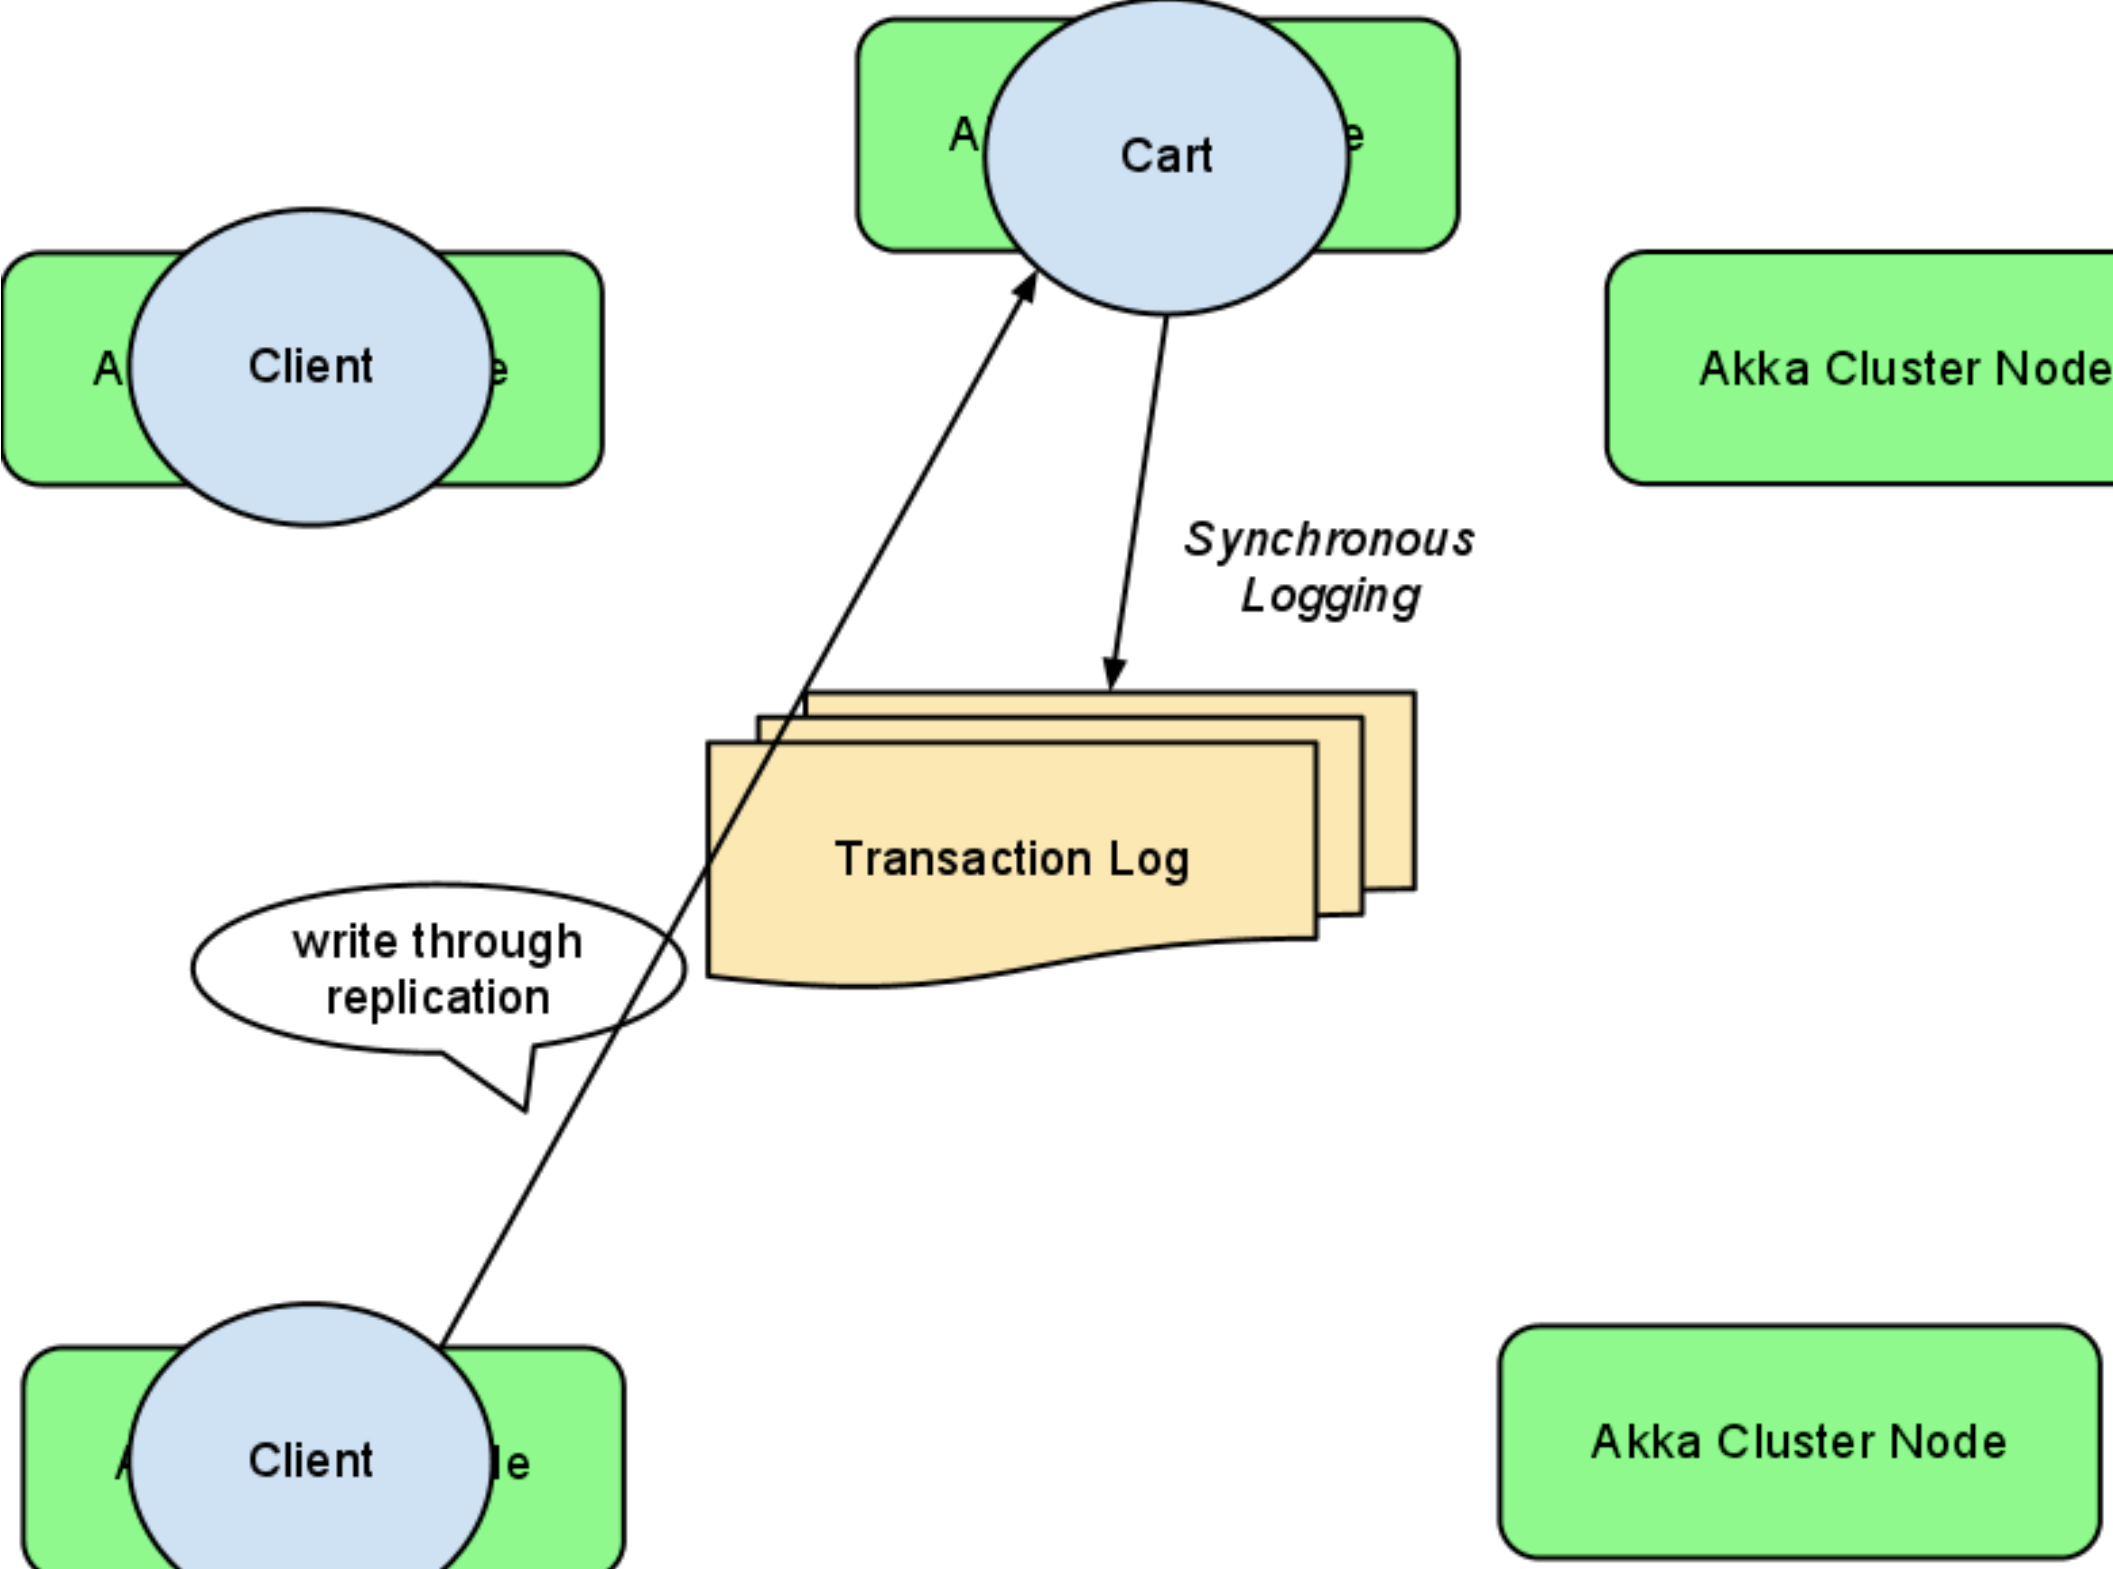
Akka Cluster Node



Akka Cluster Node







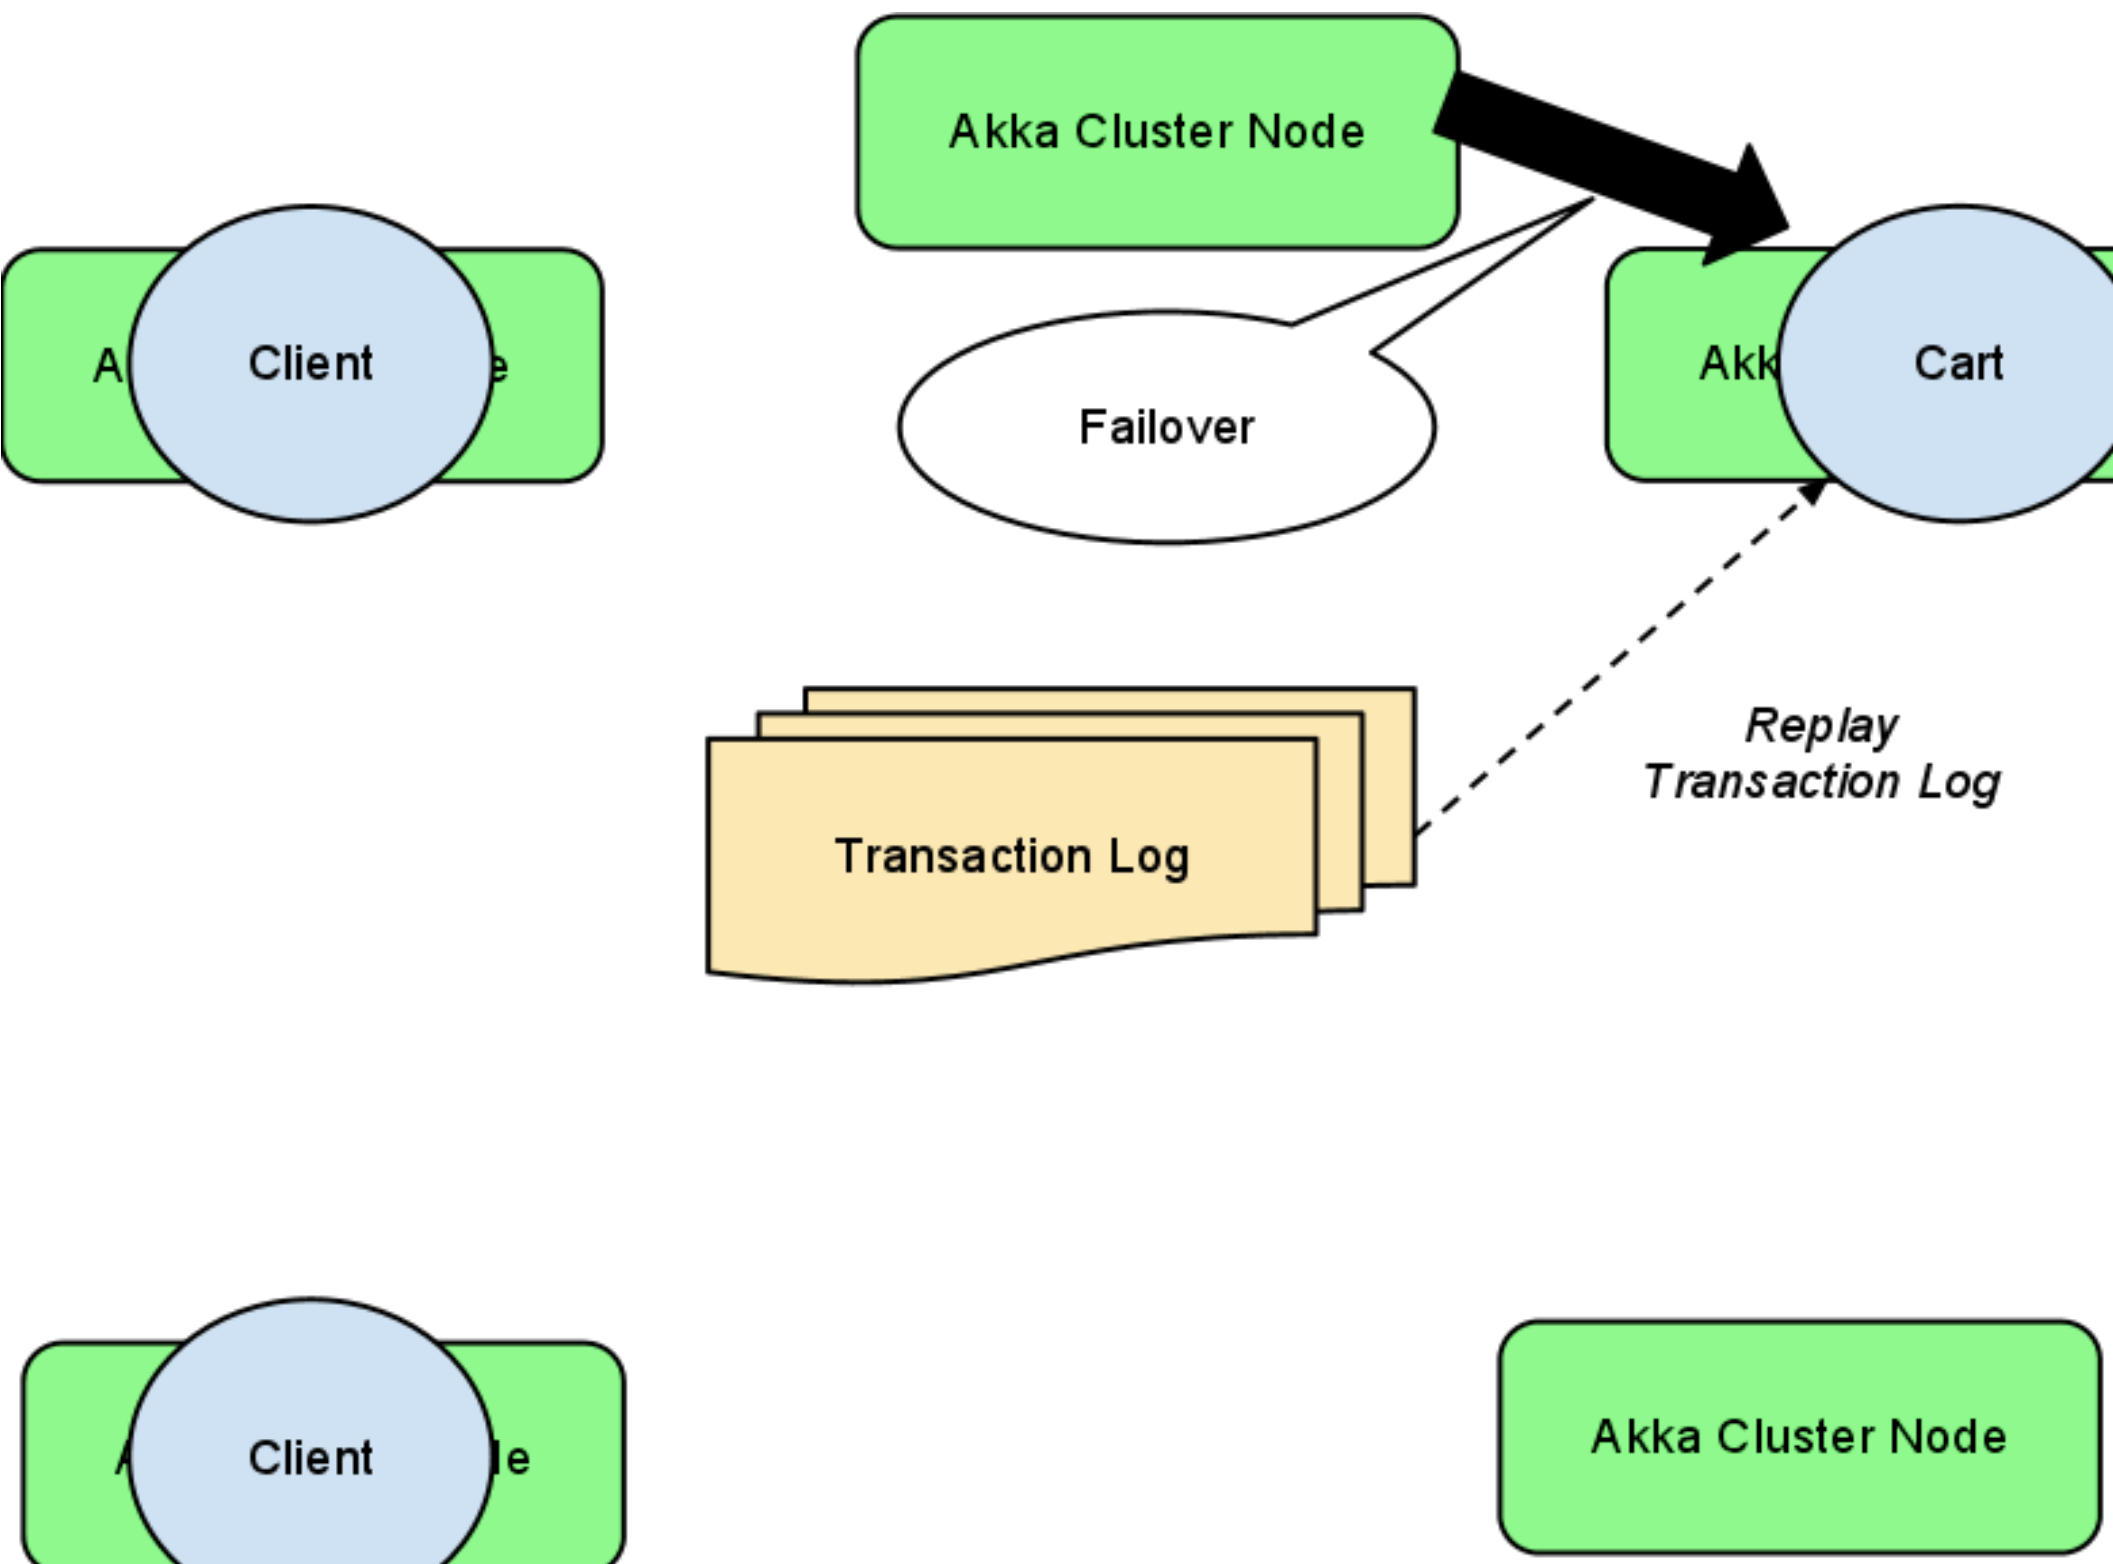








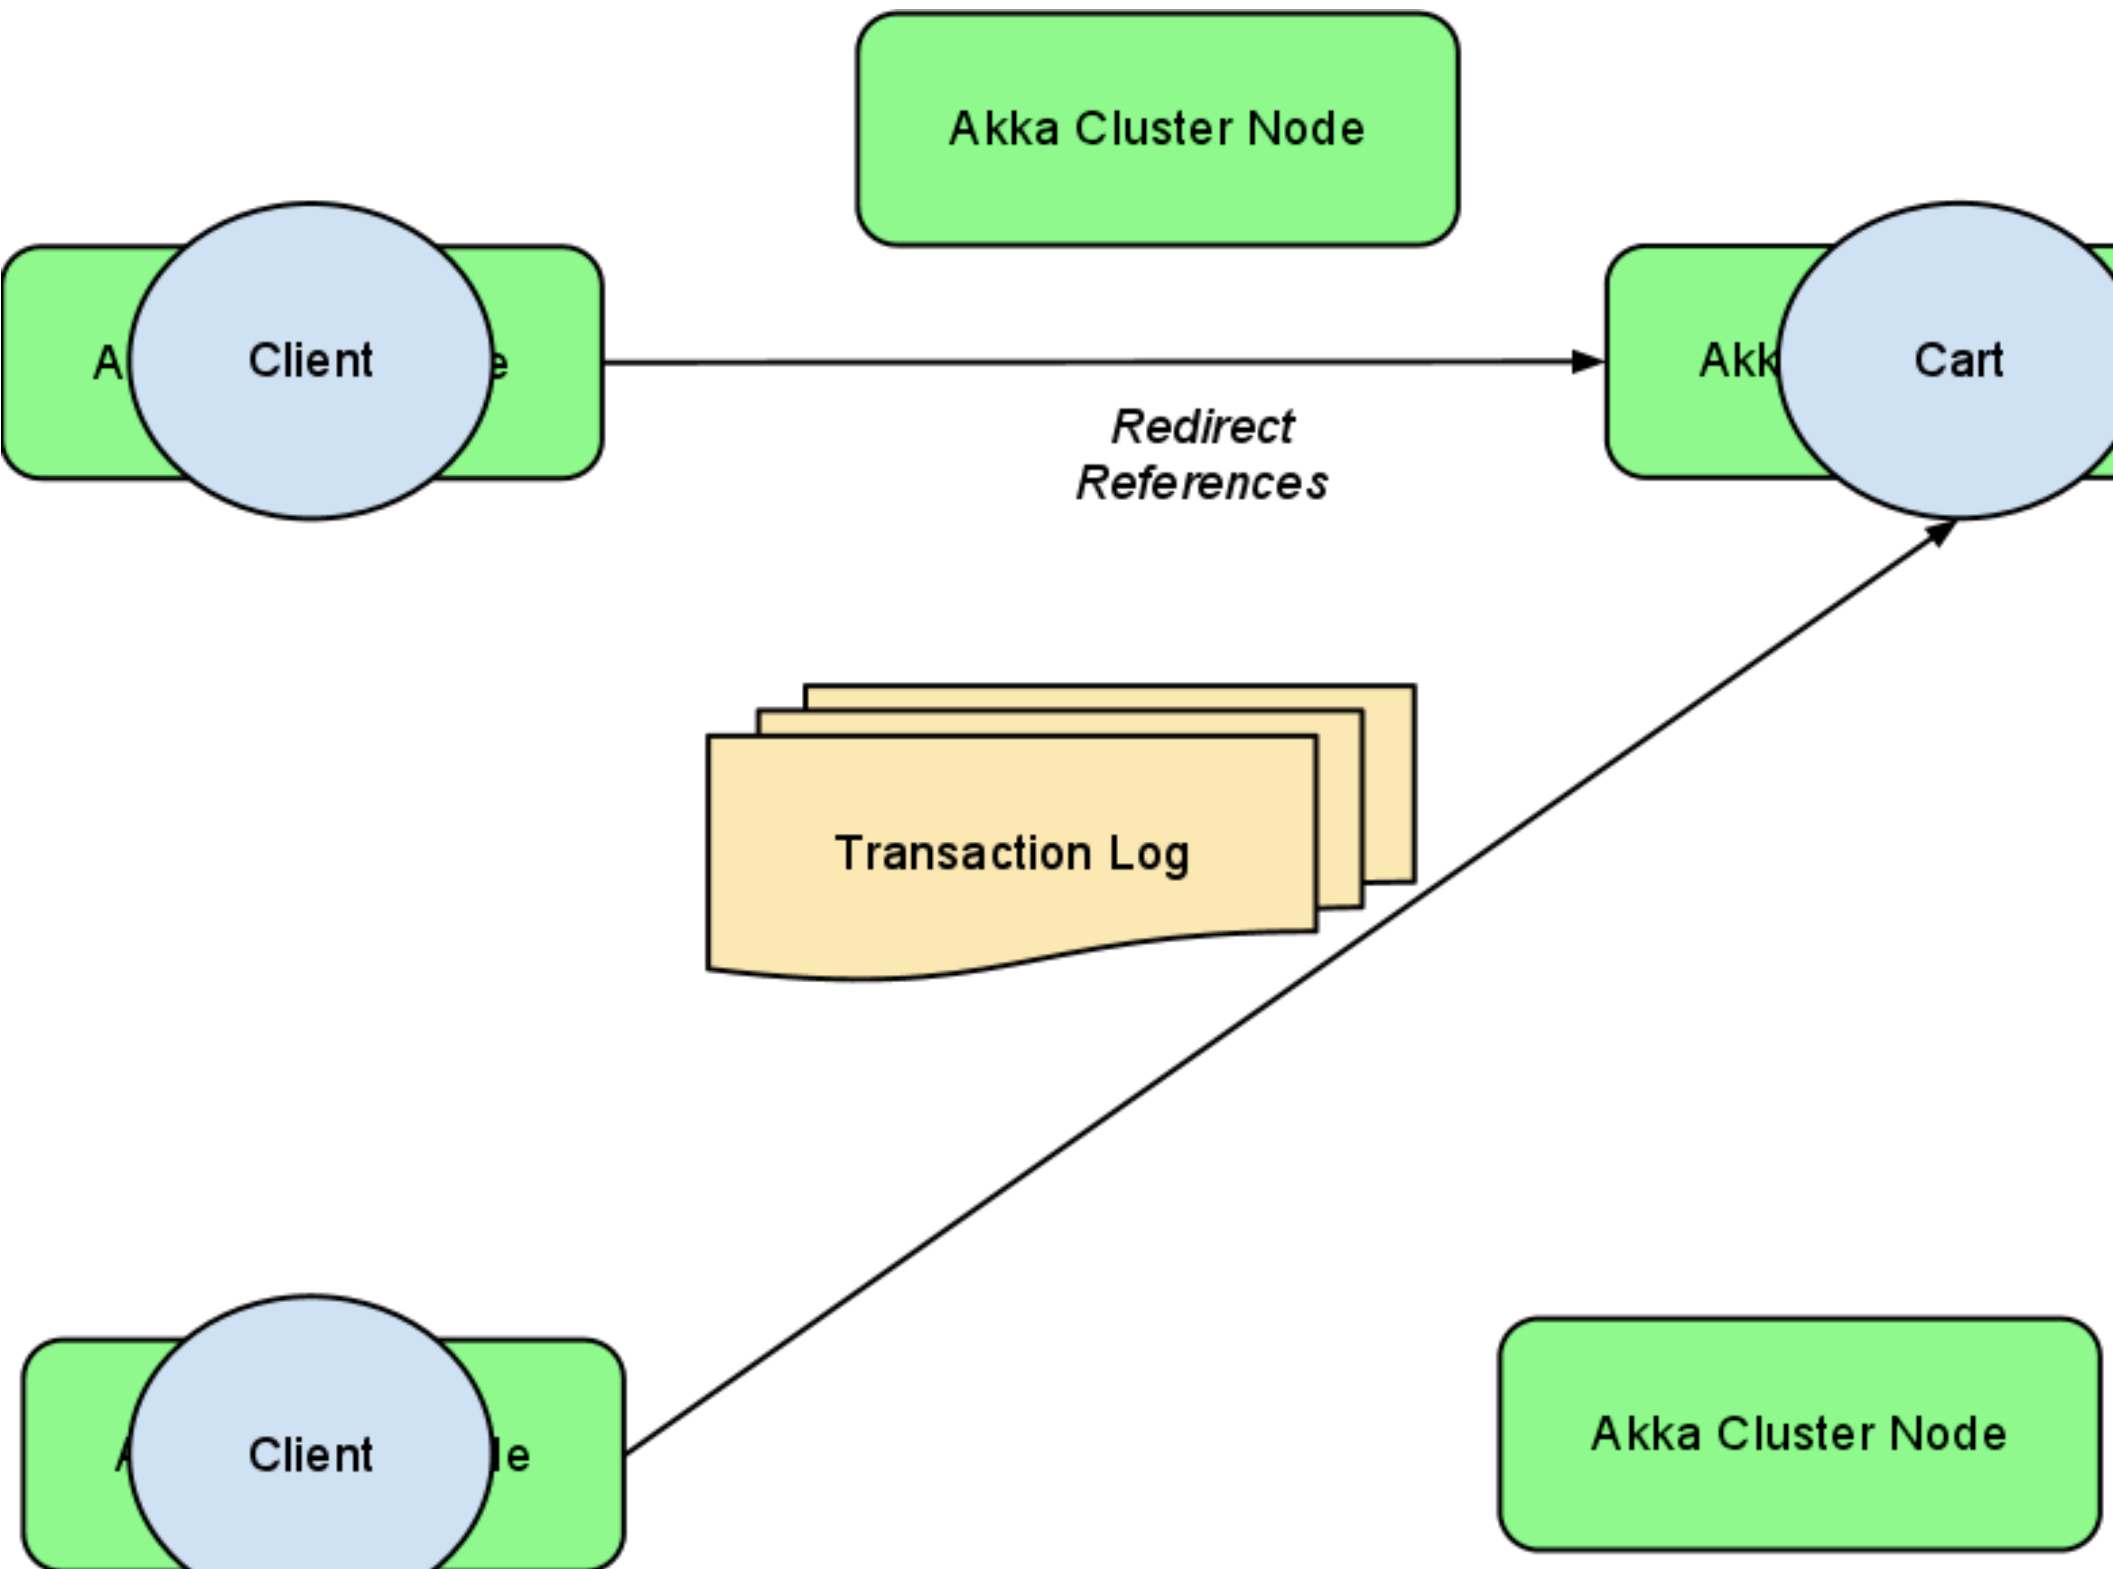














Transaction Log: Storing Messages

Transaction Log: Storing Messages




Transaction Log: Storing Messages




Transaction Log: Storing Messages




Transaction Log: Storing Messages




Transaction Log: Replaying



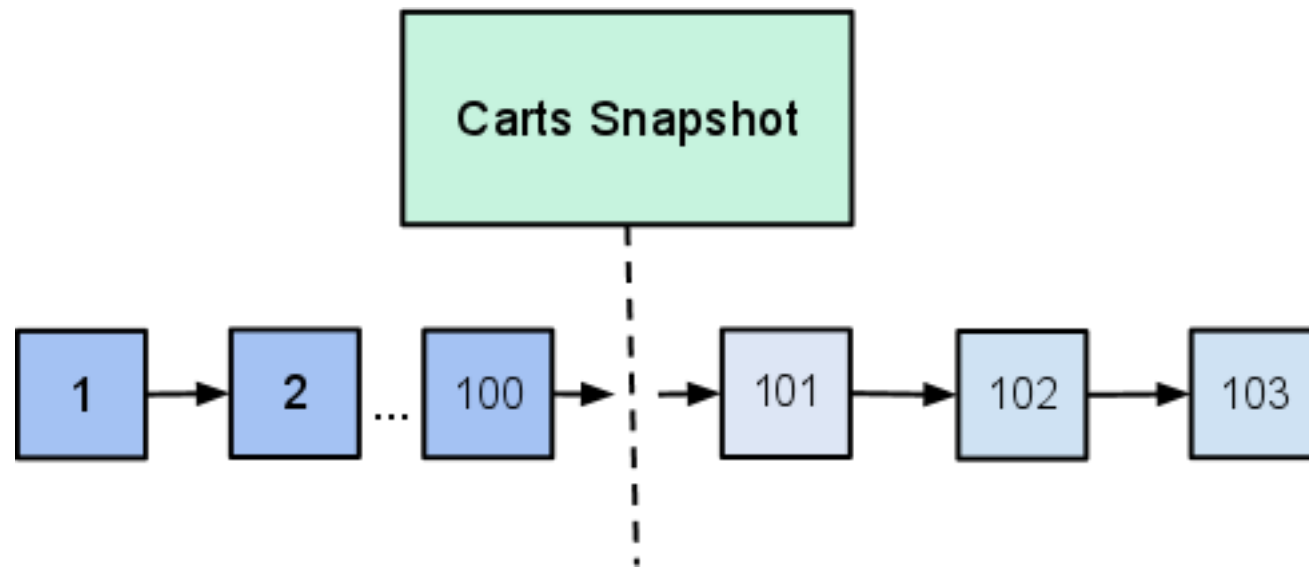
Transaction Log: Rolling Snapshot



Transaction Log: Rolling Snapshot



Transaction Log: Rolling Snapshot



Replication

2

Data Grid

Data Grid

Actor State

- Stored in 'external' Data Grid
- Transactional (distributed STM)
- Versioned
- Replicated
- Query-able

Implementations

- Custom Akka Data Grid
- SPI for third parties

Akka Cluster Node




Akka Cluster Node





Akka Cluster Node

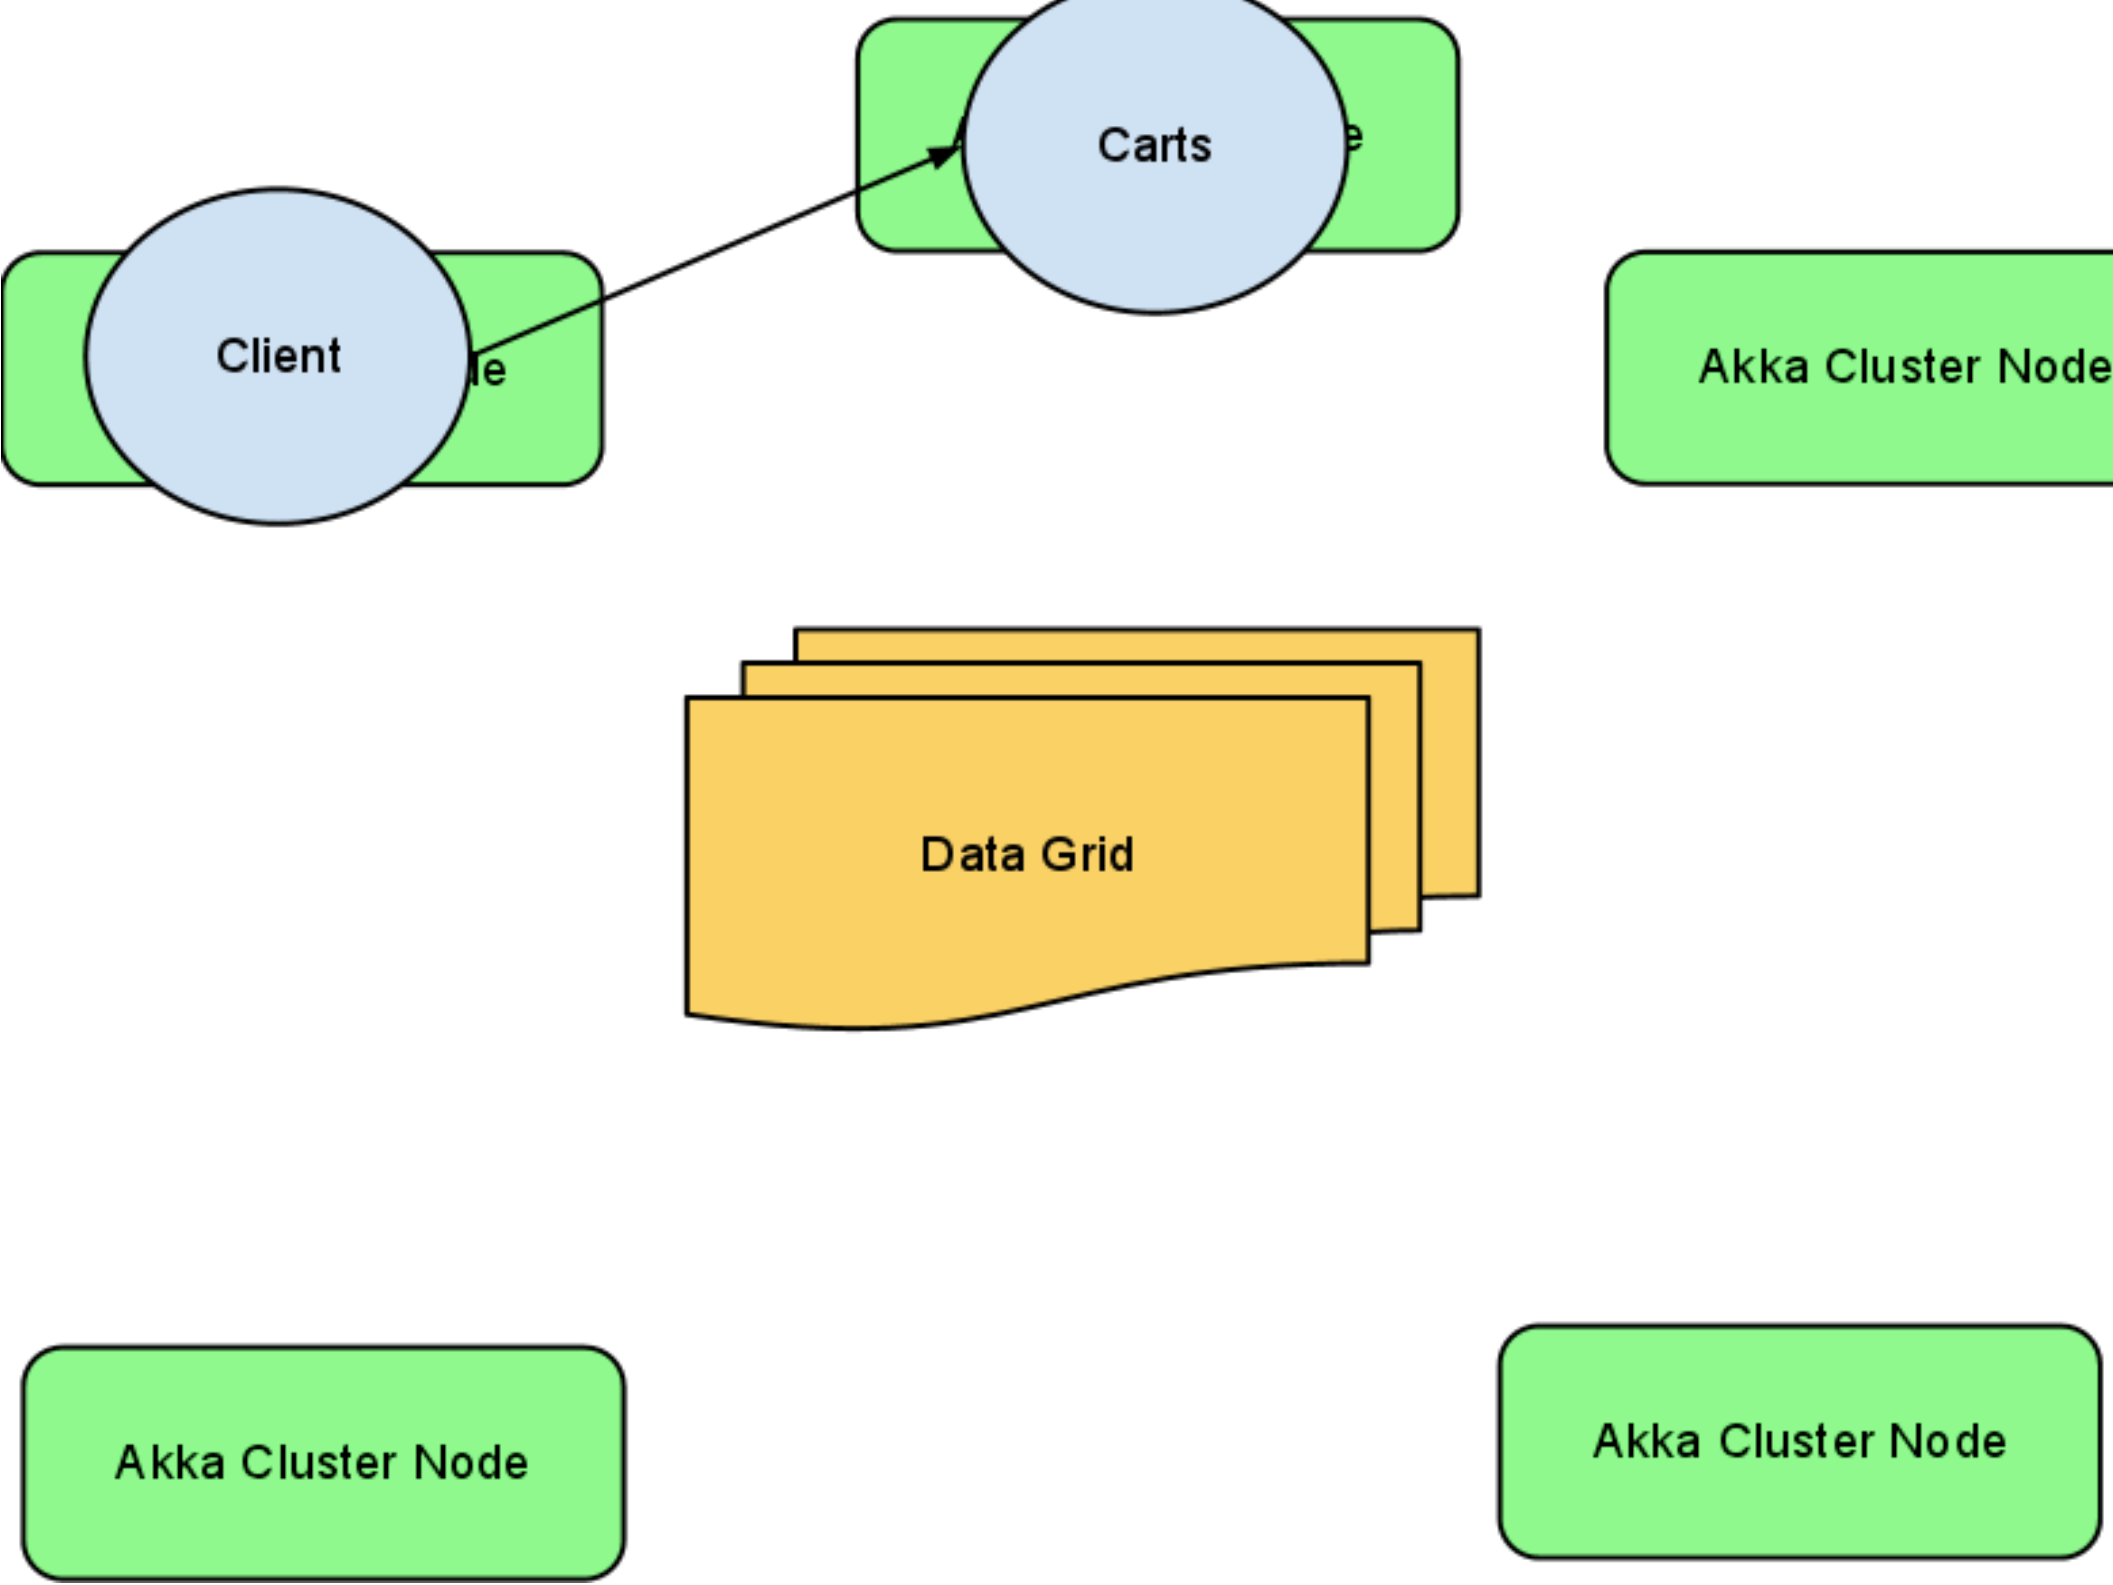
Akka Cluster Node


Akka Cluster Node




```
akka {  
  actor {  
    deployment {  
      carts {  
        clustered {  
          stateless = off  
          replicas = 3  
        }  
      }  
    }  
  }  
}
```









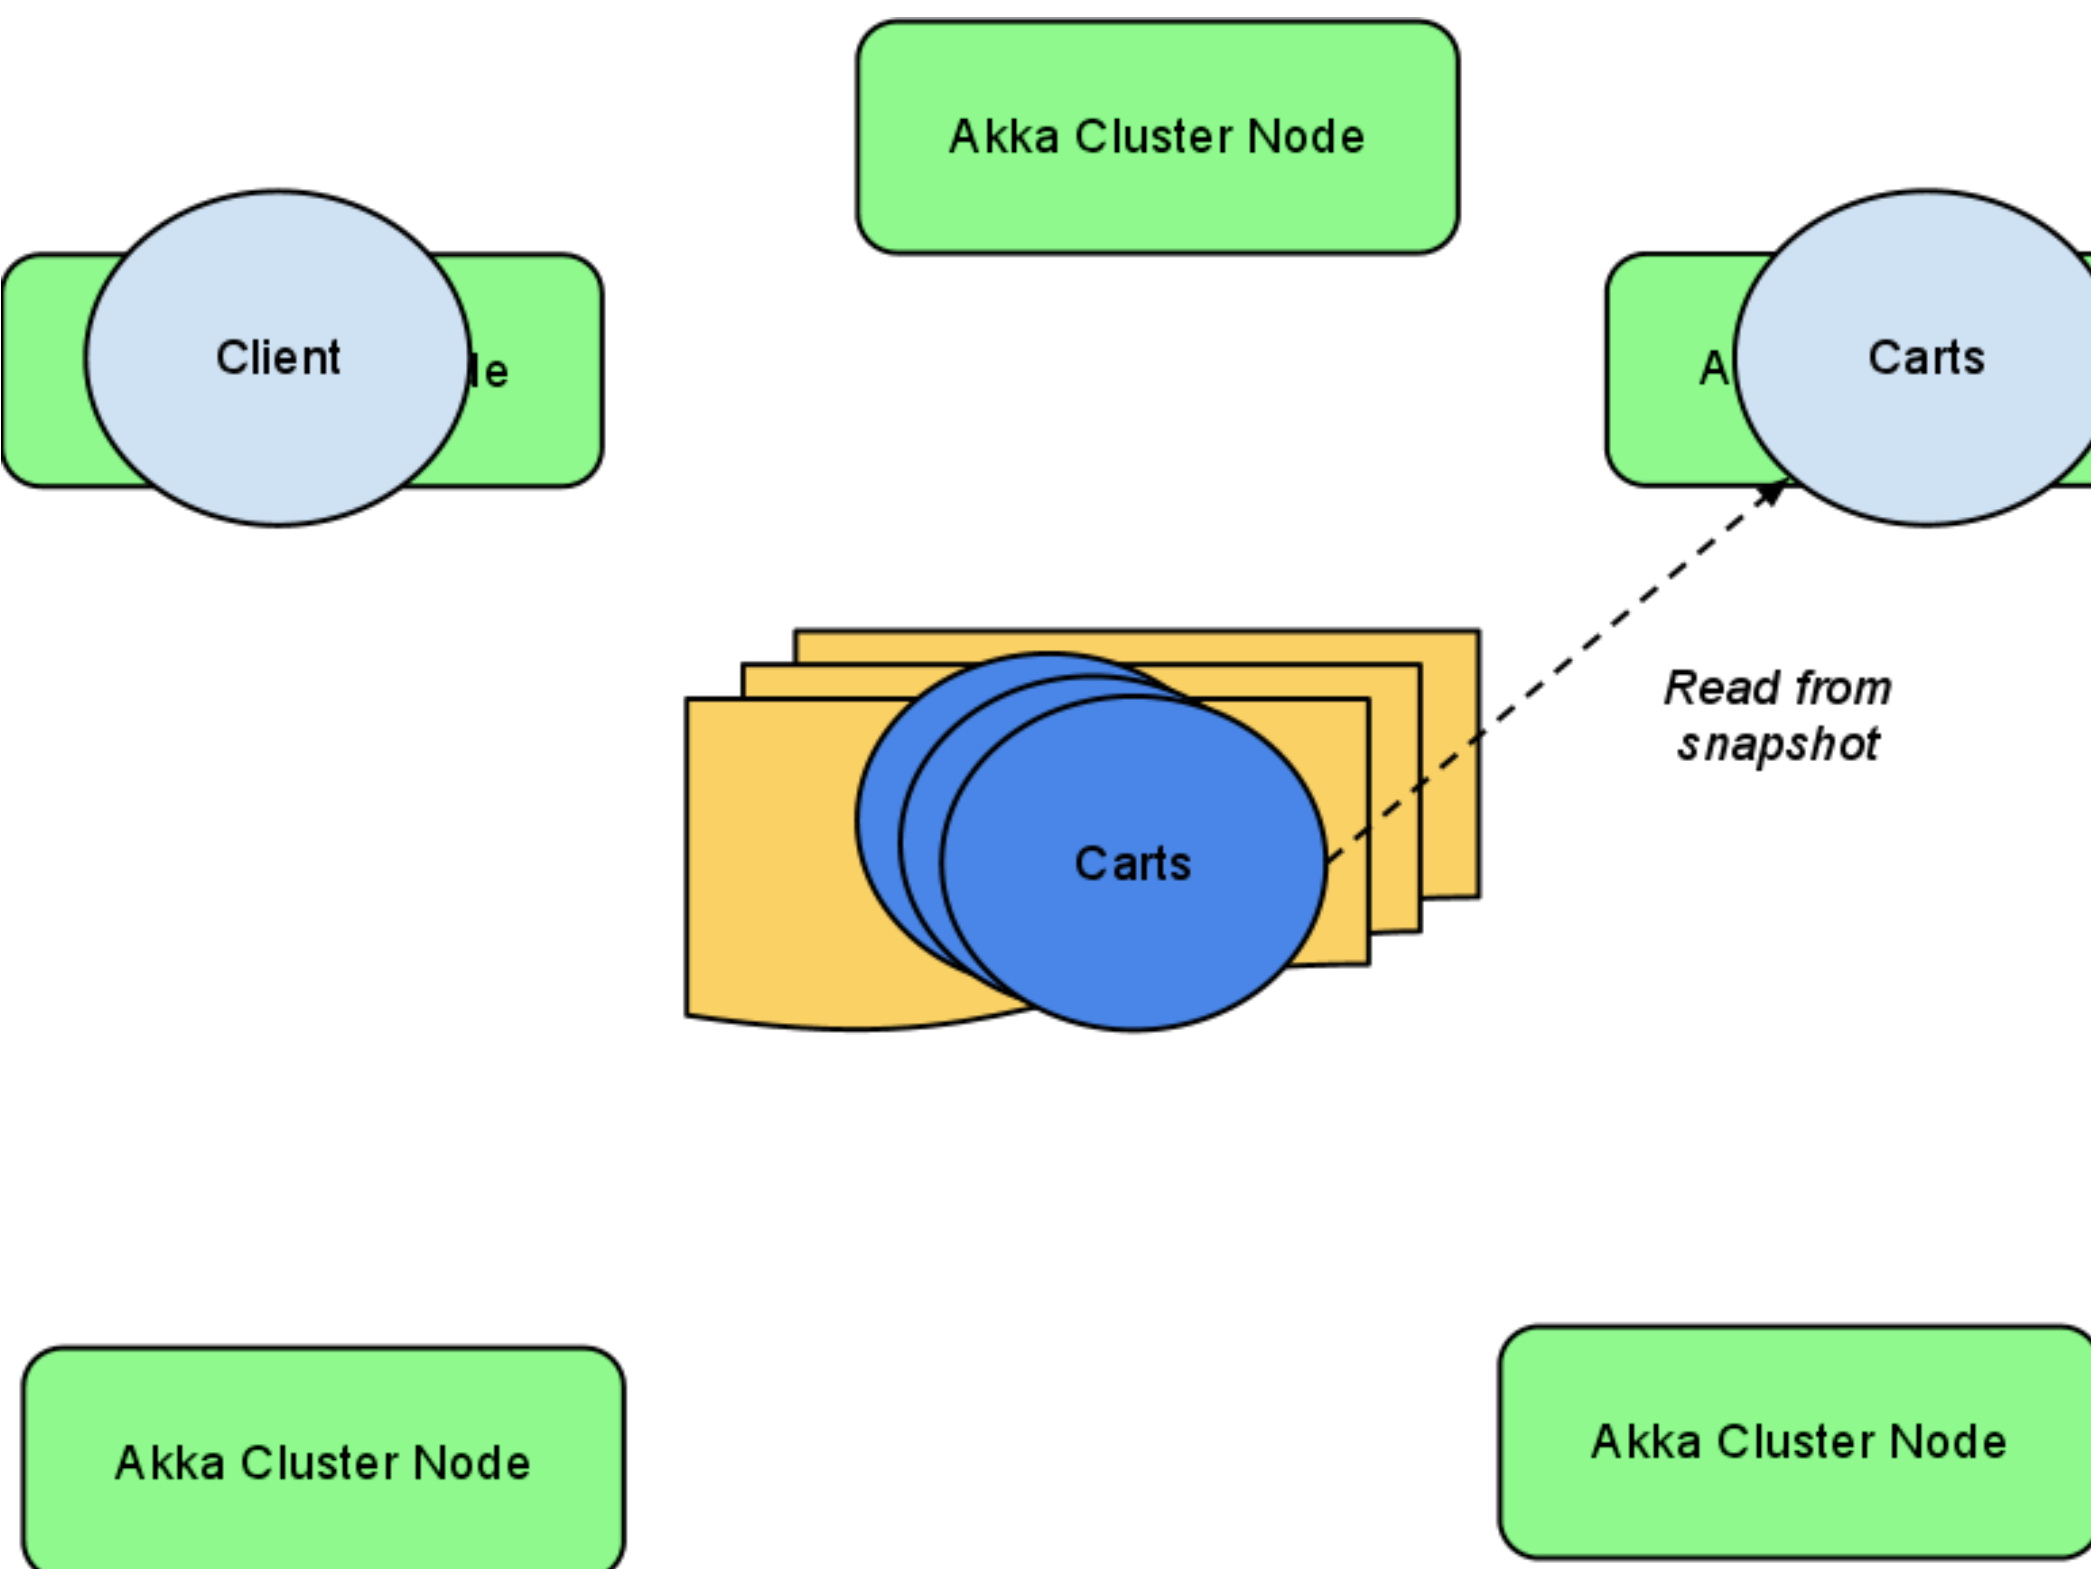


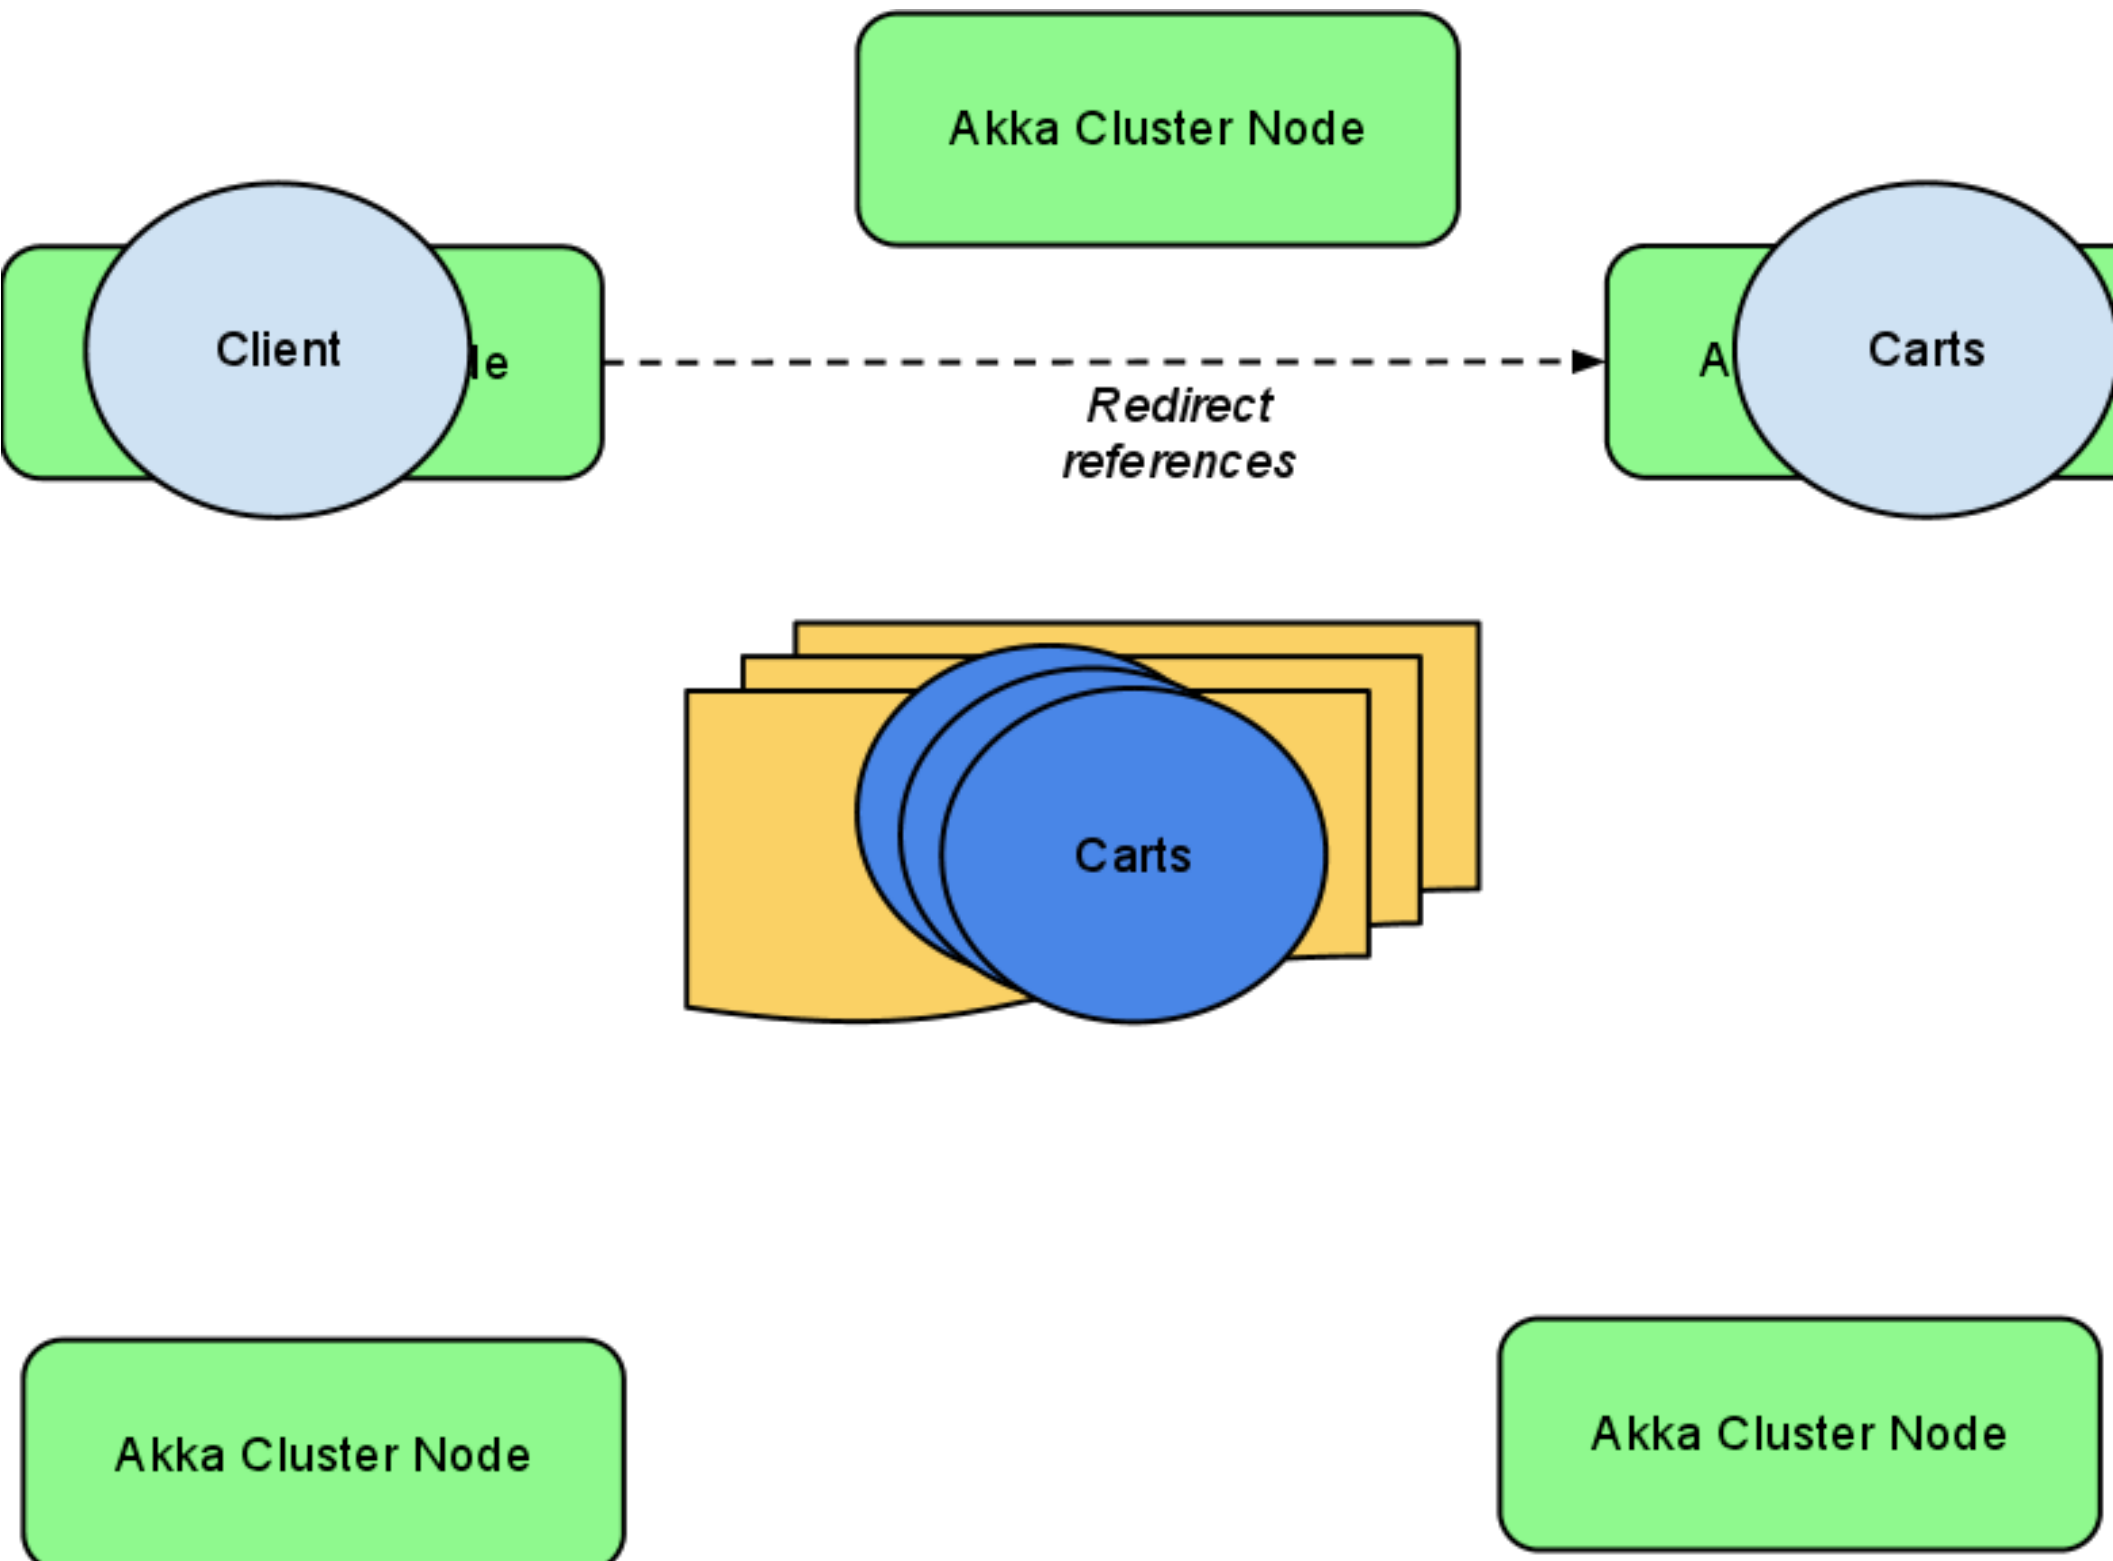












For **power** users: Cluster API

```
import Actor.cluster
```

For **power** users: Cluster API

```
import Actor.cluster  
cluster.start()
```

For **power** users: Cluster API

```
import Actor.cluster  
cluster.start()  
cluster.shutdown()
```

For **power** users: Cluster API

```
import Actor.cluster
cluster.start()
cluster.shutdown()
cluster register (new ChangeListener {
  def nodeConnected(node: String, client: ClusterNode) {
    ...
  }
  ...
})
```

For **power** users: Cluster API

```
import Actor.cluster
cluster.start()
cluster.shutdown()
cluster register (new ChangeListener {
  def nodeConnected(node: String, client: ClusterNode) {
    ...
  }
  ...
})
cluster store actorRef
cluster remove actorAddress
```

For **power** users: Cluster API

```
import Actor.cluster
cluster.start()
cluster.shutdown()
cluster.register (new ChangeListener {
  def nodeConnected(node: String, client: ClusterNode) {
    ...
  }
  ...
})
cluster.store actorRef
cluster.remove actorAddress
val actorRef = cluster.use actorAddress
val actorRef = cluster.ref (actorAddress, router)
```


For **power** users: Cluster API

```
import Actor.cluster
cluster.start()
cluster.shutdown()
cluster.register (new ChangeListener {
  def nodeConnected(node: String, client: ClusterNode) {
    ...
  }
  ...
})
cluster.store actorRef
cluster.remove actorAddress
val actorRef = cluster.use actorAddress
val actorRef = cluster.ref (actorAddress, router)
cluster.migrate (fromNode, toNode, actorAddress)
```

For **power** users: Cluster API

```
import Actor.cluster
cluster.start()
cluster.shutdown()
cluster.register (new ChangeListener {
  def nodeConnected(node: String, client: ClusterNode) {
    ...
  }
  ...
})
cluster.store actorRef
cluster.remove actorAddress
val actorRef = cluster.use actorAddress
val actorRef = cluster.ref (actorAddress, router)
cluster.migrate (fromNode, toNode, actorAddress)
cluster.send (() => { ... }, nrReplicas) map (_.result)
```

Routers

- Direct
- Random
- Round Robin
- Least CPU (soon)
- Least RAM (soon)
- Least Messages (soon)
- Custom

Durable Mailboxes

- File-based
- Redis-based
- Beanstalk-based
- MongoDB-based
- Zookeeper-based
- Cassandra-based (soon)
- AMQP-based (soon)
- JMS-based (soon)

AKKA 2.x

Hakka's Paradise

<http://akka.io>

Roadmap:

- 2.0
 - Location Transparency
 - API Cleanup
 - Configuration-based deployment
 - Failure Detection
 - Improved Supervisors + Lifecycle management
- 2.1
 - Clustered Elastic Akka
 - TO THE CLOUD!
- 2.2
 - Clustered management, replication, migration of stateful actors

EOF