

# Event-driven Programming in Clojure

Zach Tellman  
Runa  
@ztellman

# synchronous is simple

```
def handle_request(request)
    result = query_database(request)
    return generate_response(result)
```

- easy to read and reason about
- errors can be handled at any scope

# asynchronous is complex


```
def handle_request(request)
  query_database(request,
    error => {
      handle_error(error)
    },
    result => {
      response = generate_response(result)
      send_response(response)
    })
```

- difficult to read and reason about
- errors must be handled locally

# the importance of being concurrent


- we often spend more time **waiting** on data than we do **computing** our response
- we can respond more quickly if we wait for multiple things at once

# synchronous concurrency




- still pretty simple
- requires one thread per task, plus the original thread


# but threads don't grow on trees



# nested concurrency




# two tasks, one pool





# two tasks, one pool



deadlock!


# asynchronous exposes an inherent complexity

- everything is concurrent by default
- a less leaky abstraction (for highly coupled concurrent problems)

# two great tastes

- these are not mutually exclusive approaches
- interop should be easy
- we want to simplify the asynchronous approach as much as possible

# publisher/subscriber




# a naive implementation

```
subscribers = [ ]
```

```
def subscribe(callback)  
    subscribers.add(callback)
```

```
def publish(event)  
    foreach s in subscribers  
        s(event)
```

# stop the presses



The same callback can be executed on different threads with different data

# ordered vs. unordered


- mouse clicks and keystrokes need to be handled one at a time and in-order
- logging we can be less fussy about

# a software koan

If an event has no subscribers, does it have  
any effect?



# a simple example



# a simple example

```
def handle_request(request)
    if whitelisted?(request.url)
        forward_request(request)
    else
        forbid_request(request)
```

# the node.js approach

- receive headers
- send whitelist request, subscribe to response
- if it's whitelisted, forward the headers and subscribe to the body
- but what if the body arrived first?

# the erlang approach

mailbox

{whitelist, true}

{body, ...}

{header, ...}

receive {header, Header}

receive {whitelist, Response}

receive {body, Body}

# push vs. pull

- the pub/sub model is **push**
- the erlang mailbox model is **pull**
- just because we get an event doesn't mean we know what to do with it

# the story so far

- event streams can be **ordered** or **unordered**
- events can be published using **push** or consumed using **pull**

# events in clojure

we need to consider:

- immutable data structures
- software transactional memory
- thread-agnosticism

# immutability

all that, and a bag of chips



# STM

- publishing is a side-effect
- side-effects inside transactions can be hazardous

# threads

- Clojure is thread-agnostic
- what happens if an event is published before another thread subscribes?

# a wish list

- support for push and pull event consumption
- efficient mechanisms for ordered and unordered event streams
- plays nicely with transactions

# channels



- represents a **stream of messages**
- messages are consumed by callbacks
- if there are no callbacks, messages wait in the queue
- queue is transactional

# unordered messages

(`receive-all` channel callback)

- callback is invoked by thread that enqueues the message
- bypasses the queue entirely

# ordered messages

(`receive-in-order` channel callback)

(`receive` channel callback)

pull messages from the queue one at a time

# closing channels

(`close` channel)

(`on-closed` channel callback)

(`on-drained` channel callback)

- when a channel is **closed**, no further messages can be enqueued
- when a closed channel is empty, it is **drained**

# channel errors

(**error!** channel exception)

(**on-error** channel callback)

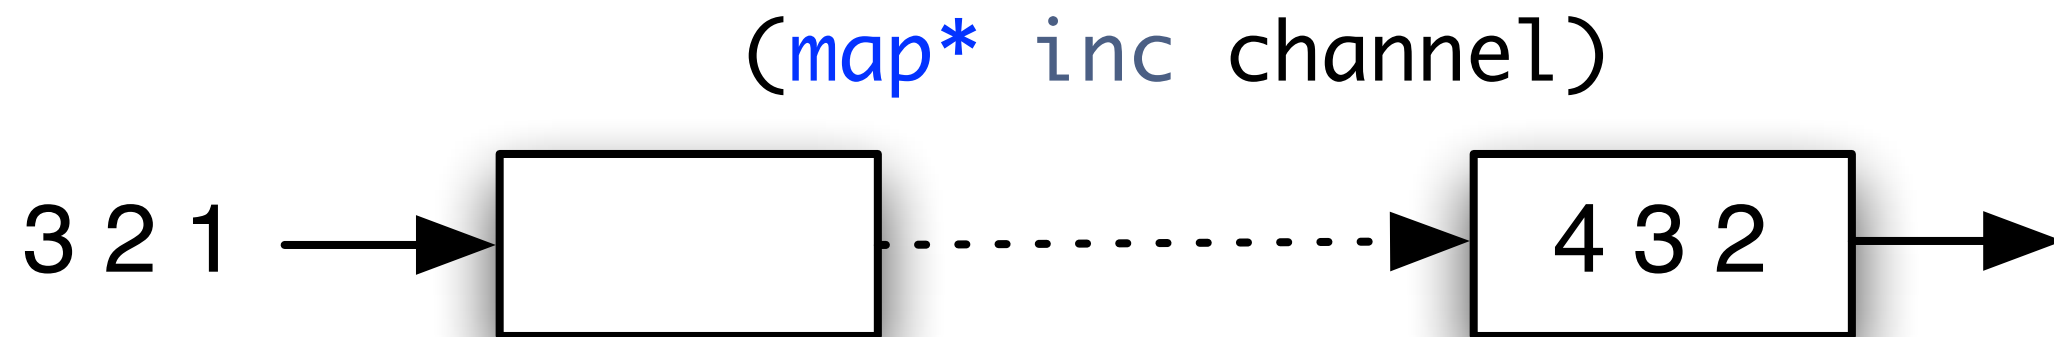
errors also close the channel



# channel operators

- familiar operators such as `map*`, `filter*`, `reduce*`, `take*`, and `partition*`
- less familiar operators such as `sample-every*`, `partition-every*`, and `siphon`

# channel operators

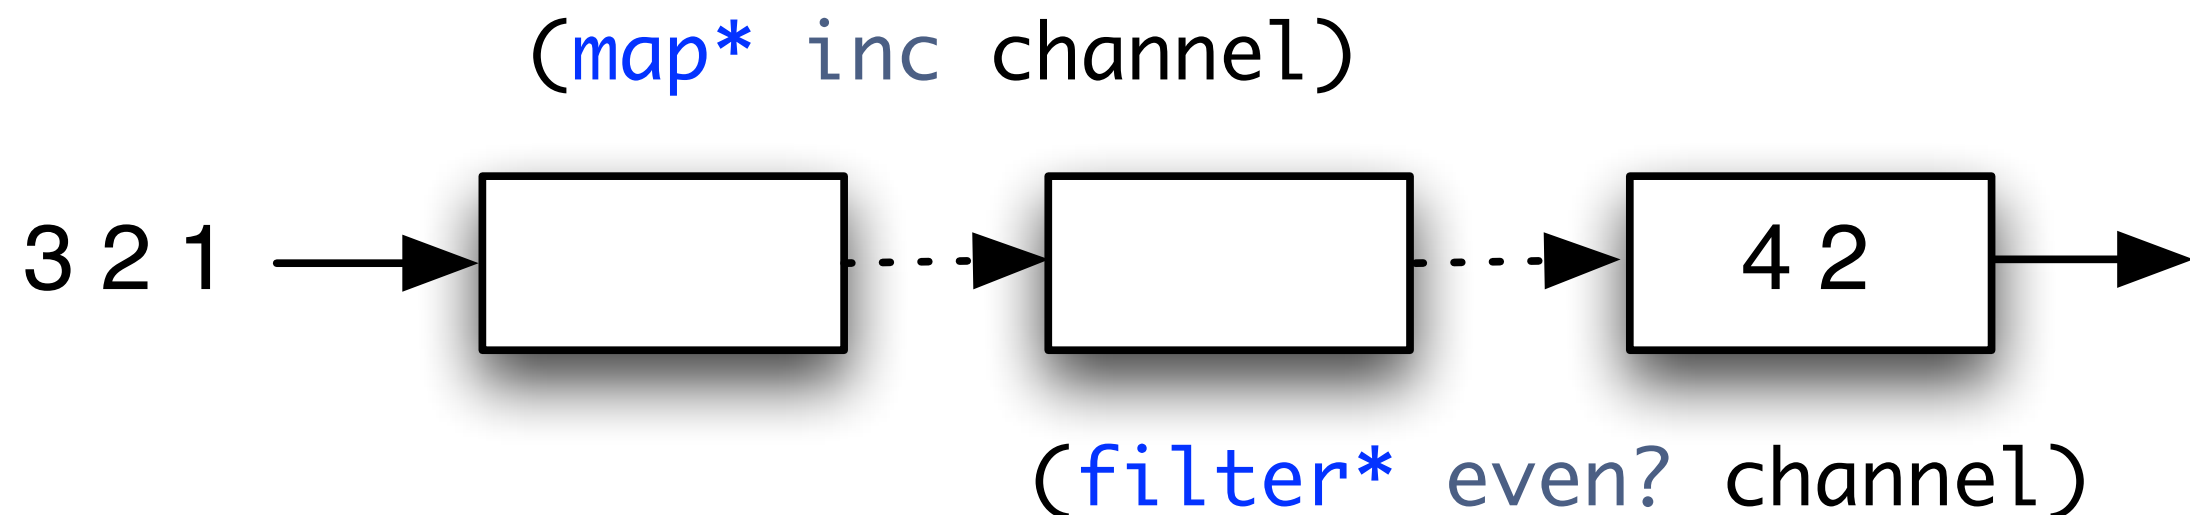


to create a derivative channel, you must consume messages from the source channel

# channel operators

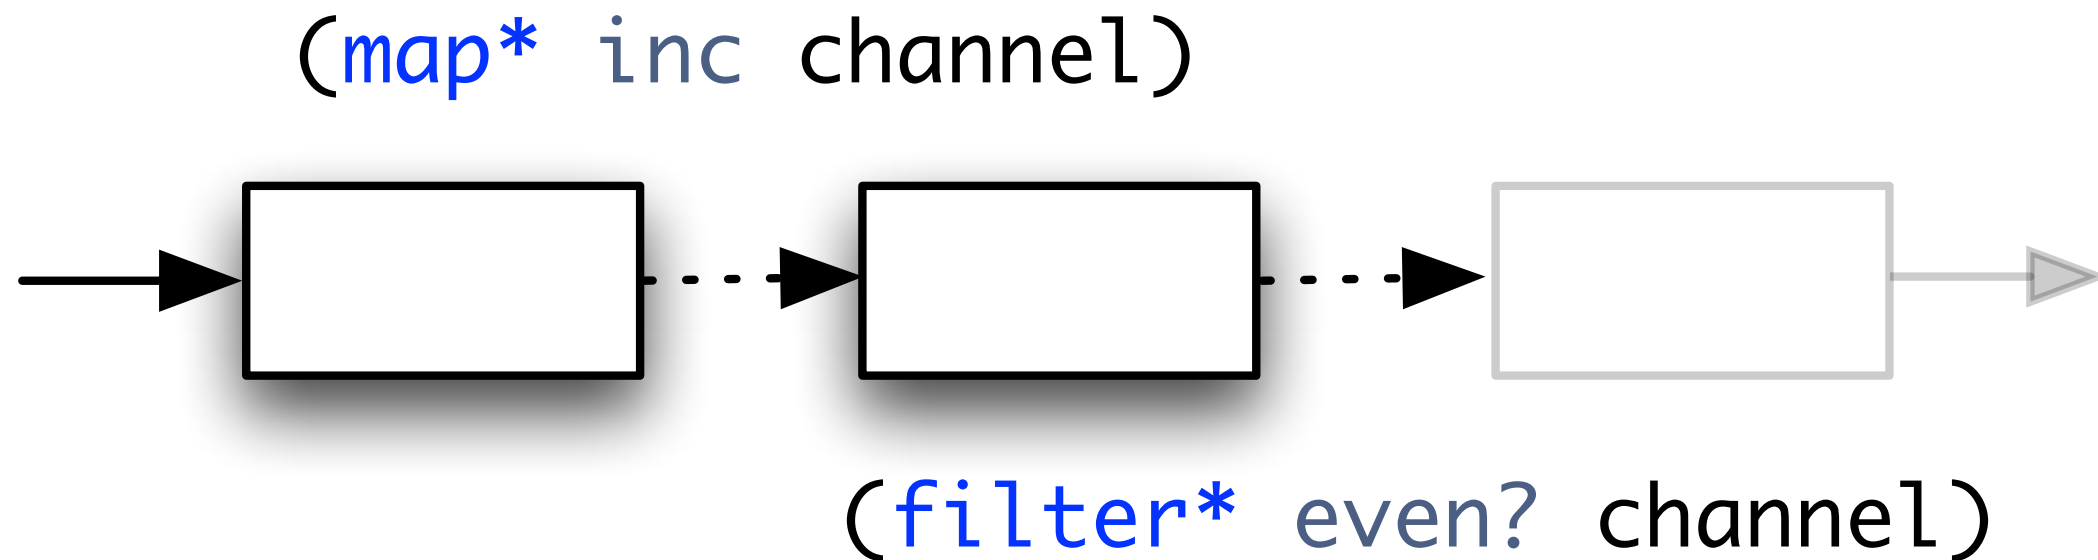
```
(->> channel (map* inc) (filter* even?))
```

# channel operators

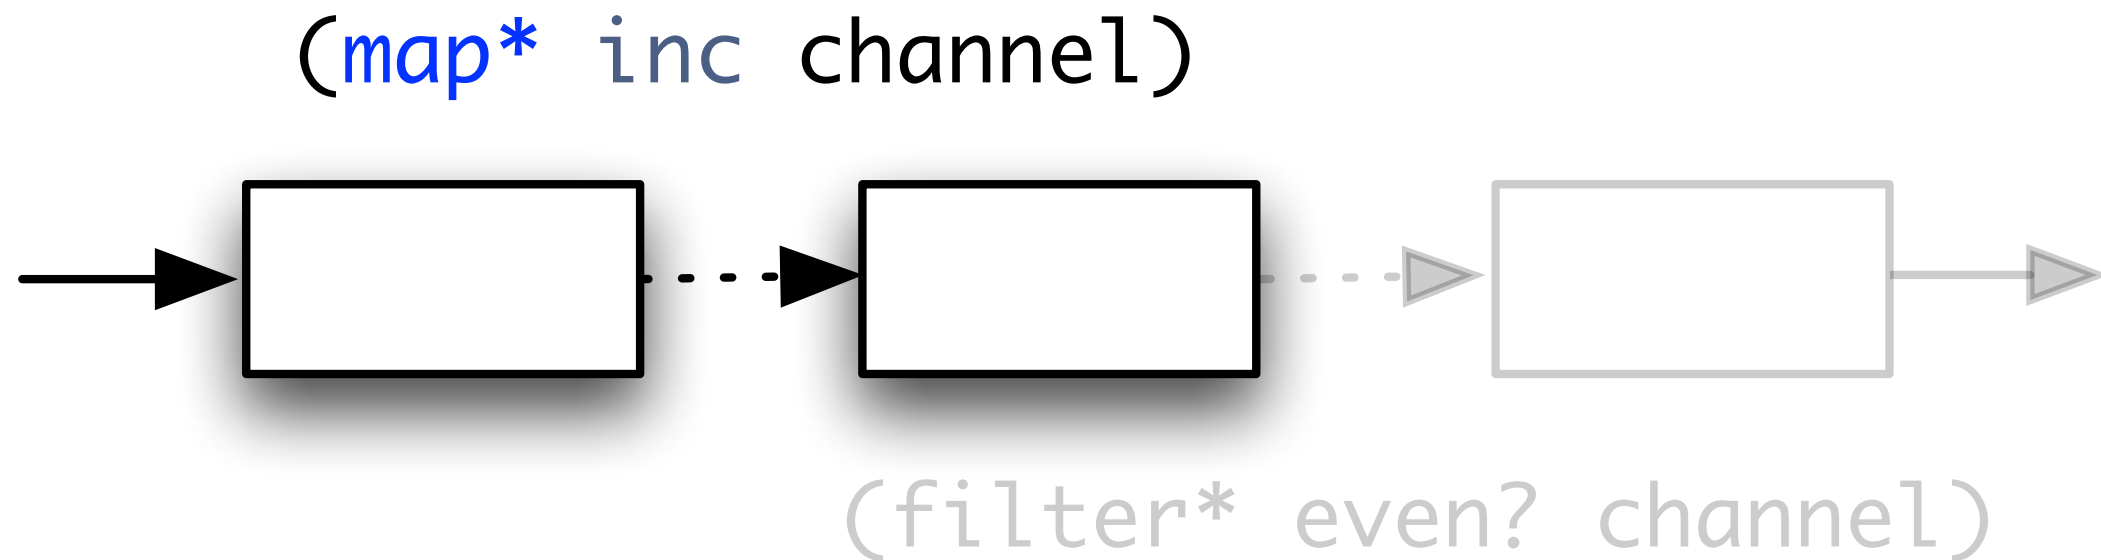


but what happens when we close the right-most channel?

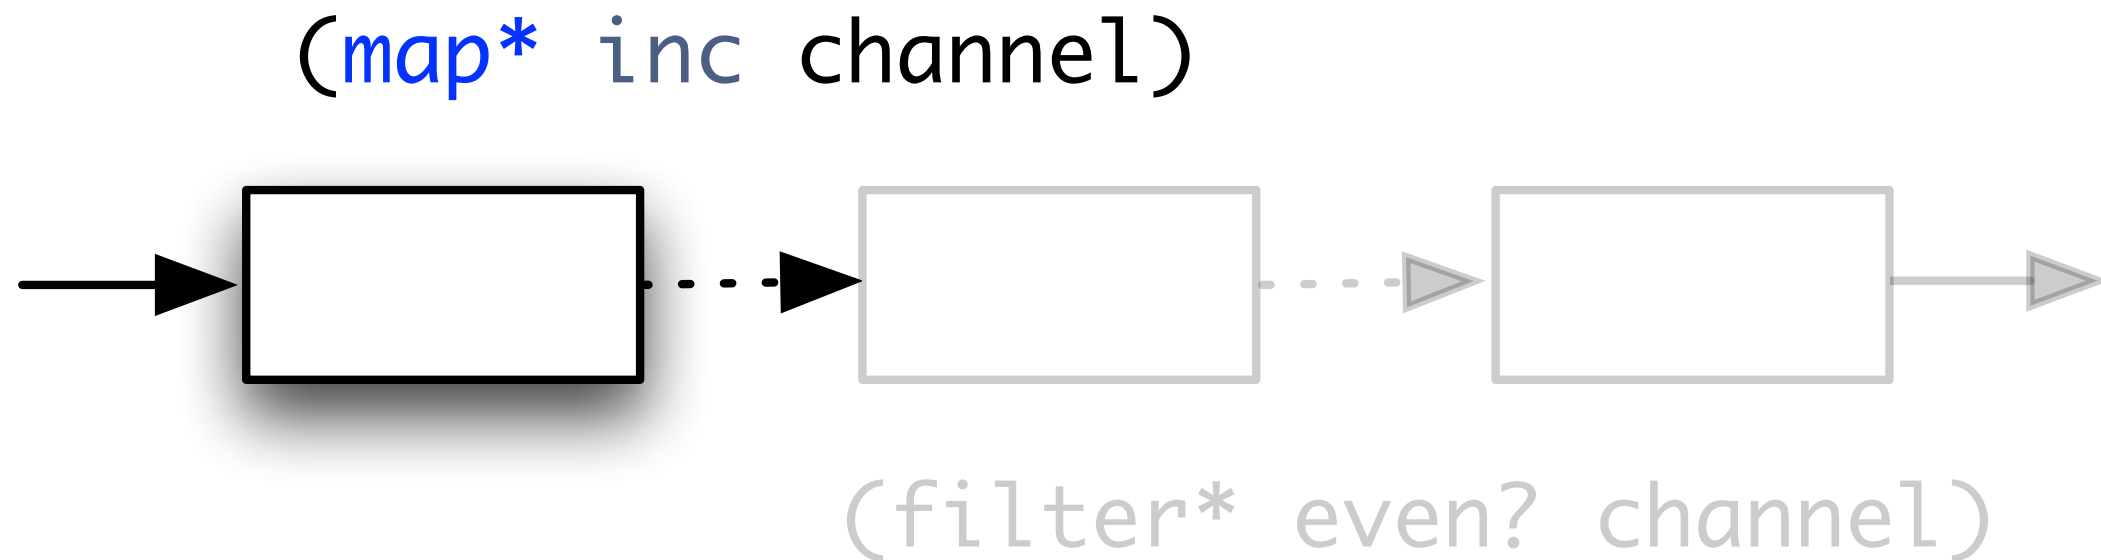
# backpropagation



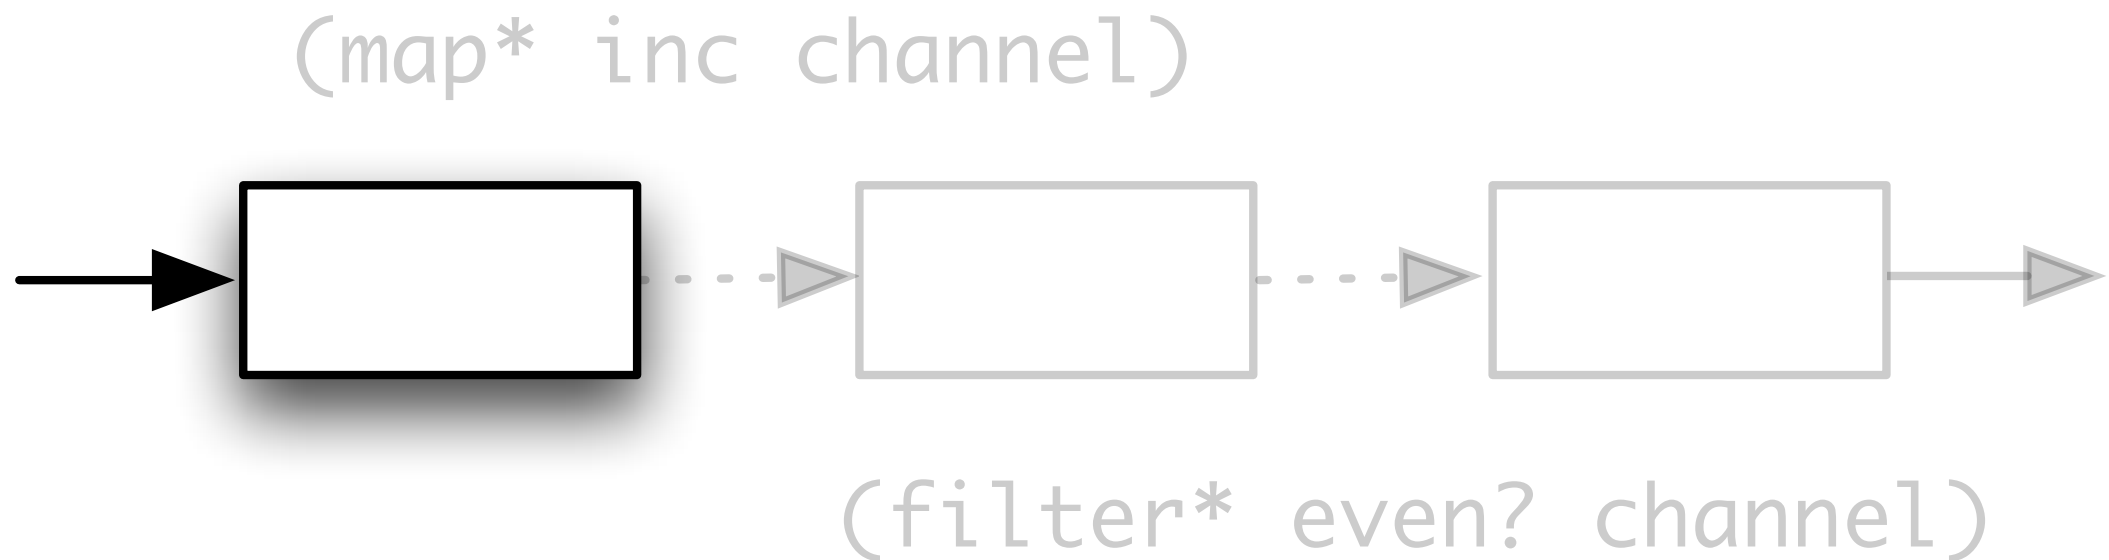
# backpropagation



# backpropagation

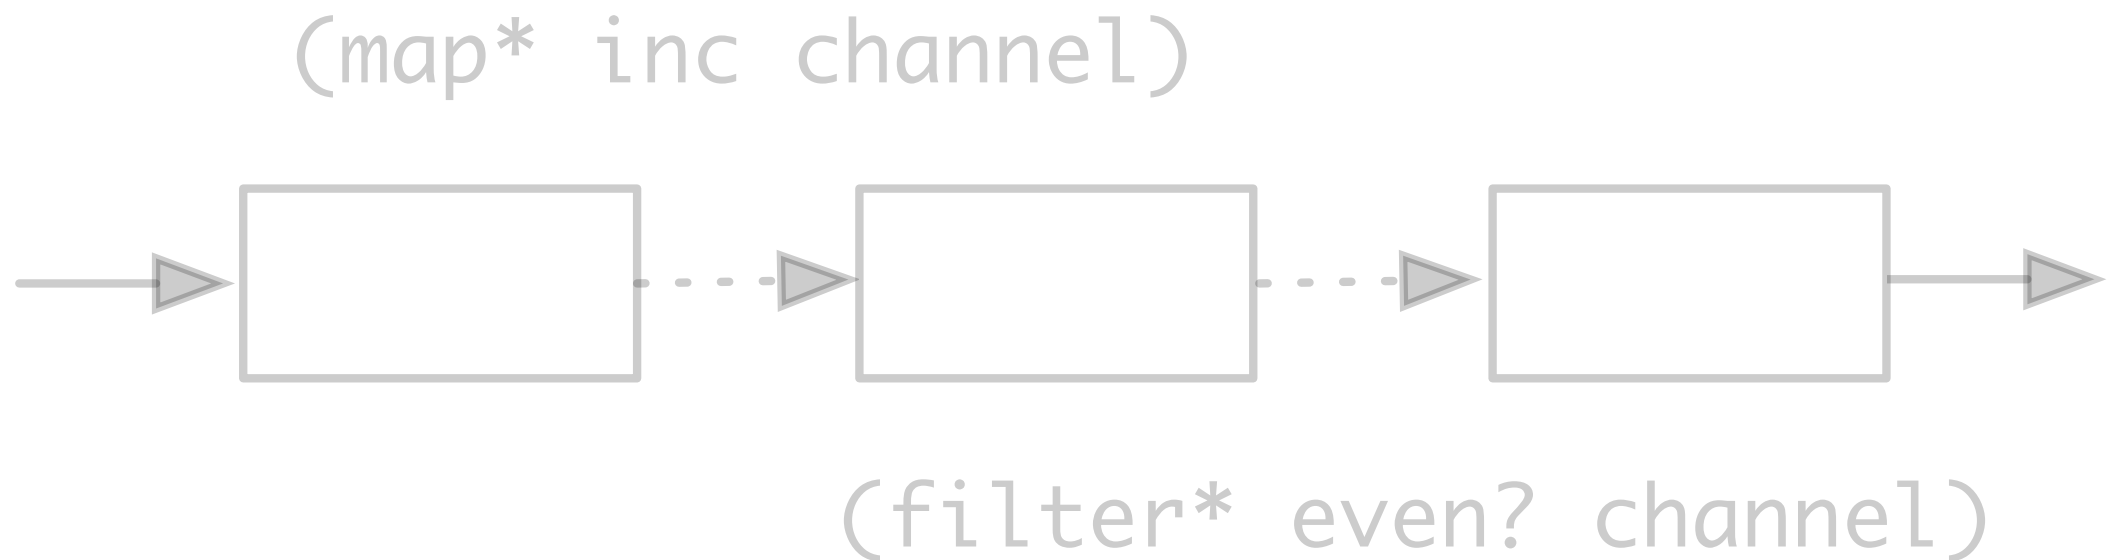


# backpropagation

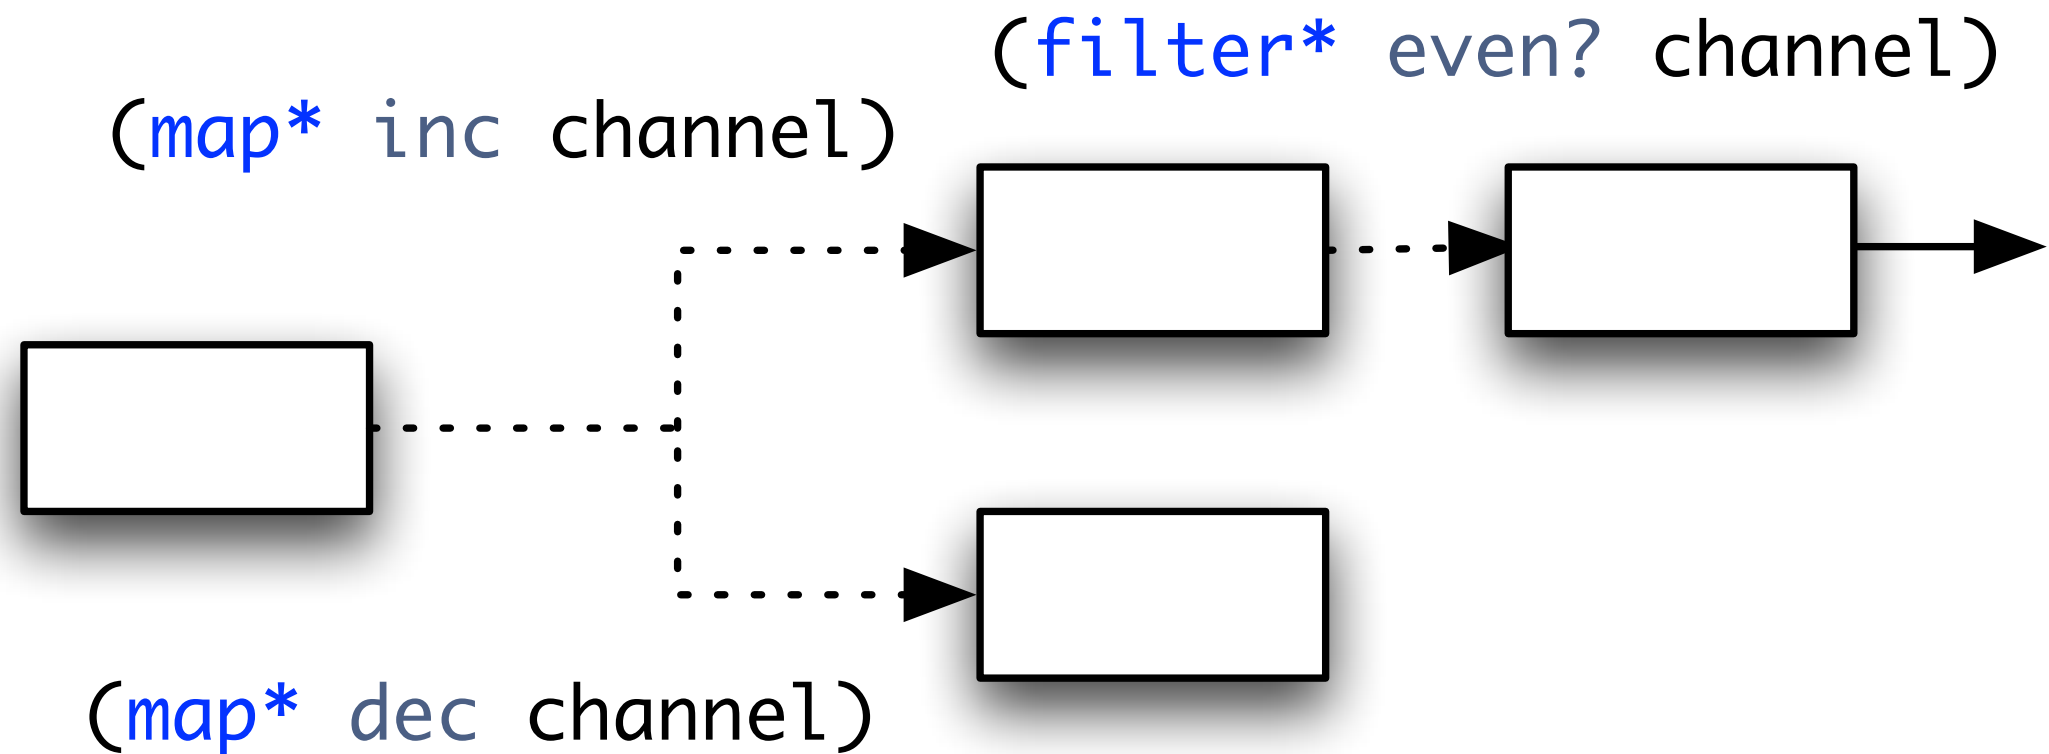




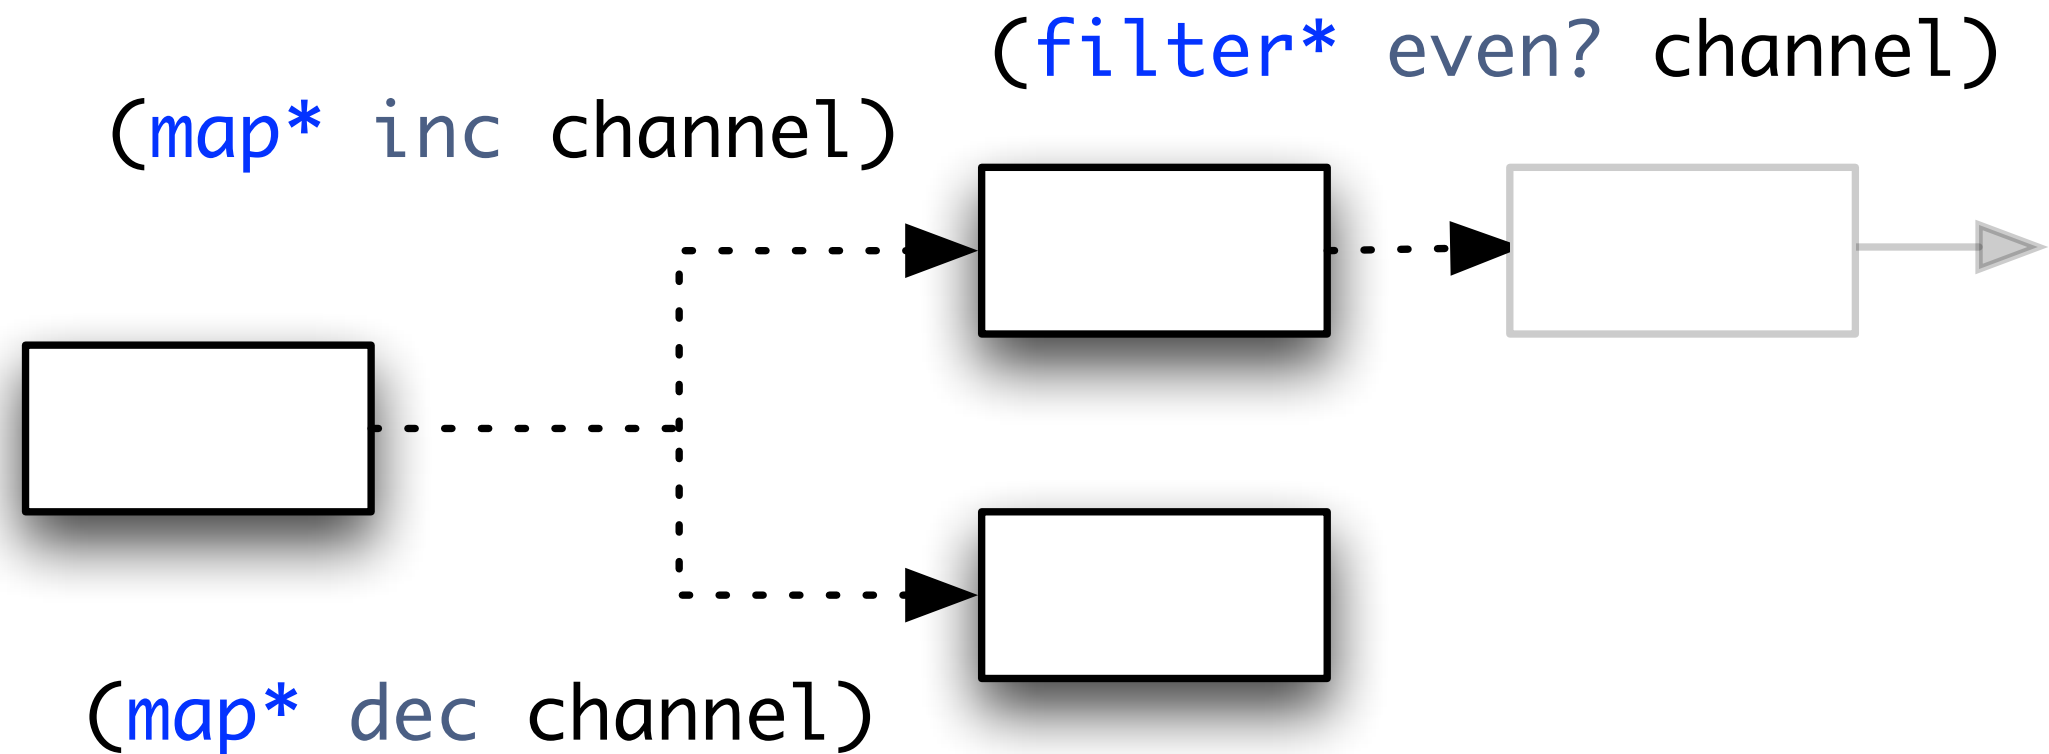
# backpropagation



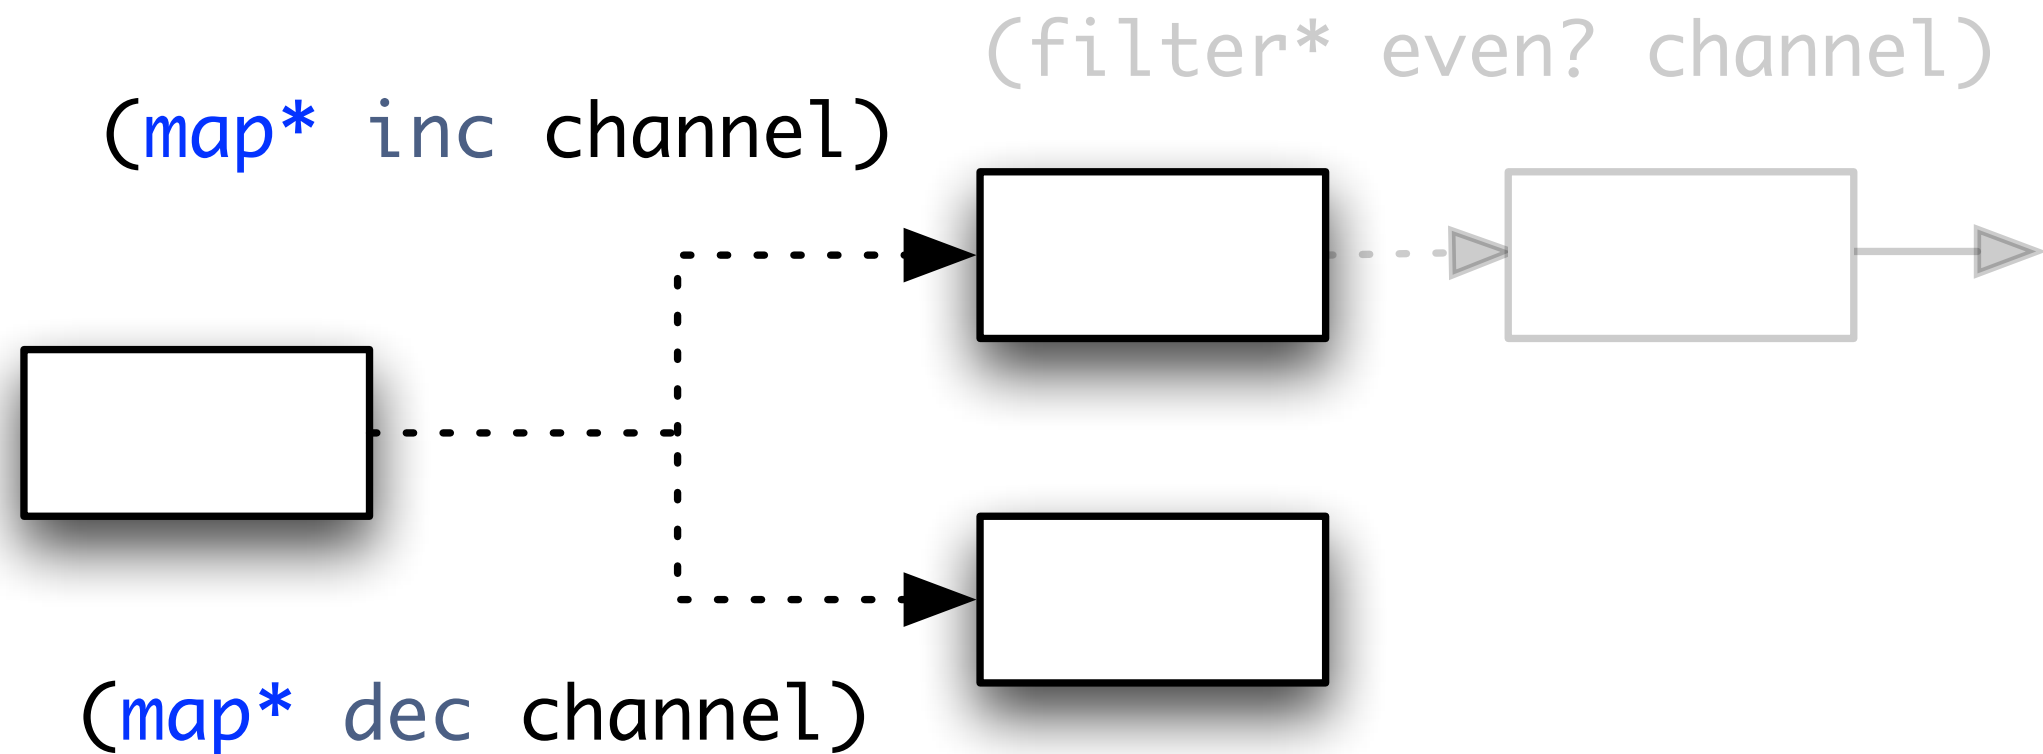
# backpropagation



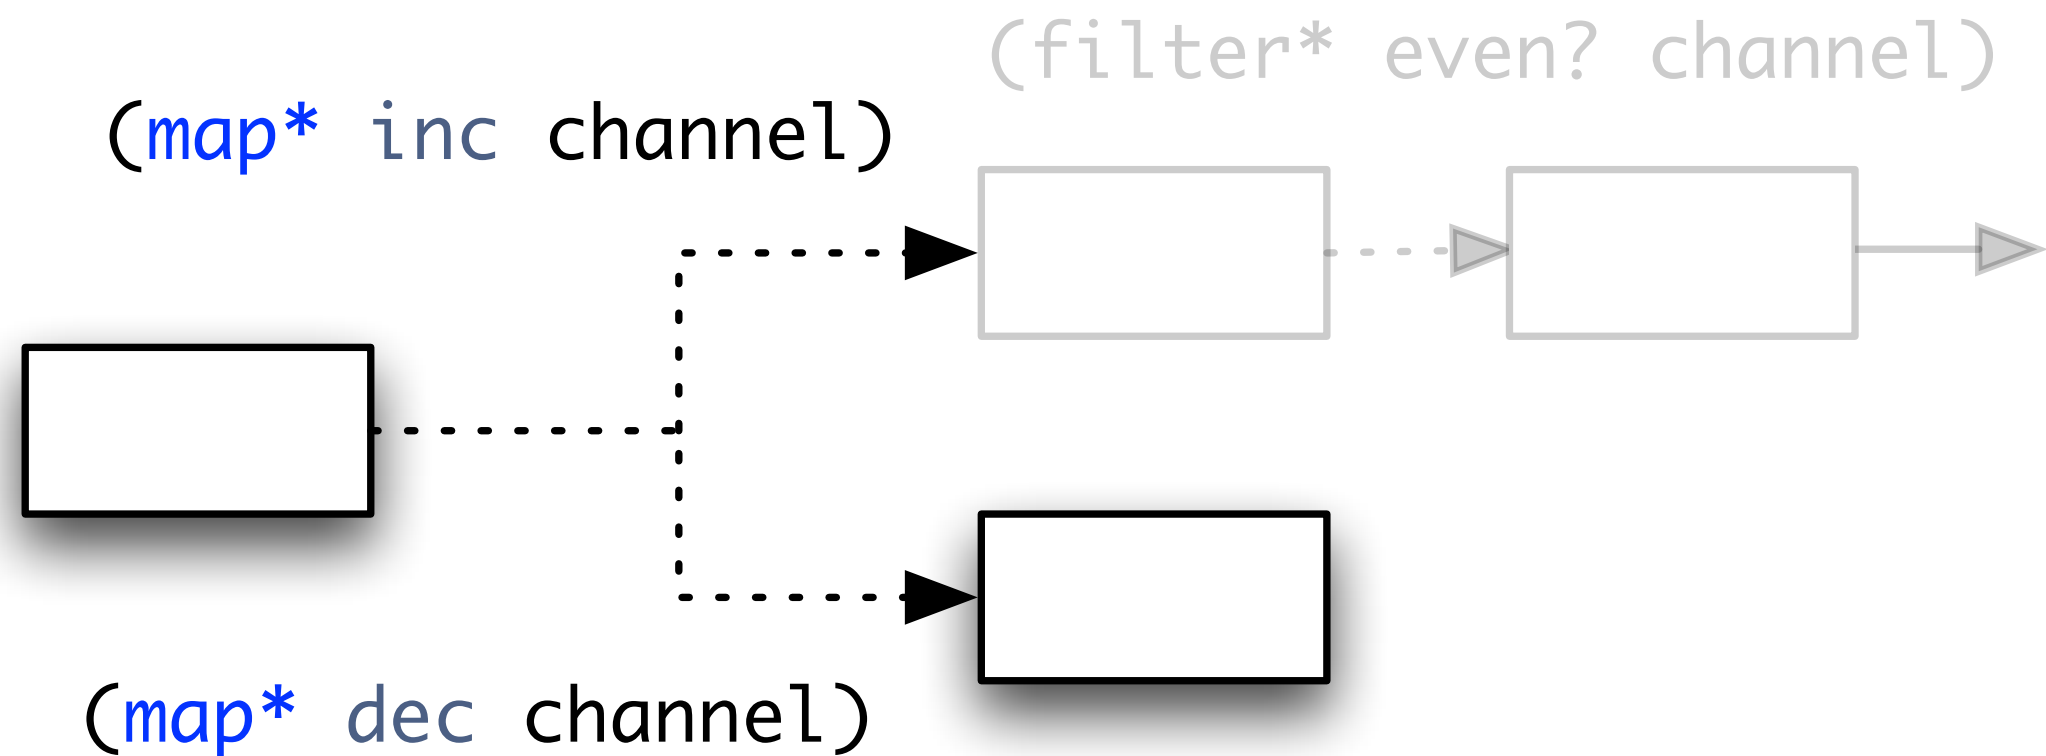
# backpropagation



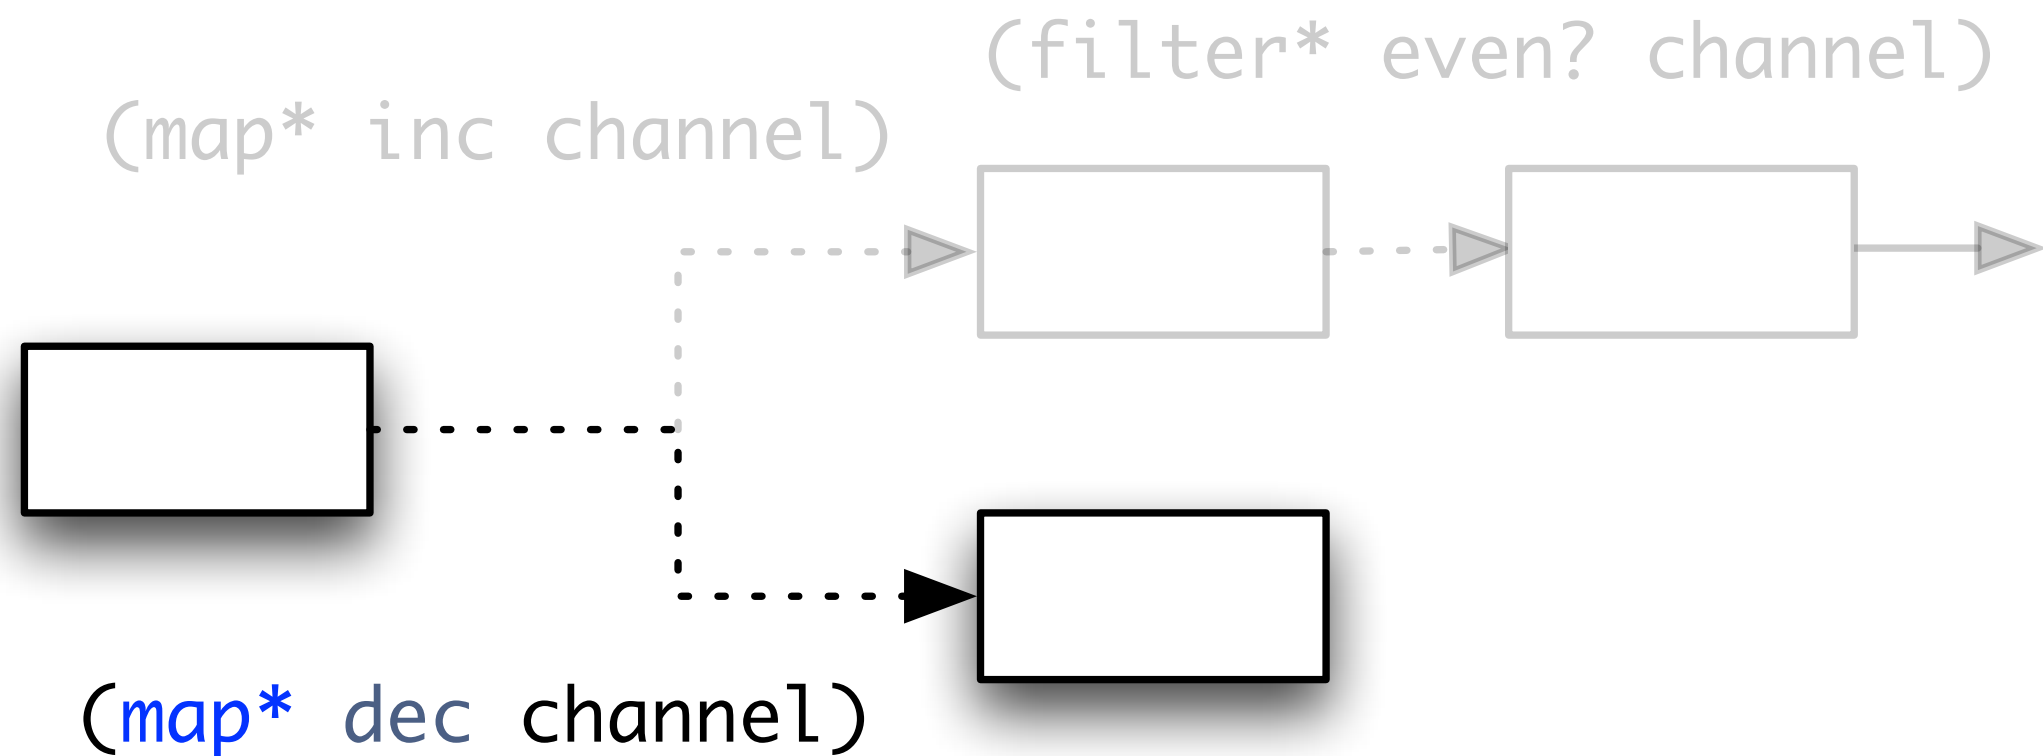
# backpropagation



# backpropagation



# backpropagation



# channels and seqs

- in channels, messages exist once they're enqueued
- in lazy-seqs, elements exist once they're consumed

# channels and seqs

- in seqs, the same operator will always return the same result (map is idempotent)
- in channels, the same operator will not always return the same result (map\* is not idempotent)



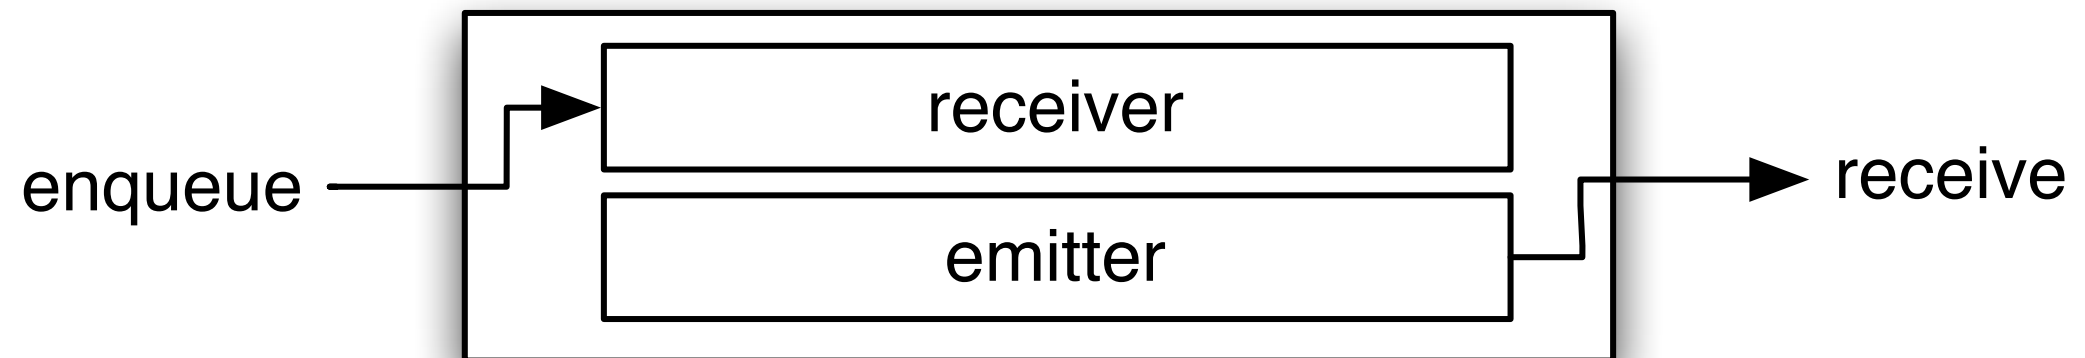
# forking channels

(`fork` channel)


- returns a duplicate channel that can be independently consumed
- closing the original channel will close the forked channel, but not vice-versa

# splicing channels

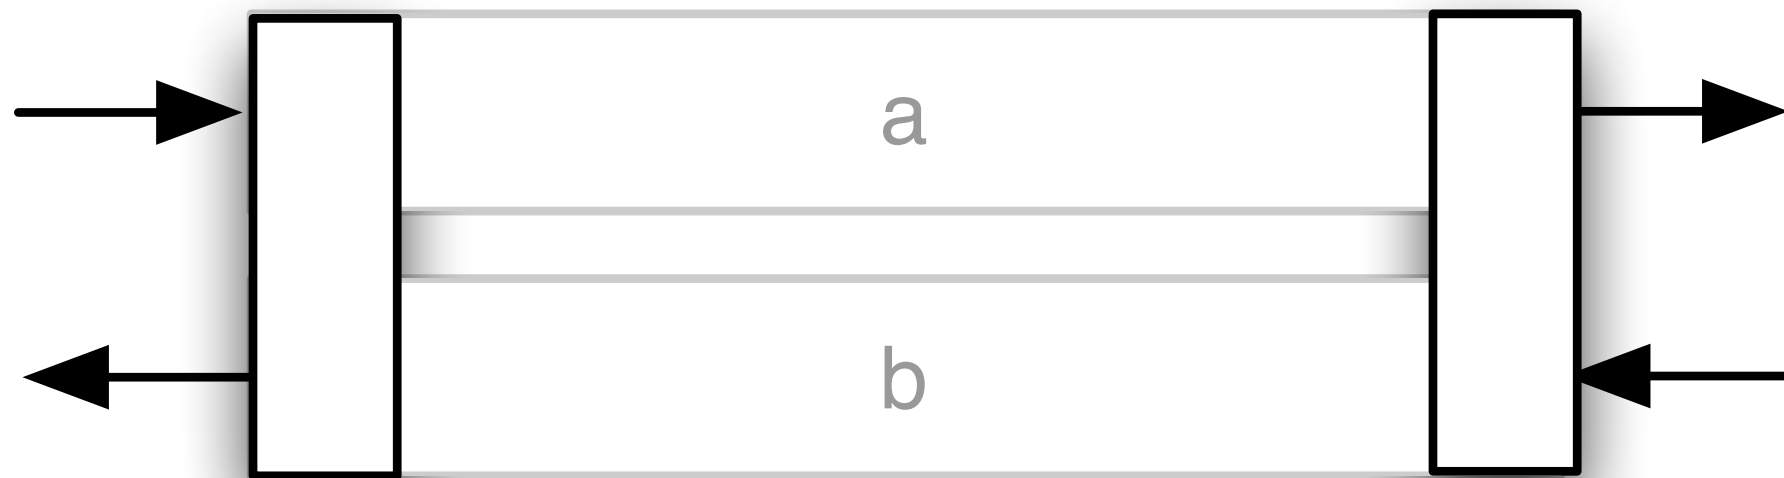
(splice emitter receiver)



# splicing channels



# splicing channels




(splice a b)




(splice b a)

# a simple chat server

broadcaster



sockets




# a simple chat server


(`siphon` broadcaster socket)

(`siphon` socket broadcaster)

# a simple chat server




# a simple chat server





# a simple chat server



# result channels

- represents a single outcome: a **result** or an **error**
- only a single value can be enqueued, but that value can be received multiple times
- this is our bridge to the synchronous world

# synchronous results

(`wait-for-result` `result-channel` `timeout`)

`@result-channel`

if the `result-channel` emits an error, synchronous  
accessors will throw an exception

# asynchronous contamination

returns a sum:

(+ 2 @(request-a-number))

# asynchronous contamination

returns a sum:

```
(+ 2 @(request-a-number))
```

returns an eventual sum:

```
(async (+ 2 (request-a-number)))
```


# a request workflow

- receive a request
- fetch the **session** data
- if the session doesn't have the **user history**, fetch it
- if the session doesn't have the **user location**, fetch it
- generate a response based on the request, session, history, and location

# a request workflow

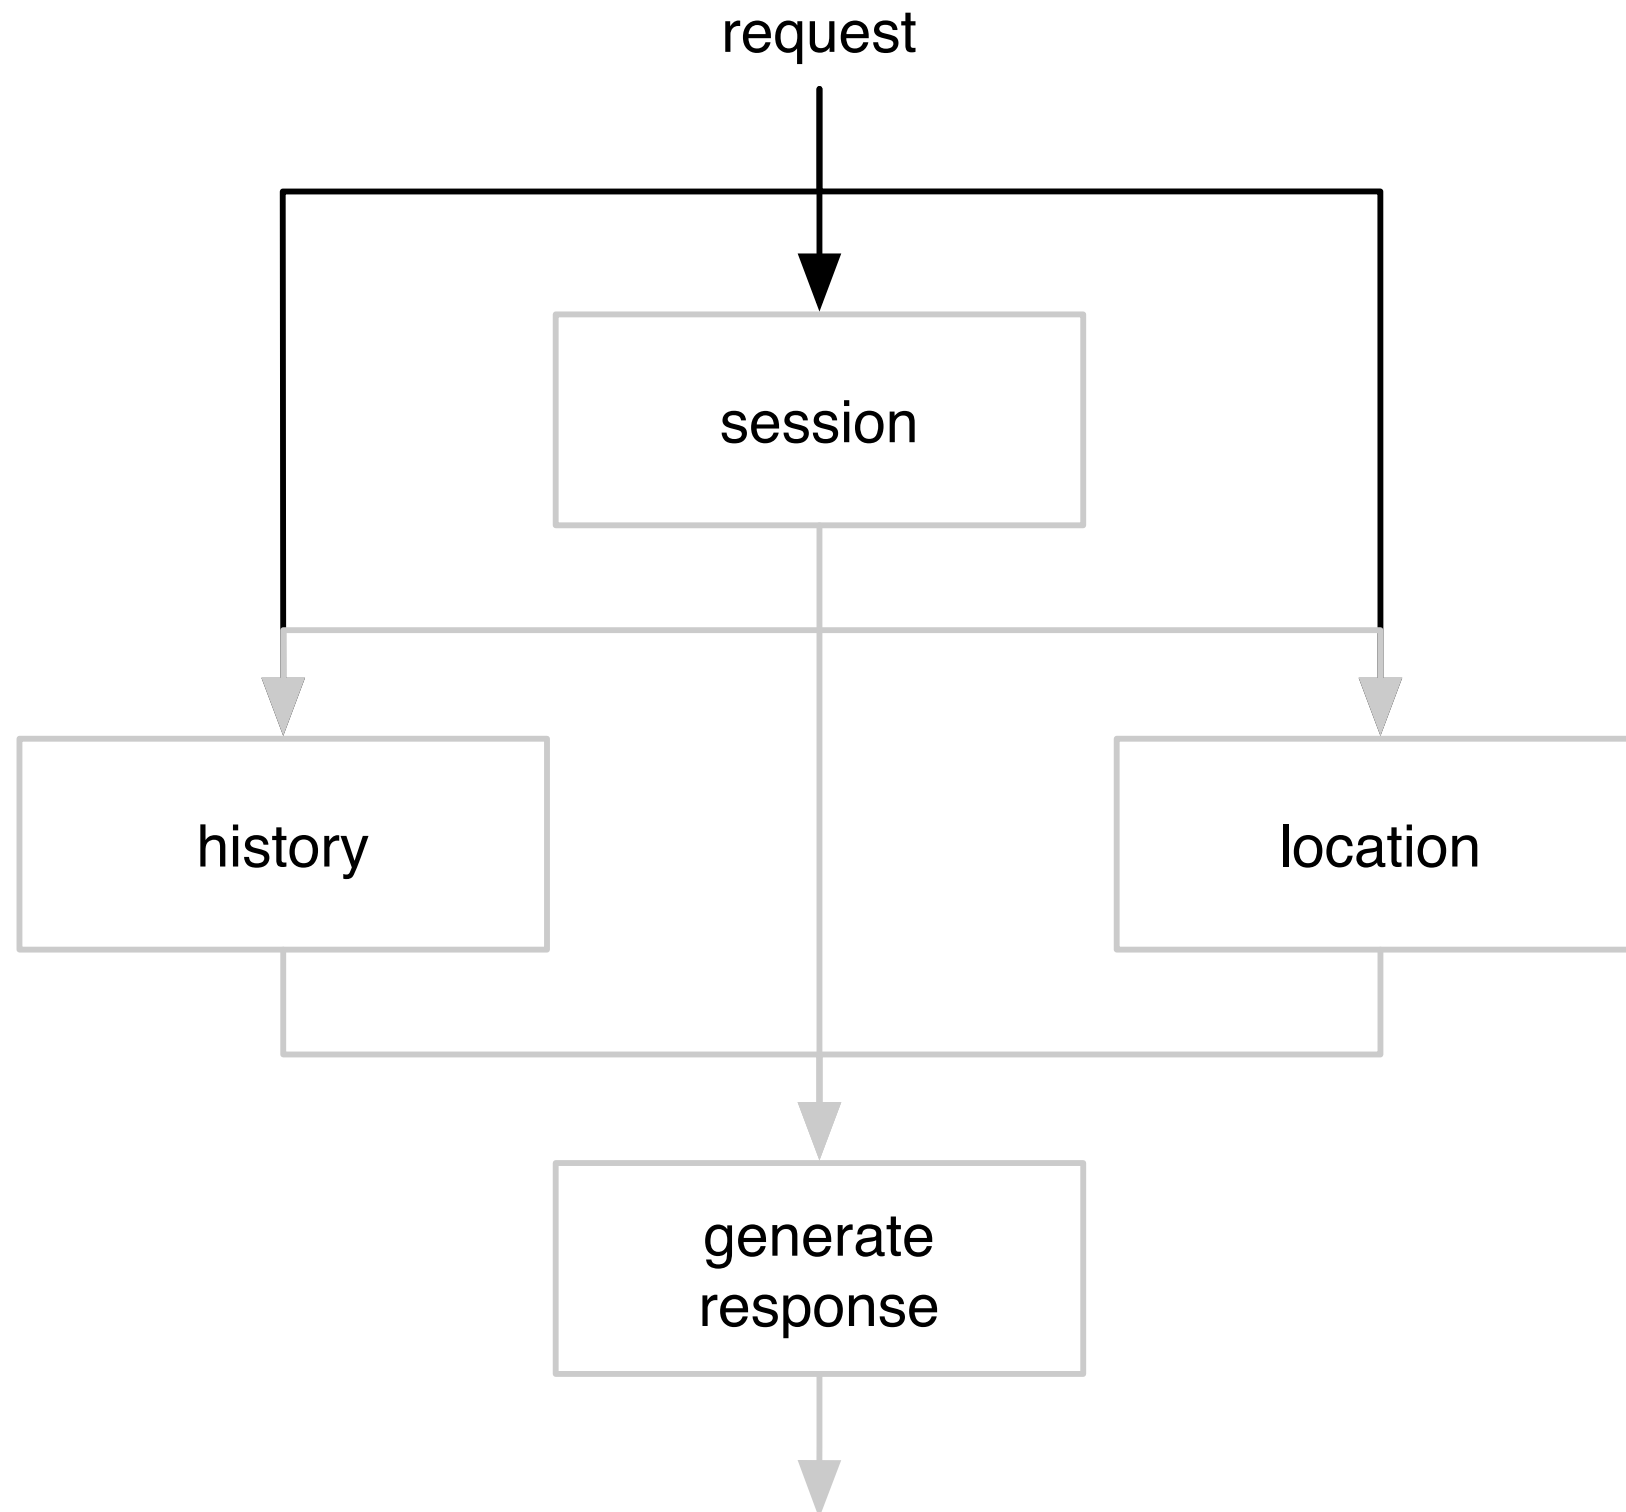
```
(defn handler [request]
  (let [session (get-session request)
        history (or (:history session)
                     (get-history request))
        location (or (:location session)
                     (get-location request))]
    (generate-response
     request
     session
     history
     location)))
```

# flow of execution

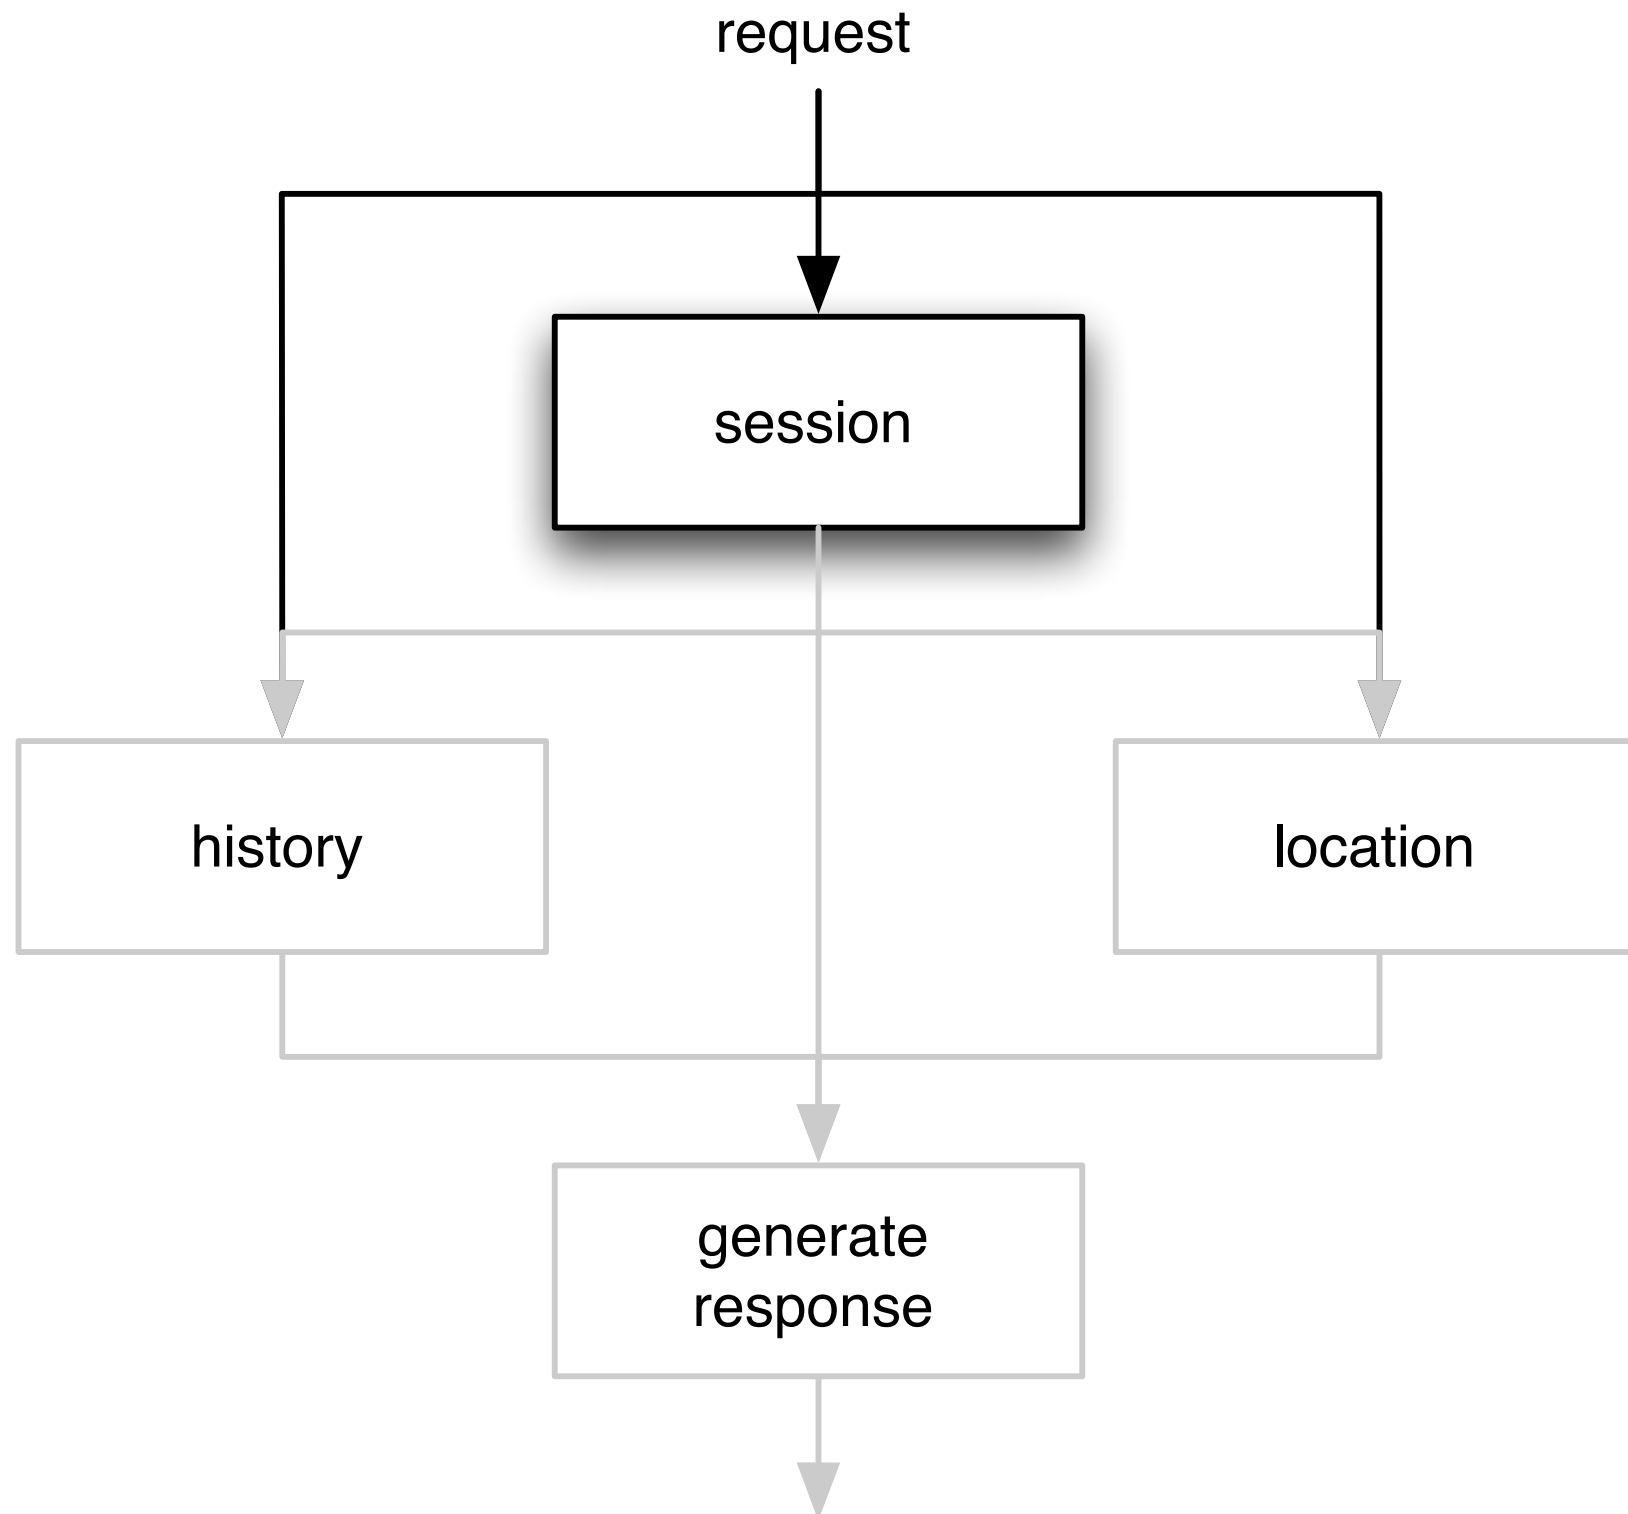




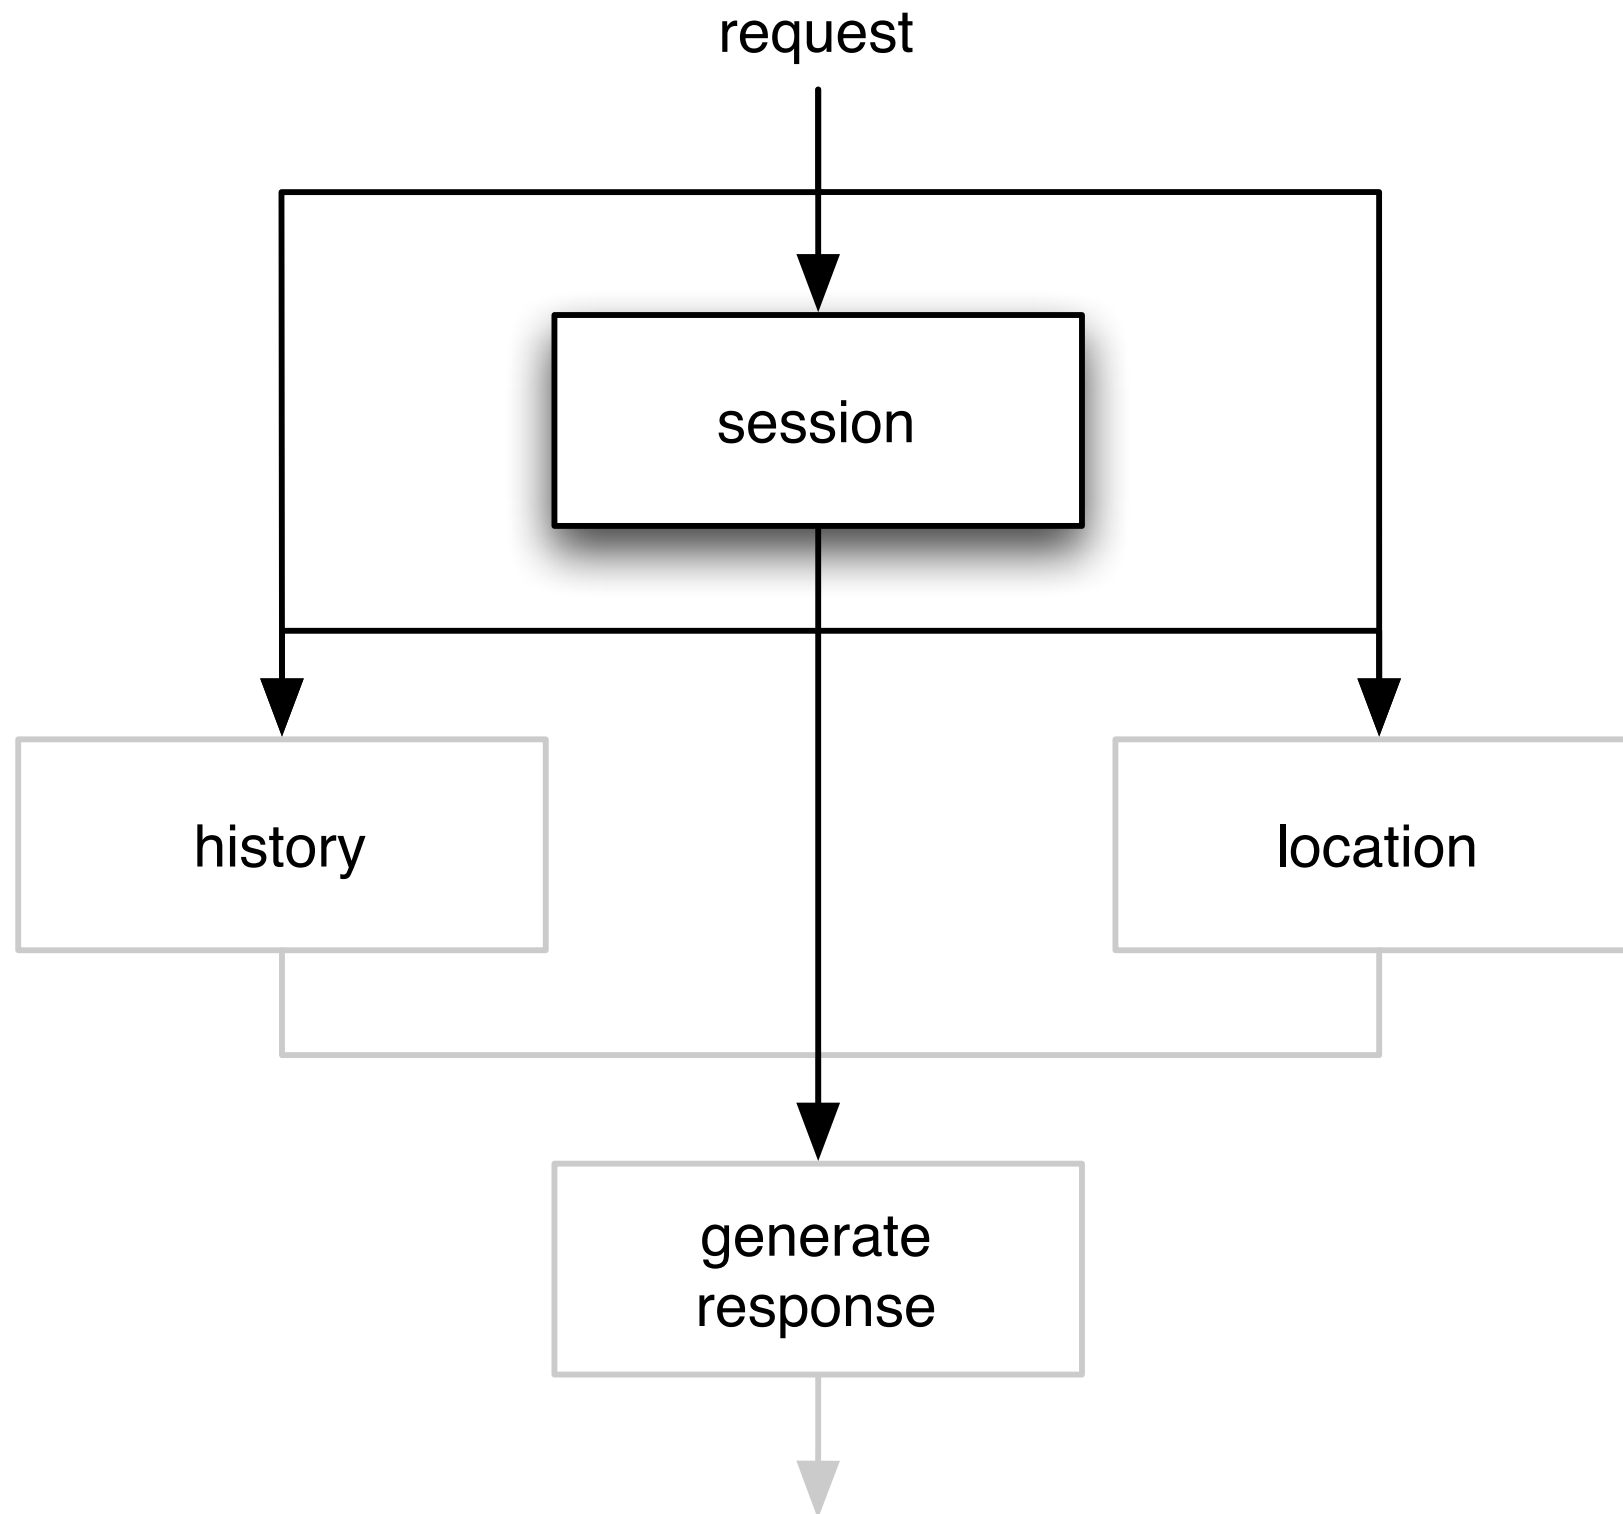
# flow of data



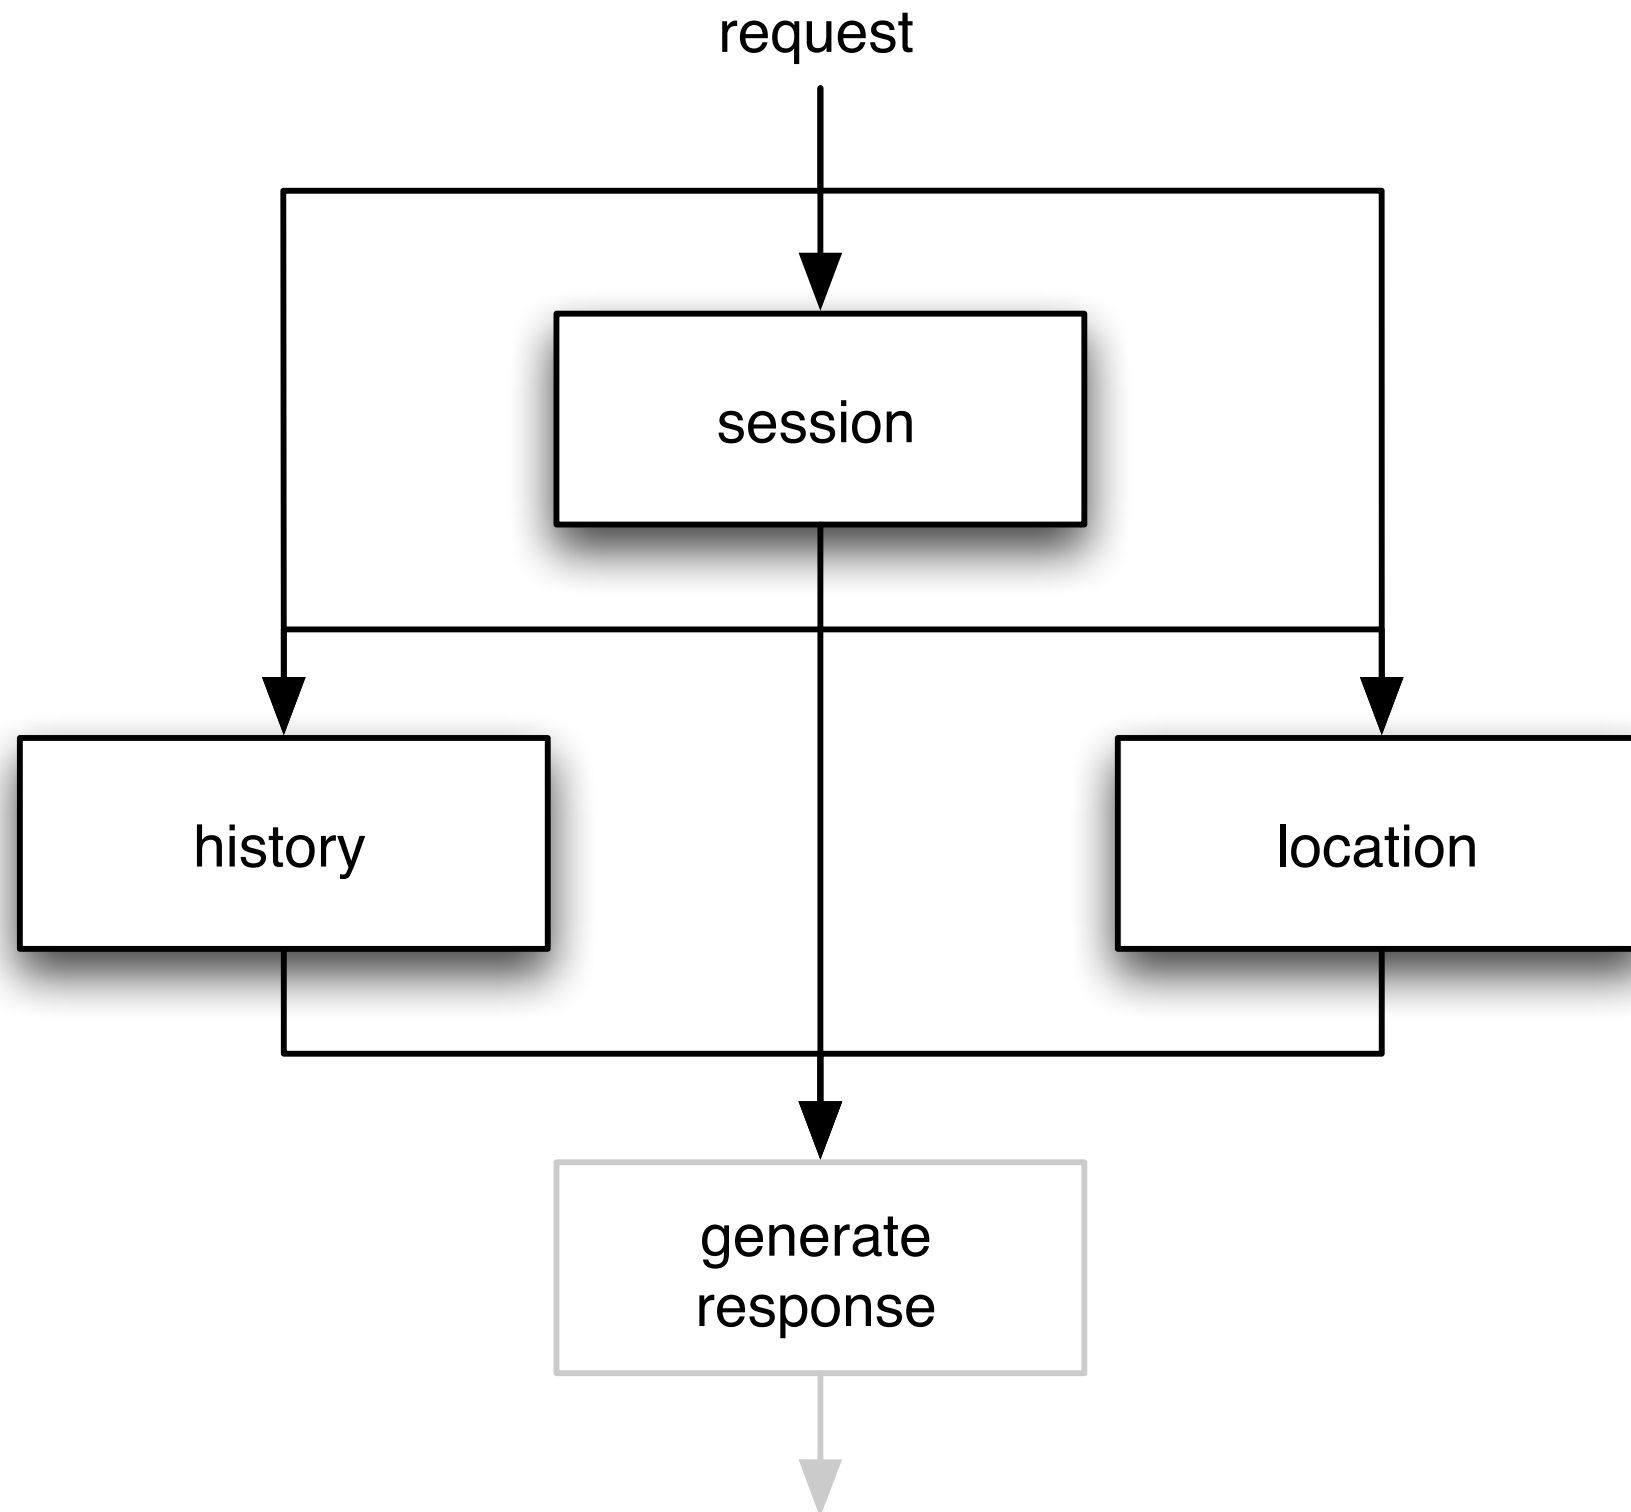
# flow of data



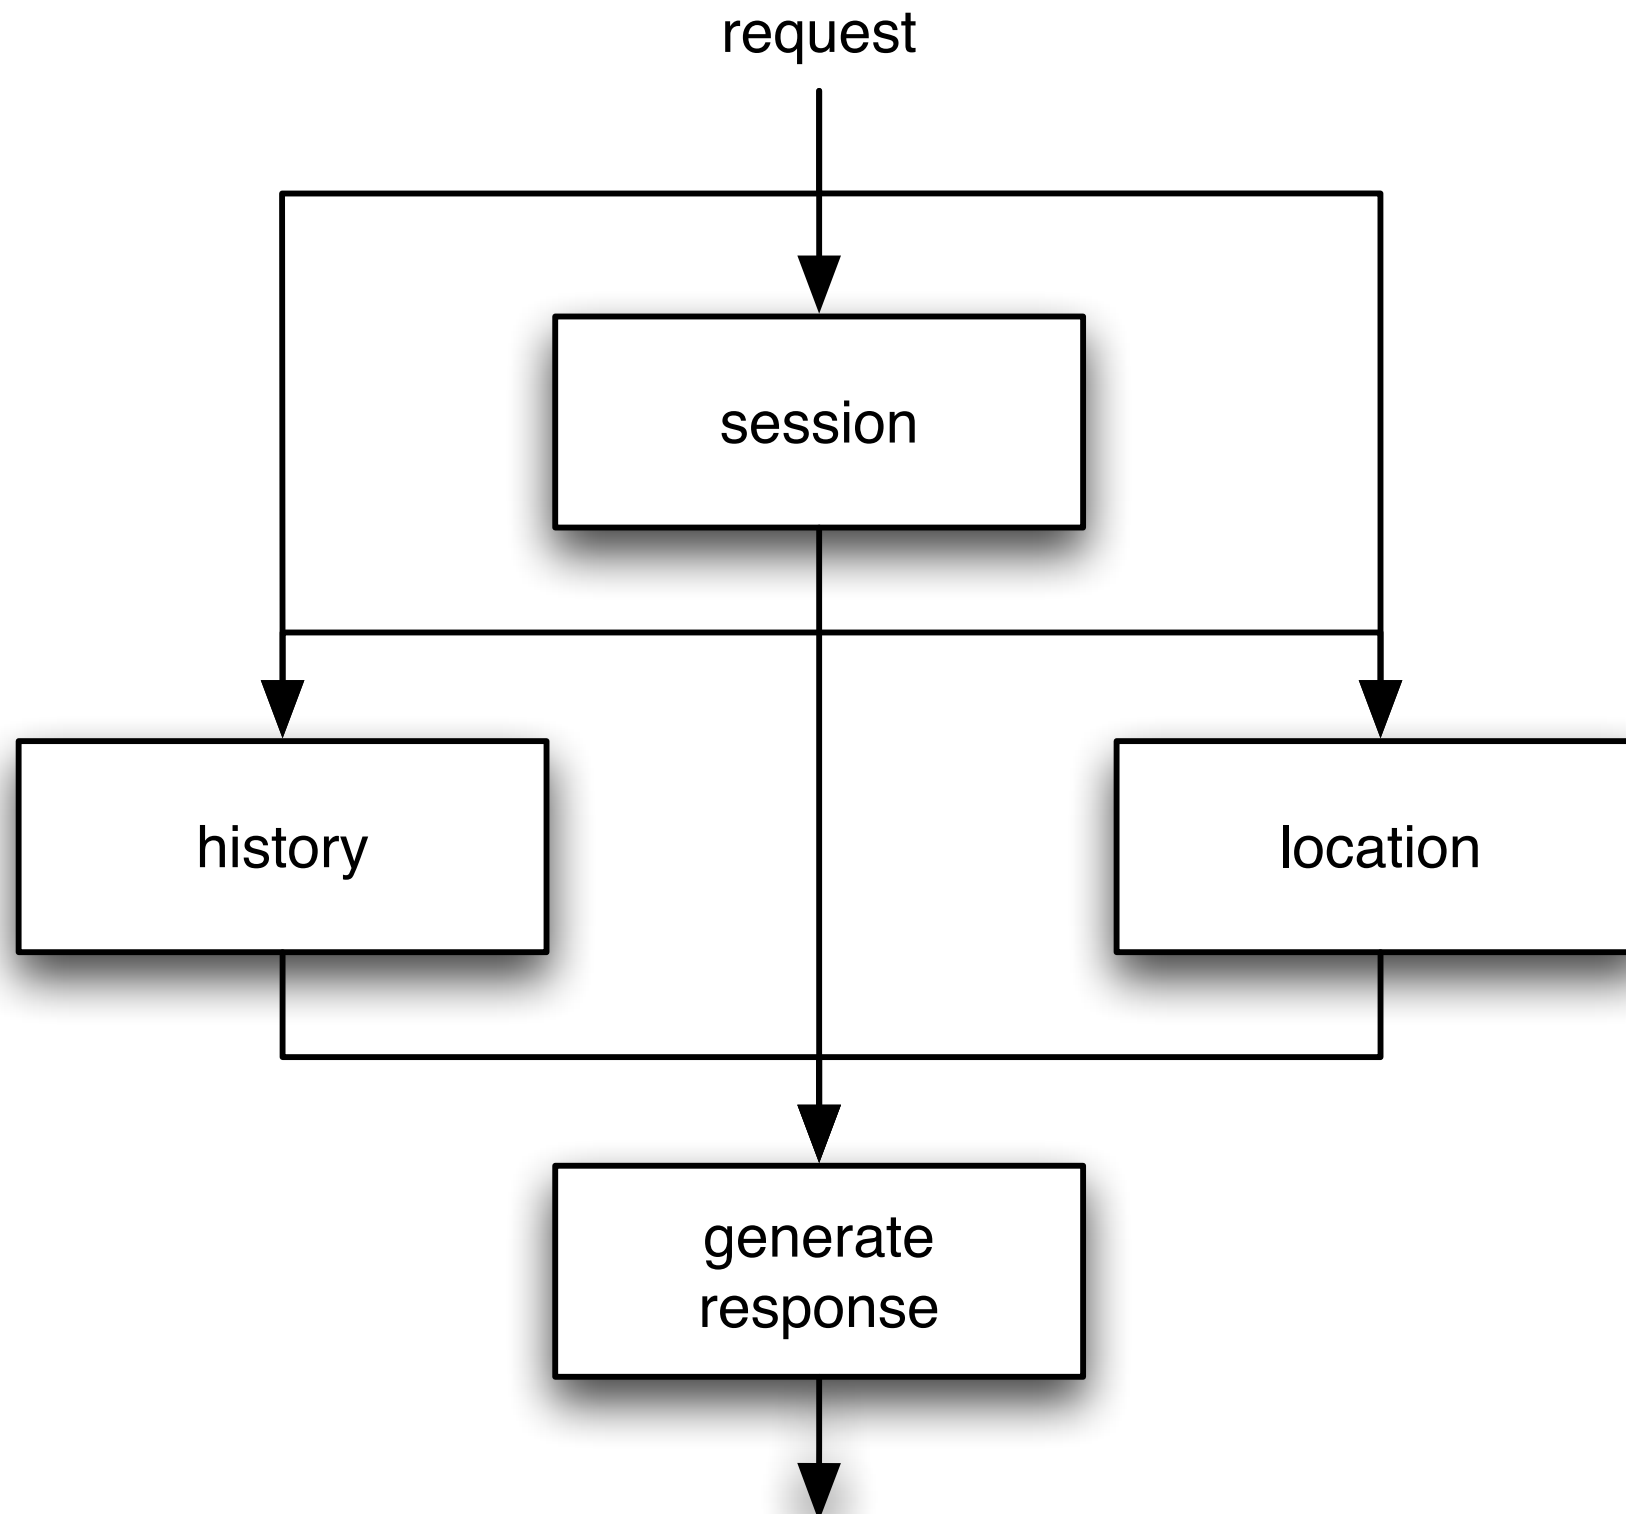
# flow of data



# flow of data



# flow of data



# flow of data

```
(defn handler [request]
  (let [session (get-session request)
        history (or (:history session)
                     (get-history request))
        location (or (:location session)
                     (get-location request))]
    (generate-response
     request
     session
     history
     location)))
```

the code is a perfect expression of the dataflow!

# the async macro

- expressions immediately return a result-channel representing the eventual result
- execution of expressions are deferred until all parameters are realized
- if a parameter emits an error, the expression emits the same error

# the async macro

```
(defn handler [request]
  (async
    (let [session (get-session request)
          history (or (:history session)
                      (get-user-history request))
          location (or (:location session)
                       (get-user-location request))]
      (generate-response
        request
        session
        history
        location))))
```



# the async macro

```
(defn handler [request]
  (async
    (let [session (get-session request)
          history (or (:history session)
                      (get-user-history request))
          location (or (:location session)
                       (get-user-location request))]
      (generate-response
        request
        session
        history
        location))))
```

# the async macro

```
(defn handler [request]
  (async
    (let [session (get-session request)
          history (or (:history session)
                      (get-user-history request))
          location (or (:location session)
                       (get-user-location request))]
      (generate-response
        request
        session
        history
        location))))
```

# the async macro

```
(defn handler [request]
  (async
    (let [session (get-session request)
          history (or (:history session)
                     (get-user-history request))
          location (or (:location session)
                      (get-user-location request))]
      (generate-response
        request
        session
        history
        location))))
```

# the async macro

```
(defn handler [request]
  (async
    (let [session (get-session request)
          history (or (:history session)
                      (get-user-history request))
          location (or (:location session)
                       (get-user-location request))]
      (generate-response
        request
        session
        history
        location))))
```

# the async macro

```
(defn handler [request]
  (async
    (let [session (get-session request)
          history (or (:history session)
                     (get-user-history request))
          location (or (:location session)
                      (get-user-location request))]
      (generate-response
        request
        session
        history
        location))))
```

# the async macro

```
(defn handler [request]
  (async
    (let [session (get-session request)
          history (or (:history session)
                     (get-user-history request))
          location (or (:location session)
                      (get-user-location request))]
      (generate-response
        request
        session
        history
        location))))
```

# the async macro

```
(defn handler [request]
  (async
    (let [session (get-session request)
          history (or (:history session)
                     (get-user-history request))
          location (or (:location session)
                      (get-user-location request))]
      (generate-response
        request
        session
        history
        location))))
```

# async try/catch

```
(async  
  (try  
    (handler request)  
    (catch Exception e  
      (handle-error e))))
```

error handling doesn't have to be local!



# task

```
(async  
  (= 3 (task (calculate-pi))))
```

- returns result-channel, executes body on thread-pool
- is just an **annotation**, not a structural change
- because async expressions will return when they can't continue, deadlock is impossible

# async side-effects

```
(async  
  (do  
    (query-moon-base)  
    (query-database)))
```

- the moon base's response isn't part of our dataflow
- if there's an error, how do we know?

# force

```
(async  
  (do  
    (force (query-moon-base))  
    (query-database)))
```

- makes an expression a pre-requisite for all subsequent expressions
- an error will cause subsequent expressions to also error out

# however, async is opaque

- no way to examine partially complete flows
- difficult to debug
- complicated flows are still complicated, under the covers

# future work

- improved debugging for async workflows
- improved instrumentation for distributed systems
- ClojureScript!

# the projects

- all functionality discussed is contained in **lamina** (<http://github.com/ztellman/lamina>)
- it is used to model network communication in **aleph** (<http://github.com/ztellman/aleph>)

questions?